

Geofinance: Following Sovereign Capital Home*

Ran Zhou[†]

May 20, 2026

Abstract

Sovereign wealth funds (SWFs) invested around \$1.28 trillion in cross-border private markets between 2002 and 2022, but the geographic effects of this capital remain poorly understood. I develop a geofinance framework in which a financier offers favorable terms to induce business toward its home country, building bilateral links that reduce frictions for outside investors. The model predicts stronger directed activity when recipients have weaker outside financing options and when the financier places greater strategic value on the sector. Following SWF private investments, firms with lower relative valuations expand more toward the SWF home country than toward non-SWF co-investor home countries: a one-standard-deviation decrease in relative valuation raises this gap by 4.8 percentage points, about 70 percent of the baseline rate. The pattern is not explained by firm scale or ownership dilution, and the response is stronger in sectors where SWF home countries have comparative disadvantages, not comparative advantages. Similar redirection appears in lending to SWF home countries by systemically important banks during the 2007–2008 financial crisis. At the aggregate level, FDI flows and cross-border banking claims toward SWF home countries rise by about 5–8 percent.

1 Introduction

In May 2025, Saudi Arabia’s Public Investment Fund (PIF) and Google Cloud announced a \$10 billion joint investment to build an artificial intelligence hub in Saudi Arabia.¹ Finance flowed one

*I thank Mikhail Chernov, Barney Hartman-Glaser, Valentin Haddad, Erica Jiang, Shohini Kundu, Francis Longstaff, Tyler Muir, Stavros Panageas, Gregor Schubert, Olav Sorenson, Ivo Welch, Yujie Liu, and Jinyuan Zhang for helpful comments and valuable discussions. I gratefully acknowledge financial support from the Fink Center for Finance and Investments and the Price Center for Entrepreneurship & Innovation.

[†]UCLA Anderson School of Management, University of California, Los Angeles. Email: ran.zhou.phd@anderson.ucla.edu.

¹Google Cloud and PIF Advance AI Hub in Saudi Arabia

way while economic activity followed the other. Such deals are increasingly common. Sovereign wealth funds (SWFs), which collectively manage over \$12 trillion, deployed roughly \$1.28 trillion in cross-border private-market investments over 2002–2022.² Does this Google–PIF pattern generalize? This paper studies when and how SWF private investments redirect recipients’ cross-border activity toward the SWFs’ home countries, and whether the bilateral links created by that redirection crowd in outside capital.

I develop a geofinance framework in which a financier offers favorable financing terms to induce cross-border activity toward its home country. The term distinguishes this mechanism from existing geoeconomic frameworks, which emphasize leverage through existing trade and financial dependencies (Clayton, Maggiori and Schreger, 2025, 2026). Geofinance instead studies how capital itself can be leveraged to induce new bilateral links. SWFs are the empirical focus, but the mechanism is not unique to sovereign capital. It applies whenever a financier is willing to trade part of its financial return for directed real activity or another strategic objective.³

The mechanism operates through a financing rent: the transfer created by favorable financing terms relative to the recipient’s outside option. The financier uses this rent to induce business toward its home market. The direction of induced activity depends on where the financier places strategic value, while the feasibility of sustaining that activity depends on the recipient’s outside financing options: recipients with weaker outside financing options value the financing rent more, so directed activity above the commercial benchmark is easier to sustain. Directed activity also adds to a stock of bilateral links. As this stock rises, bilateral frictions fall, making subsequent investment in the financier’s home country more attractive to outside investors.

The model delivers two empirical implications. First, SWF capital should induce recipients to do more business in the SWF home market than their commercial benchmark would imply, especially on the new-entry margin where bilateral links are formed. The effect should be stronger when recipients have weaker outside financing options and when the SWF places greater value on the relevant sector or market. Second, if directed activity builds bilateral links, it should crowd in third-party capital toward SWF home countries. I test these implications in four settings: private-market investments, crisis-era bank recapitalizations, bilateral foreign direct investment (FDI), and

²For scale, the \$12 trillion figure is comparable to the \$12.7 trillion in total loans and leases held by U.S. commercial banks (Federal Reserve H.8 release, March 2025). The \$1.28 trillion global figure is computed from GSWF and SWFI transaction databases for 2002–2022. SWF investments into U.S. targets alone (\$373 billion) equal about 6.3% of cumulative U.S. FDI net inflows (\$5.9 trillion, World Bank Group).

³SWFs are a natural application because funding from commodity revenues, foreign-exchange reserves, or fiscal surpluses lowers their shadow cost of capital. The mechanism also applies to other financiers, such as development finance institutions or strategic corporate investors, that accept lower financial returns in exchange for influence over real decisions.

cross-border banking claims. The first two settings test micro-level business redirection and its heterogeneity; the last two test whether those directed links are followed by broader capital-flow responses.

I first test the mechanism in 833 private-market deals. The main outcome is new market entry: whether the firm records a first post-financing transaction in the investor's home country. I compare the same firm's expansion into SWF and non-SWF investor home countries within a single financing round, holding fixed the recipient firm and financing event. Within this design, firms are not more likely on average to enter the SWF home country than non-SWF co-investor home countries. Instead, the differential appears among firms raising capital at lower valuations relative to sector-year peers: a one-standard-deviation decrease in relative valuation raises the SWF–non-SWF entry gap by 4.8 percentage points, relative to a 6.8% non-SWF baseline. This pattern is not driven by firm scale or ownership dilution, and I interpret lower relative valuations as evidence of weaker outside financing options.

The financier-side evidence reveals the strategic motive of SWFs. The expansion effect is stronger in sectors that the investing SWF had overweighted in its prior cross-border portfolio: in sectors with a 10 percentage-point prior overweight, a one-standard-deviation valuation discount implies a 6.8 percentage-point SWF–non-SWF entry gap. The effect is also larger in sectors where the SWF home country lacks export comparative advantage or relies more heavily on imports. Together, these results match the model's first implication: redirection is easier to sustain when recipients have weaker outside financing options and when financiers have a stronger reason to create activity at home.

This within-round comparison holds the financing event fixed, but it benchmarks the SWF country against same-round non-SWF investor countries, often large destination markets such as the United States and the United Kingdom. I therefore also estimate a firm-level counterfactual that holds the potential SWF destination fixed. SWF-backed firms are 6.3 percentage points more likely to enter the SWF country than similar non-SWF-backed firms. This comparison asks whether similar firms would have entered the same SWF destination absent SWF backing; the positive gap supports the model's level prediction that SWF capital induces activity above the recipient's commercial benchmark.

The bank setting provides sharper variation in the recipient-side condition behind the model: demand for outside capital. During the 2007–2008 financial crisis, banks with larger pre-crisis write-down exposure had greater demand for outside capital, and SWFs became a major source of recapitalization.⁴ Using syndicated lending data with bank-by-country, bank-by-quarter, and

⁴The International Monetary Fund's *Global Financial Stability Report* (April 2008). Public offerings and other

country-by-quarter fixed effects, I find that high-write-down banks redirected lending toward SWF countries after the crisis cutoff. A one-standard-deviation increase in write-down severity is associated with about 19% higher deal counts among bank–country–quarter cells with positive lending.

I then test the crowd-in prediction in bilateral foreign direct investment (FDI). I exploit the 2018 Foreign Investment Risk Review Modernization Act (FIRRMA), which tightened U.S. scrutiny of foreign government-linked investments ([Zhang and Zhou, 2025](#)). SWF countries with greater pre-2018 exposure to U.S. high-equity-stake deals subsequently attract more third-party FDI from non-U.S. source countries. Conditional on positive flows, observed exposure levels imply roughly 5–8% higher post-FIRRMA FDI inflows relative to SWF destinations with zero exposure. This magnitude is comparable to benchmark estimates of the bilateral FDI effect of signing a bilateral investment treaty.⁵

The banking-claims evidence points in the same direction. Using the Bank for International Settlements (BIS) Locational Banking Statistics, I trace lender-country claims on SWF home countries after SWF equity injections into banks in a staggered-cohort design. Each additional \$1 billion of SWF equity is associated with a 5.8% increase in loans and placed deposits to SWF countries. Together, the FDI and banking-claims results match the model’s second implication: SWF-backed relationships are followed by broader bilateral capital-flow responses.

Sovereign capital transactions are large, strategic, and infrequent, so no single setting can fully isolate every margin of the mechanism.⁶ I therefore read the four designs together. The common pattern is that SWF capital is followed by activity toward SWF home markets. At the micro level, the response is strongest when recipients have weaker outside financing options and when SWFs place greater value on the relevant sector or market. At the aggregate level, FDI and banking-claims outcomes subsequently move toward SWF home countries.

Related literature. The paper speaks first to geoeconomics. Recent theory shows how states project power through existing trade and financial relationships ([Clayton, Maggiori and Schreger, 2025, 2026](#)), and network accounts emphasize coercion through control over hubs ([Farrell and Newman, 2019](#)). Related empirical work shows that state influence and state-backed finance can redirect economic activity: political influence redirected trade toward American products ([Berger et al., 2013](#)), China manages foreign access to its bond market ([Clayton et al., 2025](#)), and export-

institutional investors were the other major sources of funding.

⁵[Egger and Merlo \(2007\)](#) estimate that bilateral investment treaties raise bilateral FDI stocks by about 4.8% in the short run and 8.9% in the long run.

⁶In the combined GSWF and SWFI transaction data, restricting to cross-border unlisted deals over 2002–2022 and excluding fund-of-fund rows, the median active SWF-year has two deals in two target countries.

credit agencies raise exports where financing and contracting frictions bind (Kabir et al., 2025). These papers study pressure through existing dependencies. I study formation: state capital creates new bilateral links that can later support private investment.

The second literature studies SWFs and state capitalism. Existing work describes SWF investment strategies, political motives, and differences between government-linked and private acquirers (Chhaochharia and Laeven, 2008; Bernstein, Lerner and Schoar, 2013; Knill, Lee and Mauck, 2012; Karolyi and Liao, 2017). A related strand estimates target-firm announcement, governance, and operating effects (Dewenter, Han and Malatesta, 2010; Kotter and Lel, 2011; Bortolotti, Fotak and Megginson, 2015); see Megginson and Fotak (2015) for a survey. More recent work shows that SWF investments increase home-country exports after FIRMA (Zhang and Zhou, 2025). My paper shifts the outcome from target-firm value or trade to the geography of subsequent activity. It asks when SWF capital redirects firm expansion and cross-border syndicated lending, and whether that crowds in FDI and banking claims toward the SWF home country.

The third literature is international finance. Cross-border flows are shaped by gravity, information, institutional, and control frictions (Anderson and van Wincoop, 2003; Portes and Rey, 2005; Wei, 2000; Head and Ries, 2008; Daude and Fratzscher, 2008; Pellegrino, Spolaore and Wacziarg, 2025). But these frictions need not be fixed. Chaney (2014) shows that trade networks lower search costs and make future trade easier. I bring this logic to cross-border capital allocation. SWF-backed relationships are a channel through which bilateral financial frictions can fall and later third-party capital can crowd in.

The remainder of the paper is organized as follows. Section 2 develops the geofinance model and derives the main empirical predictions. Sections 3 and 4 present the micro evidence from PitchBook and DealScan. Section 5 studies the macro crowd-in effects using FDI and BIS banking claims data. Section 6 concludes.

2 A Framework for Geofinance

I model geofinance as strategic capital used to induce real activity toward the financier's home market. Firms differ in their outside financing options: firms facing costlier outside finance value financing rents more and are therefore more willing to depart from their commercial benchmark. Financiers differ in the strategic value they attach to strategic-market activity. Geofinance arises when transferable financing surplus supports activity that the financier values but that the firm

would not choose at the same level under purely commercial finance.⁷

Ordinary investors can supply capital to the same firms, but their capital need not pull activity toward their home countries. A strategic financier can pair capital with future financing opportunities and a country-specific relationship channel. The model separates receptiveness from destination: the recipient’s outside financing options govern how valuable the financing rent is, while the financier’s strategic value governs where the induced activity is directed.

2.1 Setup

Let $\mathcal{I}$ be a finite set of potential recipient firms. I summarize firm i ’s outside financing options by an effective marginal outside financing cost $r_i > 0$: the cost of raising marginal capital from non-strategic financiers. A high r_i means the firm has weaker outside financing options: alternative financing terms are costly because of risk, opacity, urgency, limited investor competition, or other outside-option premia. It does not mean that the firm cannot raise funds. For tractability, I summarize firm i ’s ordinary activity by a pre-optimized value $A_i(r_i)$, which is additively separable from strategic-market deployment and drops out of the marginal conditions below.⁸

Let s index the financier’s home market. The firm chooses strategic-market activity, $k_{is} \geq 0$, with return $G_i(k_{is})$. This activity carries a proportional cost $\mu(Z)k_{is}$, where Z is the stock of bilateral links. Since $\mu'(Z) < 0$, accumulated links lower the cost of later activities in the financier’s home market.⁹ I assume $A_i(r_i)$ is well behaved; G_i is twice differentiable, strictly increasing, and strictly concave; $\mu(Z) > 0$ and $\mu'(Z) < 0$; and $U(Z)$ is increasing and weakly concave.

I begin with the firm’s allocation under purely commercial finance. Taking Z as given, the firm solves

$$\max_{k_{is} \geq 0} \{A_i(r_i) + G_i(k_{is}) - \mu(Z)k_{is} - r_i k_{is}\}. \quad (1)$$

⁷The framework draws on the geoeconomic setting of [Clayton, Maggiori and Schreger \(2025, 2026\)](#). They formalize geoeconomic power through trade networks and threats. My framework treats transferable financing surplus as a way to induce directed activity. I integrate the principal-multi-agent contracting structure of [Segal \(1999\)](#), the public-liquidity logic of [Holmström and Tirole \(1998\)](#), and the “large trader” coordination mechanism of [Corsetti et al. \(2004\)](#).

⁸ r_i can include the firm’s risk-adjusted cost of capital, outside-option premia from information or agency frictions, regulatory pressure, distress, or market segmentation. In the bank application, this cost reflects balance-sheet distress and regulatory leverage pressure; in the corporate application, it reflects costly outside equity or weak outside financing options. One microfoundation is $A_i(r_i) \equiv \max_{k_{ih} \geq 0} \{F_i(k_{ih}) - r_i k_{ih}\}$, where k_{ih} represents ordinary activity outside the strategic market.

⁹The model treats $\mu(Z)$ as a reduced-form bilateral relationship friction. It can reflect search, information, contracting, or coordination costs that decline as bilateral activity accumulates.

For an interior solution, the first-order condition for strategic-market activity is

$$G'_i(k_{is}^{\text{out}}) = \mu(Z) + r_i. \quad (2)$$

Activity in the financier's home market is disciplined by two costs: the bilateral operating friction $\mu(Z)$ and the firm's outside financing cost r_i . Holding Z fixed, firms with weaker outside financing options choose less strategic-market activity.¹⁰

2.2 The contract and the firm's response

A strategic financier has cost of funds $\rho \geq 0$ and can deliver a transferable financing rent. It offers a strategic-finance bundle $\{x_i, r_{is}, k_{is}\}$: capital $x_i \geq 0$, an effective package term r_{is} , and strategic-market activity k_{is} . The effective package term satisfies

$$r_{is} < r_i, \quad (3)$$

conditional on strategic-market activity. The rate notation is a convenient way to express the transfer; the model does not require r_{is} to be a separately observed investor-specific price. It can represent a lower contracted rate, favorable security terms, guarantees, follow-on access, or other package features.¹¹ With x_i units of finance, the package delivers financing rent $(r_i - r_{is})x_i$. The gross joint surplus from replacing outside finance is $(r_i - \rho)x_i$, divided between recipient rent and financier margin.¹²

The financing rent determines whether the firm accepts the bundle. The financier's home-country relationship determines where the response is directed.¹³ This relationship can include local customers, procurement, regulatory navigation, certification, local partners, or repeated sovereign financing. I summarize the realized effect of the relationship on the real allocation by the implementation wedge τ_{is} . Ordinary strategic-market activity is the limiting case in which

¹⁰Low observed activity in the strategic market does not mean that projects there have low gross returns. The home-bias literature emphasizes that foreign activity is limited by information, familiarity, monitoring, contracting, and institutional frictions. These additional frictions are summarized by $\mu(Z)$. Firms may earn positive returns abroad and price risk into those activities, but without strategic capital they may still choose a smaller foreign footprint than the financier would like.

¹¹In a same-round equity syndicate, co-investors may share a headline valuation. The model's r_{is} should be read as an effective package term that delivers financing surplus, not as a claim that the financier always pays a different observed price from commercial co-investors.

¹²The same rent could be delivered through debt, equity, guarantees, direct transfers, or future access. In the transfer-separable benchmark, the instrument does not affect the implemented allocation.

¹³Commercial co-investors may provide capital and general expertise, but they do not systematically value the same home-country link.

the financier invests on market terms, creates no strategic wedge, and the firm remains at its commercial benchmark. Strategic finance differs because the relationship changes the relative payoff to strategic-market activity and the rent makes the bundle acceptable. I interpret k_{is} as enforceable through project finance, ring-fencing, covenants, regulatory commitments, relational contracting, or other observable deployment restrictions.

The firm's payoff under the contract decomposes into operating value and a financing rent:

$$\Pi_i(x_i, r_{is}, k_{is}, Z) = \underbrace{\text{Fund}_i(k_{is}, Z)}_{\text{operating value given } k_{is}} + \underbrace{(r_i - r_{is}) x_i}_{\text{financing rent}} \quad (4)$$

where

$$\text{Fund}_i(k_{is}, Z) \equiv A_i(r_i) + G_i(k_{is}) - \mu(Z)k_{is} - r_i k_{is}. \quad (5)$$

The rent term is the transferable part of the bundle. It compensates the firm for moving away from its commercial outside option. The country-specific relationship is not a generic cash transfer; it enters the real allocation through the implementation wedge below. The firm accepts the contract whenever

$$\Pi_i(x_i, r_{is}, k_{is}, Z) \geq v_i(Z), \quad (6)$$

where the outside option under commercial finance is

$$v_i(Z) = \max_{k_{is} \geq 0} \{A_i(r_i) + G_i(k_{is}) - \mu(Z)k_{is} - r_i k_{is}\}. \quad (7)$$

The implemented activity can be represented by an implicit marginal wedge τ_{is} :

$$G'_i(k_{is}) = \mu(Z) + r_i - \tau_{is}. \quad (8)$$

At a given Z , a positive τ_{is} lowers the effective marginal cost of strategic-market activity and induces $k_{is} > k_{is}^{\text{out}}$. The wedge combines home-market access, strategic requests, and enforceability. Financing alone does not imply directed activity: if $\tau_{is} = 0$, the firm may accept capital but remains at k_{is}^{out} .

Finally, let $C \subseteq \mathcal{I}$ denote the contracted set of firms. I take this set as given in the benchmark contracting problem and characterize optimal contracts conditional on C . Firms outside C choose their commercial strategic-market activity given the network stock. The aggregate network stock is

$$Z = \sum_{i \in C} k_{is} + \sum_{j \notin C} k_{js}^*(Z), \quad (9)$$

where the second term captures crowd-in. Under the maintained assumptions, outside firms respond positively to a larger network stock:

$$\frac{\partial k_{js}^*}{\partial Z} = \frac{\mu'(Z)}{G_j''(k_{js}^*)} > 0. \quad (10)$$

For local stability of the network fixed point, I impose

$$0 \leq \sum_{j \notin C} \frac{\partial k_{js}^*}{\partial Z} < 1. \quad (11)$$

2.3 The financier's problem and main results

The financier acts as a Stackelberg leader conditional on a contracted recipient set C . It values the bilateral network stock Z , possibly with different weights across sectors. With strategic-objective weight $\theta > 0$, aggregate funding capacity $\bar{X}$, and convex implementation cost $\frac{\epsilon}{2}x_i^2$, it chooses contracts $\{x_i, r_{is}, k_{is}\}_{i \in C}$ and the resulting network stock Z :

$$\begin{aligned} \max_{\{x_i, r_{is}, k_{is}\}_{i \in C}, Z} \quad & \sum_{i \in C} [(r_{is} - \rho)x_i - \frac{\epsilon}{2}x_i^2] + \theta U(Z) \\ \text{subject to:} \quad & \Pi_i(x_i, r_{is}, k_{is}, Z) - v_i(Z) \geq 0, \quad \forall i \in C, \quad [y_i] \\ & \bar{X} - \sum_{i \in C} x_i \geq 0, \quad [\lambda] \\ & \sum_{i \in C} k_{is} + \sum_{j \notin C} k_{js}^*(Z) - Z = 0, \quad [\Psi_Z] \\ & x_i \geq 0, \quad k_{is} \geq 0, \quad \forall i \in C. \end{aligned} \quad (12)$$

All inequality multipliers are nonnegative. The equality multiplier Ψ_Z is unrestricted. I use $\Psi_Z > 0$ for the case in which the financier has a positive shadow value of expanding the bilateral network. The first-order conditions characterize contracted recipients with $x_i > 0$, interior k_{is} , and an interior contract rate, conditional on C .

In the transfer-separable benchmark, the contract rate divides surplus in the interior pricing regime. Only two terms involve r_{is} : $(r_{is} - \rho)x_i + y_i(r_i - r_{is})x_i$. The first-order condition gives $x_i(1 - y_i) = 0$, so for active funding $x_i > 0$, participation binds with $y_i = 1$. The effective financing term r_{is} allocates the joint surplus between the financier and the recipient; it does not determine funding volume or directed activity.¹⁴

¹⁴The cancellation of r_{is} does not require literal risk neutrality by the recipient. It follows from transfer separability: the contract rate and funding volume enter the recipient's payoff through the financing rent $(r_i - r_{is})x_i$. With transfer

The remaining first-order conditions yield two main results, stated below as propositions. Appendix A provides the full derivations.

Proposition 1 (Outside financing options and directed relationships). *For any contracted recipient satisfying the interior benchmark conditions,*

$$\tau_{is}^* = \Psi_Z, \quad x_i^* = \frac{r_i - \rho - \lambda}{c}. \quad (13)$$

Therefore, holding Z fixed, the interior contract induces $k_{is} > k_{is}^{out}$ whenever $\Psi_Z > 0$. Active funding for a contracted recipient requires

$$r_i > \rho + \lambda. \quad (14)$$

Weaker outside financing options make the financing rent more valuable and make the active-funding condition easier to satisfy. They make the firm more receptive to a strategic bundle, but the geography of the response comes from the financier's home-country relationship.

Intuition. The home-market relationship determines which activity the financier wants to induce. The financing package supplies the participation rent that makes the requested activity acceptable to the firm. Once participation binds, r_{is} cancels from the funding first-order condition: it allocates surplus, while funding volume depends on the value of replacing outside finance. Thus r_i is a receptiveness margin, not a destination margin. The implementation wedge is pinned down by Ψ_Z , while the level of directed activity depends on Ψ_Z , $\mu(Z)$, and the recipient's outside financing options.

Bilateral efficiency. The allocation is a Coasean outcome between the firm and the financier. The firm alone chooses k_{is}^{out} , satisfying $G'_i(k_{is}^{out}) = \mu(Z) + r_i$. A financier that values the bilateral link adds Ψ_Z , so the contracted allocation satisfies $G'_i(k_{is}^*) + \Psi_Z = \mu(Z) + r_i$. The contract moves the firm from its commercial benchmark to the firm-financier contracting optimum. The supporting transfer can be delivered through a lower effective rate, favorable equity terms, debt, guarantees, or cash; in the transfer-separable benchmark, the instrument does not affect the implemented allocation.

The same logic gives the heterogeneity tests. The directed-activity gap, $k_{is} - k_{is}^{out}$, should be larger when two forces coincide. First, the recipient has weaker outside financing options, so the financing rent can support a larger departure from the commercial benchmark. Second, the financier places greater value on activity in that sector or market, so it has more reason to request directed

separability and unrestricted transfers, the contract term affects surplus division but not the implemented allocation. The cancellation can fail when the rate changes payoff risk, default incentives, limited-liability constraints, or the marginal cost of directed activity.

activity. I therefore test whether SWF-relative entry is stronger among lower-relative-valuation firms and in sectors that the SWF has historically overweighted, or where the SWF home country has greater strategic demand.

Proposition 2 (Network multiplier and crowd-in). *Strategic finance affects outside firms through the endogenous network stock Z . Under the stability condition in Equation (11), the shadow value of network expansion is*

$$\Psi_Z = \left[1 - \sum_{j \notin C} \frac{\mu'(Z)}{G_j''(k_{js}^*)} \right]^{-1} \left[\theta U'(Z) - \mu'(Z) \sum_{i \in C} (k_{is} - k_{is}^{out}) \right]. \quad (15)$$

The multiplier term is positive and finite, and $\Psi_Z > 0$ whenever

$$\theta U'(Z) - \mu'(Z) \sum_{i \in C} (k_{is} - k_{is}^{out}) > 0. \quad (16)$$

Intuition. The first bracketed term is the crowd-in multiplier. It is finite by the stability condition. It equals one when the network does not lower outside frictions ($\mu' = 0$) and exceeds one when a larger network stock lowers outside firms' bilateral frictions ($\mu' < 0$). The second bracket contains the financier's direct value of network expansion and the friction-reducing effect of a larger network. Crowd-in is stronger when $|\mu'(Z)|$ is large and outside firms respond more elastically (small $|G_j''|$).

Extensions. Appendix A studies two extensions. In a repeated relationship, the firm complies because deviating would end future access to favorable financing. In the binding-dynamic incentive compatibility, slack-participation regime, the dynamic extension gives $\tau_{is}^{dyn} = \beta \tau_{is}^{static} = \beta \Psi_Z$, where $\beta \in (0, 1)$ is the firm's discount factor. Relational enforcement attenuates the static covenant-enforceable wedge. The appendix also studies a first-best planner, clarifying which externalities the strategic financier does and does not internalize.

2.4 From model to tests

Figure 1 summarizes the mechanism. The model object is the induced gap between strategic-market activity under the contract and the firm's commercial counterfactual, $k_{is} - k_{is}^{out}$. This gap is positive only when the financier's strategic value and the feasible financing rent support a positive implementation wedge.

In PitchBook, the commercial counterfactual is not directly observed. I therefore use two empirical benchmarks. The main test compares post-deal expansion by the same firm into the SWF




Figure 1: Geofinance: Directed Activity and Aggregate Spillover

Notes: In Panels A and B, the firm's marginal return on strategic-market activity is $G'_i(k_{is})$. The commercial benchmark is $MC = \mu(Z) + r_i$, and the strategic-finance relationship lowers the effective marginal cost to $MC_{SF} = \mu(Z) + r_i - \tau_{is}^{imp}$. The wedge combines the financier's home-market relationship value with the feasible financing rent needed to make the bundle acceptable. In the interior benchmark, $\tau_{is}^* = \Psi_Z$. Panel A illustrates a case in which feasible terms cannot support the desired directed activity, leaving the recipient at the commercial benchmark or scaling back the action. Panel B illustrates a positive wedge, $\tau_{is}^{imp} > 0$, which generates $k_{is}^* > k_{is}^{out}$. Panel C shows how the induced activity raises Z , lowers $\mu(Z)$, and expands third-party capital from K_0 to K_1 .

home country versus into non-SWF co-investor home countries, using the latter as the ordinary investor-linked benchmark. This comparison matters because commercial co-investors in the same round also provide capital. If generic capital were the mechanism, expansion should tilt similarly toward commercial co-investor home countries. A SWF-specific differential instead points to the SWF home-country relationship, not generic capital alone. A separate matched-counterfactual exercise holds the potential SWF destination fixed and compares SWF-backed deal–country pairs with observationally similar non-SWF-backed pairs.

I use the firm’s relative valuation compared with sector-year peers as a reduced-form proxy for weaker outside financing options. Lower relative valuation, or higher valuation discount, does not measure an investor-specific contract rate. It measures whether the firm is lower-valued than comparable private firms in the same sector and year. Firms with weaker outside financing options should be more receptive to strategic finance, while the SWF home-country relationship determines where that receptiveness is expressed. In the bank setting, the analogous proxy is pre-crisis write-down severity.

I use SWF portfolio tilts and home-country export disadvantage and import reliance as proxies for the financier’s strategic preferences. If SWFs use foreign investments to build links in sectors where the home economy lacks revealed comparative advantage, the directed-action gap should be largest among firms with weaker outside financing options in those sectors.

Proposition 2 implies that directed activity lowers bilateral frictions for other investors and crowds in third-party capital. I study this implication using bilateral FDI inflows and BIS cross-border banking claims.

3 Directed Expansion in Private-Market Co-Investments

When SWFs and other institutional investors invest in the same private firm, I find that SWF participation alone does not predict directed expansion toward the SWF home country relative to the co-investor home countries. Instead, the SWF–non-SWF expansion differential appears among firms raising equity at a larger valuation discount relative to sector–year peers, consistent with directed activity becoming easier to support when the recipient has weaker outside financing options. The effect is strongest in sectors that SWFs had historically overweighted, especially when the SWF home country has low export comparative advantage or high import reliance in the sector.

3.1 Data and empirical design

Data. Using PitchBook, I construct a 1998–2025 sample of co-investment rounds in which at least one SWF and one non-SWF institutional investor participate in the same deal. Appendix B.1 documents the data construction and cleaning steps.

Sample. The unit of observation is the investor–deal pair. I impose three restrictions. First, I drop same-country pairs, where the target’s headquarters country matches the investor’s home country, to avoid classifying home-market activity as directed expansion. Second, I drop non-SWF investors located in the same country as the deal’s SWF, because if the SWF induces entry into its home market, those co-investors mechanically receive $Y = 1$ without independent evidence of redirection, which contaminates the control group.¹⁵ Third, for companies with multiple rounds, I drop later-round rows where the same investor already appeared in an earlier round to avoid double-counting persistent investor–target relationships across rounds. The restrictions produce a broader cleaned co-investment panel of 9,646 investor–deal pairs across 2,362 deals and 2,308 companies, with SWF investors from 24 home countries. The main tests require a non-missing valuation discount, the relative-valuation proxy introduced below. This valuation-discount sample contains 3,456 investor–deal pairs across 833 deals and 802 companies, with SWF investors from 20 home countries.¹⁶

Outcome variables. I identify expansion events using two PitchBook relationship routes after linking focal companies to PitchBook investor IDs with an exact normalized-name bridge (Appendix B.1).¹⁷ An expansion event is a subsequent PitchBook-recorded transaction by the target in the investor’s home country, including M&A, buyouts, corporate asset purchases, joint ventures, PE growth/expansion, corporate transactions, seed rounds, early-stage VC, later-stage VC, and PIPE. The main outcome is new market entry: $Y_{\text{ext}} = 1$ if the target expanded into the investor’s home country within five years after the deal but had no prior expansion there in PitchBook’s recorded

¹⁵Appendix Table 22 shows that this restriction is conservative. Re-adding the dropped rows leaves the $\text{SWF} \times \text{Discount}$ coefficient virtually unchanged, and dropping the full deal whenever the SWF’s home country has more than one investor in the round strengthens the estimate.

¹⁶In the valuation-discount sample, the typical deal has one SWF and roughly three institutional co-investors: the mean is 4.15 investor rows per deal and the median is 3; 36.6% of deals pair a single SWF with a single non-SWF investor; 89.6% involve only one SWF.

¹⁷The bridge requires the lower-cased, trimmed company name to equal the lower-cased, trimmed investor name. I do not use fuzzy matching, alias expansion, or manual predecessor-name stitching. As a result, the measure may miss expansion activity conducted under alternate legal names, subsidiaries, or affiliates that are not linked to the focal firm in PitchBook. This creates potential attenuation from missed true expansions rather than false positives from incorrect matches.

history.¹⁸

Key regressor: valuation discount. The model predicts that, conditional on SWF financing, the SWF-country expansion gap should be larger when the target has weaker outside financing options. It does not require high-discount firms to expand more in absolute terms. For each deal d in sector s and year t , I construct a valuation discount as the negative of the deal’s log post-money valuation residual relative to the sector×year mean¹⁹:

$$\text{ValDiscount}_d = - \left(\log(\text{PostVal}_d) - \overline{\log(\text{PostVal})}_{s(d),t(d)} \right), \quad (17)$$

where the benchmark is computed across all completed deals in sector s and year t from the full PitchBook deal universe.²⁰ Higher values indicate lower valuations relative to sector×year peers. I interpret this discount as a reduced-form proxy for weaker outside financing options, not as a direct measure of financing access. Lower relative valuation can reflect costly outside finance, higher perceived risk, less investor competition, or other outside-option premia. The robustness tests use alternative scale proxies and ownership-share proxies, including revenue, employees, residual valuation discount, and dilution measures.²¹

Composition in SWF-backed rounds. As supporting evidence on the outside-financing margin, Appendix Table 15 asks whether higher-discount rounds contain more SWF financing presence relative to non-SWF co-investors. Within the retained co-investment sample, a higher valuation discount is associated with a larger SWF share of institutional investor slots, and dilution gives a similar pattern.²² I interpret this as a composition check, not as a direct test of x_i^* : PitchBook rarely reports investor-level ticket amounts, and the sample already conditions on at least one SWF

¹⁸For later financing rounds whose five-year post-deal window extends beyond the latest observed PitchBook expansion data, outcomes are measured through the end of the observed expansion window. I do not condition on firm survival or continued PitchBook coverage; the issue is only calendar observability. Because the regressions include deal-by-year fixed effects, implemented as deal fixed effects, all investor-country rows within a financing round share the same available post-deal window.

¹⁹I use post-money rather than pre-money valuation because PitchBook reports post-money for roughly twice as many completed deals (21.2% versus 11.0% of the universe). The two values are correlated at over 95%.

²⁰The measure is winsorized at the 1st and 99th percentiles and standardized to unit standard deviation.

²¹As an additional check, I also considered size, age, and round-history proxies commonly used for financing conditions. Hadlock and Pierce (2010) argue that firm size and age are among the most useful observable predictors of outside-financing conditions. In this sample, smaller-revenue and smaller-employee targets show the same standalone pattern as valuation discount, while firm age and down-round status are weaker and less precisely measured. When these proxies are horse-raced against valuation discount, they do not add independent explanatory power. Appendix Table 23 reports the related scale, residual-discount, and control-share checks.

²²I use dilution, measured as deal size over post-money valuation, as a robustness proxy for financing intensity and control rights: higher dilution means the firm sells a larger ownership stake in the round.

investor. The main test below therefore asks a separate question: conditional on the round’s investor composition, does the SWF–non-SWF directed-expansion gap load on valuation discount?

Specification. I estimate

$$Y_{idt} = \beta_1 \text{SWF}_{idt} + \beta_2 (\text{SWF}_{idt} \times \text{ValDiscount}_{dt}) + \alpha_{dt} + \alpha_{it} + \alpha_{i \times h(d)} + \varepsilon_{idt}, \quad (18)$$

where i indexes the investor’s home country, d indexes the financing round, t is the deal year, and $h(d)$ is the target firm’s headquarters country. Y_{idt} is the post-deal expansion outcome into country i for deal d . β_1 captures the within-deal SWF–non-SWF expansion differential at the mean discount. β_2 captures whether that differential is larger in higher-discount deals. The deal-by-year fixed effects α_{dt} absorb firm-, round-, and deal-year shocks; in estimation they are implemented as deal fixed effects because each deal occurs in one year. The investor-home-country×deal-year fixed effects α_{it} absorb time-varying differences in destination-country attractiveness, and the investor-home-country×target-HQ-country fixed effects $\alpha_{i \times h(d)}$ absorb stable bilateral geography between investor home countries and target home countries.²³

3.2 Directed Expansion

Main results. Table 1 reports the main within-deal estimates. Column (1) shows no evidence that SWF-backed firms expand more into SWF markets on average. This does not mean that SWF backing has no effect relative to firms without an SWF-country investor anchor. The within-deal benchmark compares the SWF home country with other investor home countries in the same round, which also carry investor linkages and are often large destination markets such as the United States and the United Kingdom.

The effect appears only in higher-discount rounds. The SWF×valuation-discount interaction is positive in columns (2)–(4). Column (3) adds controls for the target’s pre-deal geographic orientation toward each investor country, and the coefficient changes little; the result is therefore not explained by firms already doing business in SWF countries before the financing round. Column (4) adds investor-home-country×deal-year and investor-home-country×target-HQ-country fixed effects, absorbing both time-varying destination-country attractiveness and stable bilateral geography between investor and target home countries.²⁴ The coefficient rises to 0.048. A one-

²³The within-deal investor-level design follows the within-borrower design of Khwaja and Mian (2008), applied here to investor–deal pairs.

²⁴A more saturated specification with investor-home-country×target-HQ-country×deal-year fixed effects gives a similar coefficient on SWF×valuation discount, 0.0488 (SE 0.0207, $p = 0.019$). I use the column (4) specification

standard-deviation increase in valuation discount therefore raises the within-deal SWF–non-SWF differential in new-entry probability by 4.8 percentage points, about 70% of the non-SWF control mean of 6.8%.²⁵

Table 1: Valuation Discount, SWF Co-Investment, and Business Expansion

	New Entry			
	(1)	(2)	(3)	(4)
SWF × Val. Discount		0.0281*** (0.0094)	0.0325*** (0.0099)	0.0481*** (0.0164)
SWF	0.0036 (0.0290)	-0.0231** (0.0092)	-0.0261*** (0.0095)	-0.0144 (0.0298)
Pre-acq. share			-0.0109* (0.0056)	
Pre-acq. share × Val. Disc.			0.0145** (0.0058)	
Deal-by-year FE (deal FE)	✓	✓	✓	✓
Investor-home-country × deal-year FE	✓	–	–	✓
Investor-home-country × target-HQ-country FE	✓	–	–	✓
R ²	0.761	0.548	0.552	0.766
SE clustered by company	✓	✓	✓	✓
Control mean	0.0679	0.0679	0.0679	0.0679
Observations	3,456	3,456	3,456	3,456

Note: Investor–deal level regressions. Column (1) regresses new entry on the SWF dummy under deal-by-year fixed effects, investor-home-country×deal-year fixed effects, and investor-home-country×target-HQ-country fixed effects, estimating the average SWF–non-SWF differential without valuation-discount heterogeneity; deal-by-year fixed effects are implemented as deal fixed effects because each deal occurs in one year. Columns (2)–(4) add the SWF×valuation discount interaction under progressively richer controls: deal-by-year fixed effects only (2); deal-by-year fixed effects plus pre-deal geographic orientation controls (3); deal-by-year fixed effects plus investor-home-country×deal-year and investor-home-country×target-HQ-country fixed effects (4). The new-entry outcome equals one if the target had no prior expansion in the investor’s home country before the deal but did expand within five years after the deal, or by the end of the observed PitchBook expansion window for later deals. Control mean is the mean of new entry for non-SWF co-investors. SE clustered by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Total effect and the marginal firm. The interaction $\hat{\beta}_2$ shows how the within-deal SWF gap varies with valuation discount.²⁶ It is a differential slope: it does not say that high-discount firms expand more into SWF countries in absolute terms, but that the SWF–non-SWF expansion gap

as the baseline because the saturated cells leave only 51 investor-home-country×target-HQ-country×deal-year cells with both SWF and non-SWF observations, covering 322 rows. It is therefore useful as a robustness check, while the baseline already absorbs destination-year attractiveness and stable bilateral geography with broader support.

²⁵Out of 100 non-SWF investor–deal pairs in the sample, the target expands into that non-SWF investor’s home country within five years in roughly 6.8 cases. The low baseline reflects the sparsity of bilateral cross-border expansion. Non-SWF co-investors are concentrated in deep markets, while SWF investors are concentrated in smaller economies. For example, Singapore accounts for 20 new-entry events in the sample, compared with 121 for the United States.

²⁶With deal-by-year fixed effects alone, implemented as deal fixed effects, this is the weighted covariance between the within-deal SWF–non-SWF expansion gap and the deal’s relative-valuation proxy.

risers as valuation discount increases. In column (4), the SWF indicator is -0.0144 and statistically indistinguishable from zero. The total within-deal SWF effect, $-0.0144 + 0.0481 \times \text{ValDiscount}$, turns positive at 0.30 standard deviations above the mean.²⁷ That 0.30-SD target is the threshold firm. Because the co-investment sample is concentrated among large private firms, the threshold still lies above the full PitchBook sector-year benchmark: in raw terms, its valuation discount is -2.80 log points, equivalent to a post-money valuation about 16.5 times the sector-year geometric mean. It is nevertheless lower-valued than the average co-investment target. This is the empirical counterpart to Proposition 1: the SWF-country gap turns positive only when the financing-rent benefit is large enough to offset the cost of entering the SWF market.²⁸

A destination-fixed matched comparison. The within-deal estimates use non-SWF co-investor home countries as the benchmark for ordinary investor expansion. That comparison is useful because it holds the financing round fixed, but it does not compare entry into the same SWF country. Appendix Table 24 provides a destination-fixed comparison using exact-cell regressions that compare SWF-backed deal-country pairs with similar non-SWF-backed pairs for the same potential SWF destination.

In the baseline control pool, SWF-backed firms are more likely to expand into SWF home countries than comparable non-SWF-backed firms in the same SWF-country, sector, deal-year-bin, and headquarters-region cells. After adding firm and deal controls, the SWF-backed coefficient is 6.32 percentage points. A tighter control pool then requires the control deal to have at least one non-SWF investor from the same country as the SWF investor in the treated pair. This raises the control entry rate to 7.65%, indicating that part of the gross level gap reflects a general investor-country anchor. The SWF coefficient nevertheless remains positive at 3.07 percentage points (SE 1.31).

This matched comparison helps reconcile the negative SWF level effect in Table 1. The main within-deal coefficient compares the SWF country with same-round non-SWF investor countries, which are often larger and more common expansion destinations such as the United States, the United Kingdom, and France. The matched comparison instead holds the SWF destination fixed and compares SWF-backed pairs with observably similar non-SWF-backed pairs. In that comparison,

²⁷Because the outcome is binary and new entry is rare, the linear probability model can generate fitted values outside $[0, 1]$, especially in fixed-effect cells with very low observed entry rates. I therefore interpret the estimates as linear-projection effects on the SWF–non-SWF entry differential, not as fitted probabilities for individual investor–deal cells. The SWF level term and the SWF $\times$ valuation-discount slope are consistently estimated under the standard zero-conditional-mean condition.

²⁸The crossover identifies the model’s marginal firm, the threshold at which the transfer just compensates the cost of entering the SWF’s market. The key empirical pattern is the positive slope, not a positive unconditional SWF level effect.

the level pattern is positive once the destination-country benchmark is held fixed. The model's sharper prediction remains the slope result in Table 1: the SWF-relative expansion gap rises with the recipient's relative-valuation proxy.

SWF sector tilts. The model predicts that directed expansion should be most valuable in sectors where the financier places high value on the bilateral link. I measure that value with portfolio tilts: an SWF's cross-border portfolio share in a sector minus the corresponding share for non-government deep-pocket private investors.²⁹ Appendix Figure 5 shows that these tilts vary sharply across SWF home countries. Appendix Figure 6 then benchmarks the high-tilt cells against early export revealed comparative advantage (RCA) in the mapped sector.³⁰ Most high-tilt cells have export RCA below one, so SWFs often overweight sectors in which the home country is not already an export specialist. This pattern points away from a mechanical comparative-advantage story and toward sectors where sovereign investors appear to value new foreign linkages.

Panel A of Table 2 asks whether these sector preferences explain away the valuation-discount result or instead sharpen it. They sharpen it. Across all four columns, the SWF $\times$ valuation-discount coefficient remains positive and statistically significant after adding individual-SWF sector-tilt controls. The level terms are negative at the mean valuation discount, which means favored sectors are not simply places where SWF-backed firms enter more often unconditionally. The interaction is the key object: a 10 percentage-point higher prior SWF-investor tilt raises the SWF $\times$ valuation-discount slope by 1.54 percentage points, and the implied slope at a +10 percentage-point tilt is 6.84 percentage points. The positive-tilt specification gives a similar implied slope of 5.35 percentage points. Sector preferences therefore do not replace the outside-financing channel; they identify where that channel is strongest.

Panel B connects the same mechanism to the SWF home country's trade position. I split SWF rows into sectors where the home country has export RCA below one, import RCA above one, and both conditions, always retaining non-SWF co-investors from the same deals as the within-deal benchmark. The SWF $\times$ valuation-discount coefficient is 6.63 percentage points in export-disadvantage sectors and 8.01 percentage points in import-intensive sectors. These are the main trade-position diagnostics: the SWF-relative valuation-discount slope is larger when the home

²⁹For Table 2, the tilt is computed at the individual-SWF level using only portfolio history before the focal deal year. The descriptive heatmap in Appendix Figure 5 reports pooled 2005–2025 home-country tilts. Both measures are count-based rather than value-weighted.

³⁰For country c and PitchBook sector s , RCA is $(X_{c,s}/X_c)/(X_{W,s}/X_W)$, where $X_{c,s}$ is the country's mapped goods-plus-services exports in sector s , X_c is its total goods-plus-services exports, and W denotes the corresponding world totals. RCA above one means the sector accounts for a larger share of the country's exports than of world exports. I use the geometric mean of annual RCA over 2005–2007. The appendix documents the non-overlapping goods-services mapping and source-level checks on world totals.

country appears to value new foreign links because it lacks export strength or relies on imports. The cell satisfying both conditions has a large point estimate, 10.27 percentage points, but too few observations to identify the effect separately with precision. I therefore read the joint cell as descriptive support rather than as the headline result.

Follow-on financing. The repeated-relationship extension in Appendix A gives another reason why directed expansion can be sustained after the initial financing round: a firm that values future SWF financing or relationship access has an incentive to maintain the relationship. Appendix Table 16 examines this continuation margin directly. SWF investors are more likely than same-deal non-SWF co-investors to return to the same portfolio company within five years, even after controlling for each investor’s prior follow-on propensity and same-home-country co-investors in the original round. Firms that receive SWF funding again have higher post-money valuations, raise more later rounds, and show the strongest descriptive link between valuation discount and subsequent expansion.

3.3 Robustness

The robustness checks address remaining selection and measurement concerns.

Valuation-discount robustness. Because valuation discount is constructed from post-money valuation, it may partly capture firm scale, maturity, or ownership stakes in the round rather than outside financing options. Appendix Table 23 addresses this concern directly. Smaller firms, measured by employees or revenue, show the expected standalone pattern, but those scale interactions lose explanatory power once valuation discount enters the same regression. The residual component of valuation discount, after partialling out log employees, log revenue, log firm age, and dilution, still predicts directed entry. Total dilution, equal-share dilution, and SWF-bloc dilution also point in the predicted direction when entered alone, but become small and statistically insignificant once valuation discount is included. These checks suggest that the relative valuation measure contains information about outside financing options that is not exhausted by simple scale or control-share proxies.

Round-level pricing diagnostics. A related concern is that valuation discount may capture generic round terms rather than outside financing options. Appendix C.5 addresses this directly. SWF-backed rounds have higher valuation multiples on average, but observable round-level pricing measures do not explain the directed-expansion gradient. In Table 18, $\log(P/\text{Revenue})$,

Table 2: SWF New Entry, Portfolio Tilts, and Sector Specialization

	New Entry			
	Tilt controls		Discount interactions	
	(1) Prior tilt	(2) Positive tilt	(3) Prior tilt	(4) Positive tilt
<i>Panel A. Individual-SWF portfolio tilt</i>				
SWF × Val. Discount	0.0505*** (0.0168)	0.0492*** (0.0166)	0.0529*** (0.0170)	0.0451*** (0.0163)
SWF × Prior SWF-investor tilt (/10 pp)	-0.0136 (0.0112)		-0.0159 (0.0114)	
SWF × Positive tilt cell		-0.0251 (0.0269)		-0.0258 (0.0278)
SWF × Val. Discount × Prior SWF-investor tilt			0.0154* (0.0086)	
SWF × Val. Discount × Positive tilt				0.0084 (0.0215)
Implied SWF × Val. Discount slope			0.0684*** (0.0226)	0.0535** (0.0228)
Deal-by-year FE (deal FE)	Yes	Yes	Yes	Yes
Investor-home-country × deal-year FE	Yes	Yes	Yes	Yes
Investor-home-country × target-HQ-country FE	Yes	Yes	Yes	Yes
SE	Company	Company	Company	Company
Mean dep. var.	0.056	0.056	0.056	0.056
Observations	3,456	3,456	3,456	3,456
<i>Panel B. Trade specialization split samples</i>				
	Export RCA < 1	Import RCA > 1	Export RCA < 1 & Import RCA > 1	
SWF × Val. Discount	0.0663*** (0.0251)	0.0801** (0.0328)	0.1027 (0.0748)	
Deal-by-year FE (deal FE)	Yes	Yes	Yes	
Investor-home-country × deal-year FE	Yes	Yes	Yes	
Investor-home-country × target-HQ-country FE	Yes	Yes	Yes	
SE	Company	Company	Company	
Mean dep. var.	0.055	0.065	0.047	
Observations	2,110	1,379	513	
Deals	522	294	120	

Notes: The outcome is new entry into the investor's home country within five years. Panel A uses observed investor-deal rows. Columns (1)–(2) add level controls for prior individual-SWF sector tilt and positive tilt; columns (3)–(4) add the corresponding valuation-discount interactions. Sector tilt is the individual SWF's prior cumulative PitchBook cross-border sector share minus the prior non-government deep-pocket benchmark share; one unit equals 10 percentage points. Panel B reports split samples for SWF rows in sectors with export RCA below one, import RCA above one, and both conditions, retaining non-SWF co-investors from the same deals as the within-deal benchmark. Export and import RCA are geometric means of mapped goods-plus-services RCA over 2005–2007; RCA above one means the sector's trade share exceeds the world benchmark. Fixed effects are shown in the table; deal-by-year fixed effects are implemented as deal fixed effects and absorb sector fixed effects. SE clustered by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

$\log(\text{DealSize}/\text{Revenue})$, and up-round indicators have little independent power once valuation discount is included, while the SWF $\times$ valuation-discount coefficient remains positive.

Leave-one-out. Appendix Figure 8 drops each SWF country in turn and re-estimates the main new-entry specification. The coefficient remains positive across all 20 SWF countries and stays close to the full-sample baseline. Appendix Table 17 shows that the new-entry coefficient remains positive under every sector omission, but it is most sensitive to omitting Consumer Products and Services (B2C).³¹ Appendix Tables 20 and 21 show that the result is stable to alternative winsorization and remains positive under 3-, 5-, and 7-year outcome windows. The window estimates rise modestly with the horizon, from 0.042 at three years to 0.050 at seven years.

Permutation inference. Appendix Table 19 asks whether the PitchBook result could arise from the particular investors labeled as SWFs within a financing round. In each permutation, I randomly reassign the SWF labels among same-deal investors. Panel A applies this exercise to a simple statistic: the SWF–non-SWF difference in new-entry rates after scaling each group by its baseline country deal share. Panel B repeats the exercise for the main regression coefficient on SWF $\times$ valuation discount. In both cases, the observed statistic lies outside the placebo distribution from 1,000 within-deal reassignments, so the result is unlikely to come from random same-round investor labels.

The PitchBook evidence supports the micro-allocation prediction: SWF co-investments are followed by a relative tilt toward SWF home countries, compared with non-SWF co-investor home countries, but only when firms have weaker outside financing options and especially in sectors where SWF home countries have relative disadvantages. The design absorbs deal-level shocks, fixes the destination benchmark in the matched comparison, and rejects random same-round SWF-label assignment as an explanation. It does not rule out selection on unobserved firm expansion plans that coincide with SWF funding. I next turn to crisis-era bank recapitalizations, where pre-crisis write-down exposure generates quasi-exogenous cross-bank variation in demand for sovereign capital.

³¹PitchBook's primary industry sectors in the valuation-discount co-investment sample are Business Products and Services (B2B), Consumer Products and Services (B2C), Energy, Financial Services, Healthcare, Information Technology, and Materials and Resources. B2C refers to consumer-facing products and services, including retail, restaurants, leisure and hospitality, media, travel, personal services, and consumer goods.

4 Directed Lending after Bank Recapitalizations During the Crisis

Similar bank lending redirection appears during the 2007–2008 financial crisis. The setting uses variation in banks’ pre-crisis need for outside capital, proxied by write-down exposure on mortgage-backed and related structured securities. In a bank–country–quarter syndicated loan panel built from the 2011 Financial Stability Board G-SIB universe, I find that high-write-down banks increased syndicated-loan deals toward SWF countries after 2007Q3, relative to banks with smaller write-downs.

4.1 Data and Empirical Design

Data. I assemble a bank–country–quarter panel of cross-border syndicated lending from Refinitiv DealScan for 2002–2015. Appendix B.2 details the panel construction, parent-bank assignment, and loan-term variables. The bank universe is externally defined by the 2011 Financial Stability Board list of global systemically important banks. The bank-level first-stage sample contains 30 banks, five of which received SWF equity during the 2007–2008 financial crisis. In the realized DealScan lending panel, Wachovia does not appear as a separate parent after its late-2008 acquisition by Wells Fargo, so the lending regressions use 29 bank parents: five treated and 24 controls. For each bank–country–quarter cell I measure whether lending begins and the number of syndicated-loan deals. The realized panel contains 163 borrower destination countries, of which 14 are effective SWF-status destinations.

SWF equity injections. Five banks received equity injections from eight SWFs across seven SWF countries, for a total of ten events. I define 2007Q3–2008Q1 as the clean recapitalization window. Banks first reported large structured-credit write-downs in 2007Q3, and SWF equity injections began in 2007Q4. I end the treatment window in 2008Q1 to avoid contamination from later government recapitalization programs and late-crisis private placements.³² This period also corresponds to the interval in which SWFs were a major source of external recapitalization, providing roughly 39% of all bank capital raised ([International Monetary Fund, 2008](#)).

³²The U.S. Troubled Asset Relief Program and the U.K. bank rescue scheme both began in 2008Q4. Once government equity became available, it became difficult to separate the lending response to SWF capital from the response to public recapitalization. Barclays is the main example: its late-2008 capital raise falls outside the clean window. I therefore keep Barclays in the control group rather than treat it as exposed, and the results are robust to dropping it.

Key regressors. I measure bank capital need with WriteDownRatio_b , defined as cumulative structured-credit write-downs in 2007Q3 and 2007Q4 divided by 2006 common equity.³³ Higher values indicate a larger crisis-era hit to bank capital. SWFStatus_c indicates whether country c was home to an active cross-border SWF as of 2006. This is an ex ante destination-level treatment: it classifies countries by the pre-crisis presence of sovereign investors, not by whether a particular bank ultimately received capital from that country. The baseline set contains 15 countries: the seven countries that injected equity into treated banks plus eight additional countries with SWFs that have documented pre-2007 cross-border private-market activity.³⁴ Post_t equals one in quarters at or after 2007Q3.³⁵ Table 3 summarizes the injection events and the bank write-down ratios.

Specification and identification. SWF equity receipt is endogenous at the bank level, so I estimate a reduced-form design that interacts bank write-down severity, destination-country SWF status, and the crisis cutoff. Write-downs reflect pre-crisis portfolio exposure and mark-to-market losses rather than country-specific lending choices, so WriteDownRatio_b is plausibly exogenous to the cross-country allocation of lending. SWFStatus_c is defined ex ante using countries with active pre-crisis SWFs. The design is therefore an intent-to-treat specification: it uses broad SWF-destination exposure rather than only the ten realized bank–SWF-country injection pairs.³⁶ The estimating equation on the bank–country–quarter panel is

$$Y_{bct} = \beta (\text{WriteDownRatio}_b \times \text{SWFStatus}_c \times \text{Post}_t) + \sum_{\ell} \gamma_{\ell} (\text{WriteDownRatio}_b \times X_{bc}^{\ell} \times \text{Post}_t) + \alpha_{bc} + \alpha_{bt} + \alpha_{ct} + \varepsilon_{bct}. \quad (19)$$

³³The Bloomberg write-down tracker compiles quarterly mark-to-market losses on subprime RMBS, CDOs, CMBS, leveraged loans, Alt-A securities, credit default swaps, and SIVs reported by financial institutions in their public filings.

³⁴The seven injection countries are Singapore, the UAE, Kuwait, Qatar, Saudi Arabia, China, and South Korea. The eight additional SWF countries (FY06 AUM $\geq$ \$10bn) are Bahrain, Australia, Oman, Malaysia, New Zealand, Brunei, Libya, and Algeria. Norway (public-market-only mandate), Kazakhstan (no pre-2007 private-market activity), and the United States (home country of three treated banks) are excluded from the SWF-status destination definition. Libya is retained in the ex ante SWF-country definition, but it has no realized syndicated-loan observations in the DealScan panel, so the estimating sample contains 14 effective SWF-status destination countries. Appendix Table 25 re-estimates the specification under the 20 PitchBook SWF home countries that contribute to the valuation-discount regression sample; the intensive-margin results remain positive and significant.

³⁵I use 2007Q3 as the post cutoff because large write-downs first became public in that quarter and recapitalization negotiations began before most SWF injections formally closed in 2007Q4 or 2008Q1. This dating captures the onset of the crisis shock without conditioning on the realized timing of bank–SWF matches. The main results are robust to using 2007Q4 or 2008Q1.

³⁶In the estimation window, the realized G-SIB panel contains 2,664 bank–country pairs, including 346 pairs involving SWF-status destinations. The full-controls specification uses 2,465 bank–country pairs and retains all 346 SWF-status pairs.

where Y_{bct} is new entry or log deal count on cells with positive lending. I do not use allocated loan volume as an outcome because DealScan reports tranche commitments at the syndicate level; mapping dollar amounts to parent banks requires allocation rules, whereas bank participation and deal counts are directly observed. The log-count specification conditions on positive lending, so it describes intensive-margin changes among active bank–country–quarter cells; the new-entry outcome captures the extensive margin.

β asks whether, after 2007Q3, high-write-down banks reallocate lending toward SWF destinations relative to low-write-down banks. $\mathbf{X}_{bc}$ is a vector of pair-level controls at their 2006 values, including the pre-crisis lending share, FTA, common language, and political distance, each interacted with WriteDownRatio_b and Post_t . Bank×country fixed effects absorb time-invariant bilateral relationships, bank×quarter fixed effects absorb bank-specific supply shocks, and country×quarter fixed effects absorb destination-specific demand shocks.³⁷ Appendix Figure 9 shows that WriteDownRatio is balanced on pre-crisis bank characteristics: none of the five standardized covariates is significant at the 5% level. The identification assumption is that, conditional on these fixed effects and controls, high-write-down banks would not otherwise have shifted lending disproportionately toward SWF countries after 2007Q3.

4.2 Directed Lending

Bank write-downs and SWF investments. Table 4 shows that banks with larger crisis write-downs were more likely to receive SWF equity. In the 30-bank G-SIB cross-section, a one-standard-deviation increase in WriteDownRatio ($\text{SD} \approx 0.149$) is associated with about a 31-percentage-point increase in the probability of receiving an SWF injection.³⁸ This result supports using write-down severity as a bank-level proxy for demand for sovereign capital.

Main results. Table 5 reports the continuous DiD estimates on the $[-8, +5]$ quarter window around 2007Q3. High-write-down banks increase lending activity toward SWF countries after the crisis cutoff, relative to low-write-down banks. The response is strongest on the extensive and deal-frequency margins. Columns (2) and (5) add $\text{WriteDownRatio} \times \text{Pre-share} \times \text{Post}$, where Pre-share is the bank–country lending share before the cutoff. This directly addresses the concern

³⁷Without country×quarter fixed effects, the coefficient on $\text{SWFStatus} \times \text{Post}$ is negative, reflecting the post-crisis contraction in cross-border bank lending (Cetorelli and Goldberg, 2011; De Haas and Van Horen, 2013). The full specification absorbs this aggregate pullback and isolates the relative shift toward SWF countries among high-write-down banks.

³⁸As in any linear probability model, fitted probabilities can exceed the unit interval at extreme write-down values in this small sample. The 0.149 standard deviation is computed across the 30-bank first-stage universe; the corresponding realized-panel bank-level standard deviation is 0.152.

Table 3: SWF Equity Injections and Bank Write-Down Ratios

Panel A: Treatment Events							
Bank	SWF	SWF ctry	Qtr	Inj. (\$bn)	Term price	D-1 stock	Prem./ Disc.
Citi	ADIA	UAE	2007Q4	7.50	\$31.83–\$37.24	\$31.70	Prem.
Merrill Lynch	Temasek	Singapore	2007Q4	4.40	\$48.00	\$55.51	Disc.
Morgan Stanley	CIC	China	2007Q4	5.60	\$48.07–\$57.68	\$48.07	Mixed
UBS	GIC	Singapore	2007Q4	9.70	CHF 47.68	CHF 50.87	Disc.
UBS	Saudi	Saudi Arabia	2007Q4	1.75	CHF 47.68	CHF 50.87	Disc.
Citi	GIC	Singapore	2008Q1	6.88	\$31.62	\$29.06	Prem.
Citi	KIA	Kuwait	2008Q1	3.00	\$31.62	\$29.06	Prem.
Credit Suisse	QIA	Qatar	2008Q1	3.00	n.a.	CHF 47.31	n.a.
Merrill Lynch	KIA	Kuwait	2008Q1	2.00	\$52.40–\$61.31	\$55.97	Mixed
Merrill Lynch	KIC	South Korea	2008Q1	2.00	\$52.40–\$61.31	\$55.97	Mixed
Panel B: Write-down Ratio Construction							
Bank	WD ratio	2006 Eq. (\$bn)	Loss (\$bn)				
Merrill Lynch	0.703	39.00	27.40				
UBS	0.435	44.62	19.40				
Morgan Stanley	0.292	35.30	10.30				
Credit Suisse	0.255	34.87	8.90				
Citi	0.199	119.80	23.80				

Note: Panel A: SWF recapitalization events with pricing benchmarks. Injection size is in USD bn; term price is the negotiated issue price on convertible or preferred securities; D-1 stock is the day-before common stock close; Prem./Disc. flags whether the term price was above or below the D-1 close. Panel B: write-down ratio used in the exposure term. Full summary statistics are in Appendix D.

Table 4: Bank Write-Downs and SWF Investments

	P(Injection) (1)	asinh(Total Amt.) (2)
WriteDownRatio	2.095** (0.956)	5.830** (2.570)
Banks (total)	30	30
Treated banks	5	5
R^2	0.682	0.661
SE	HC3	HC3

Note: Cross-sectional bank-level OLS on the 30-bank G-SIB universe from the Financial Stability Board's 2011 list. Unit of observation is a bank. Outcome in column (1): indicator for receiving any SWF equity injection (LPM). Outcome in column (2): asinh(total SWF injection amount, USD bn) across the clean 2007Q4–2008Q1 treatment window. WriteDownRatio is in raw units (0–1 scale). Standard errors are heteroskedasticity-robust (HC3), used because bank-level clustering is uninformative in this small cross-section. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

that the result simply reflects high-write-down banks expanding in destinations where they were already more active before the crisis. The new-entry and log-count estimates remain positive and significant after this adjustment.

Columns (3) and (6) add bilateral FTA, common-language, and political-distance controls, fixed at 2006 and each interacted with WriteDownRatio and Post. In the full-control specifications, the new-entry coefficient is 0.028 and the log-count coefficient is 1.144.³⁹ A one-standard-deviation increase in WriteDownRatio at 0.152 implies about 0.17 log points, or roughly 19%, higher deal counts among positive-lending cells. As in PitchBook, the most stable response appears in whether and how often recipients direct activity toward SWF countries.

Table 5: Bank Write-Down Ratios, SWF Countries and Redirected Syndicated Loans

	New Entry			log(Count)		
	(1)	(2)	(3)	(4)	(5)	(6)
WD × SWFStatus × Post	0.037*** (0.007)	0.036*** (0.007)	0.028*** (0.008)	1.114*** (0.308)	1.010*** (0.330)	1.144*** (0.381)
WD × Pre-share × Post	–	-0.027* (0.013)	-0.024 (0.048)	–	-1.635** (0.771)	-1.508* (0.824)
Add. controls (FTA, language, pol. dist.)	–	–	✓	–	–	✓
Bank×Country FE	✓	✓	✓	✓	✓	✓
Bank×Quarter FE	✓	✓	✓	✓	✓	✓
Country×Quarter FE	✓	✓	✓	✓	✓	✓
SE clustered by bank	✓	✓	✓	✓	✓	✓
Observations	37,296	37,296	34,510	10,761	10,761	10,010
R ²	0.260	0.260	0.268	0.809	0.809	0.813

Note: Bank–country–quarter panel within $[-8, +5]$ quarters of 2007Q3. Columns progress within each outcome group: (1)/(4) baseline; (2)/(5) adds $\text{WriteDownRatio}_b \times \text{Pre-share} \times \text{Post}$, where Pre-share is the bank–country average lending share over 12 quarters before 2007Q3; (3)/(6) adds the full control bundle: Pre-share, FTA, political distance, and common language fixed at 2006. Outcome groups: new entry LPM (cols 1–3, equals 1 only in the first post-cutoff quarter with positive deals for pairs that had zero deals in all eight pre-period quarters) and log(deal count) on cells with positive deal count (cols 4–6). Three-way FE: bank×country, bank×quarter, country×quarter. SE clustered by bank. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Dynamic evidence. Figure 2 reports the binned event study for log deal count using the baseline specification, with the quarter immediately before the crisis cutoff as the omitted bin.⁴⁰ The pre-period coefficients are statistically indistinguishable from zero and do not show a systematic upward trend before the cutoff. After the cutoff, the coefficients turn positive: 0.755 in the first post-crisis

³⁹A placebo-destination exercise that draws 1,000 non-SWF country sets gives a right-tail randomization-inference p -value of 0.074 and a two-sided p -value of 0.138 for the full-controls log-count coefficient. I therefore read the bank evidence as directionally consistent but more finite-sample sensitive than the PitchBook co-investment evidence.

⁴⁰I focus the main dynamic figure on log deal count because it is the most stable intensive-margin response. I aggregate individual quarters into two-quarter bins because quarterly bank–country lending is sparse, making quarter-by-quarter estimates noisy (Schmidheiny and Sieglöcher, 2023).

bin, 1.655 in quarters 2–3, and 1.203 in quarters 4–5. The pattern points to a sustained post-crisis increase in deal frequency toward SWF-status destinations rather than a pre-existing upward trend.




Figure 2: Binned Event Study: Deal Count

Notes: The figure shows $\text{WriteDownRatio}_b \times \text{SWFStatus}$ coefficients by binned event-time period relative to 2007Q3. Outcome: $\log(\text{deal count})$, restricted to cells with positive deal count. Reference bin: $t = -1$. Error bars: 95% CIs; SE clustered by bank. Baseline specification; three-way FEs (Bank $\times$ Country, Bank $\times$ Quarter, Country $\times$ Quarter).

Leave-one-out. Appendix Figure 10 drops each SWF destination country in turn and re-estimates the full-controls log-count specification. The coefficient remains positive in every omission and remains statistically significant except when Australia is omitted. Because Australia is part of the ex ante SWF-destination set but is not one of the realized sovereign capital providers in the crisis recapitalization window, this sensitivity is best read as a precision issue rather than evidence that a single capital-provider country drives the result.

The two micro settings point to the same financing-rent margin: recipient activity tilts toward SWF countries when recipients' outside financing options are weak. The next question is whether aggregate capital-flow outcomes move in a way consistent with lower bilateral frictions. Section 5 turns to that macro prediction by testing whether SWF-backed activity is followed by broader capital-flow responses: third-party FDI flows in FIRRMA and aggregate banking claims in BIS.

5 Bilateral Capital-Flow Responses

5.1 FDI Responses to FIRRMA

In 2018, the Foreign Investment Risk Review Modernization Act (FIRRMA) expanded CFIUS authority over foreign investments in sensitive U.S. businesses, raising regulatory risk for SWF investments in the United States. I find that SWF countries with greater pre-2018 high regulatory-risk exposure to FIRRMA subsequently attract more FDI from non-U.S. source countries, with the response concentrated among continuing bilateral relationships rather than new source-country entry.

Institutional setting. FIRRMA, enacted in August 2018, expanded the authority of the Committee on Foreign Investment in the United States (CFIUS) to review and restrict foreign acquisitions in sensitive sectors. For SWFs, the relevant margin was large-equity transactions in the United States. Appendix Figure 12 documents the corresponding non-U.S. deal margin: non-U.S. high-dilution SWF deal counts rise after 2018, including among countries with positive pre-FIRRMA U.S. high-dilution exposure. I use this pattern as institutional background rather than as a first-stage regression, because the post-2018 increase is not differential across exposure groups. The reduced-form test below therefore asks whether the countries most exposed to the U.S. shock subsequently attract more third-party FDI.

Exposure construction. The main exposure measure is constructed using the PitchBook data in Section 3 as the 2014–2017 share of a country’s SWF co-investment deals that both targeted the United States and involved above-median dilution:

$$\text{Exposed}_i^{\text{HD}} = \frac{\#\{\text{US-targeted, high-dilution deals by SWFs from } i\}_{2014-2017}}{\#\{\text{deals by SWFs from } i \text{ with non-missing dilution}\}_{2014-2017}}. \quad (20)$$

High-dilution deals are those with dilution above the median in the SWF co-investment sample. This margin maps to the policy shock because larger equity stakes are more likely to face tighter CFIUS scrutiny. I use the 2014–2017 window because SWF co-investment is thin and volatile before 2014. Under this measure, 9 of the 23 non-U.S. PitchBook SWF home countries have positive exposure;⁴¹ the remaining SWF destinations form the zero-exposure comparison group.

⁴¹The exposure figure is constructed for the full non-U.S. PitchBook SWF-home-country universe. After merging to the WBG-HBFDI panel and imposing the control coverage used in the main regressions, the estimating sample contains 22 effective SWF-home destinations, as reported in Table 6.

Appendix Figure 13 plots the cross-country distribution of the high-dilution FIRRMA exposure measure.

Specification and identification. I estimate whether SWF countries with greater pre-FIRRMA U.S. exposure subsequently attract more third-party FDI from non-U.S. source countries:

$$\text{asinh}(\text{FDI}_{jit}) = \beta (\text{Exposed}_i \times \text{Post}_t) + \mathbf{X}'_{jit}\gamma + \alpha_{jt} + \alpha_{ji} + \varepsilon_{jit}, \quad (21)$$

where j indexes source countries, i destination countries, and t years. The outcome is bilateral FDI inflows from the World Bank Harmonized Bilateral FDI Database.⁴² The coefficient β compares post-2018 changes in FDI inflows toward more- and less-exposed destinations. Source $\times$ year fixed effects absorb shocks to each source country's outward FDI, pair fixed effects absorb time-invariant bilateral affinity, and $\mathbf{X}_{jit}$ adds standard gravity controls like those in the bank setting.⁴³ Appendix Figure 11 shows that exposure is not strongly related to destination GDP, GDP growth, pre-2018 FDI growth, or the number of source partners. It is higher for SWF destinations with greater pre-2018 U.S. trade exposure and lower pre-2018 FDI inflow levels, so the main table reports specifications that interact the post-FIRRMA period with both U.S. trade share and standardized pre-2018 FDI level. The identifying assumption is that, absent FIRRMA, more-exposed SWF countries would not have experienced systematically different post-2018 changes in FDI from the same source countries.

Main results. Table 6 reports the main FDI estimates. The high-dilution exposure coefficient is positive and stable across specifications in the main sample. With the gravity controls (column 2), the estimate is 0.643. Column (3) adds the two balance controls motivated by Appendix Figure 11: the destination's pre-2018 U.S. trade share interacted with post-FIRRMA and its standardized pre-2018 FDI inflow level interacted with post-FIRRMA. The high-dilution coefficient remains positive at 0.663, while both balance-control interactions are small and statistically insignificant.

Columns (4)–(6) test whether the result depends on the source-country set or on the exposure definition. Column (4) restricts FDI source countries to the 47 non-U.S. countries that appear as targets of high-dilution SWF co-investment deals in PitchBook; the coefficient rises to 1.574. Columns (5)–(6) replace the baseline high-dilution exposure with two broader measures of pre-FIRRMA U.S. exposure: the U.S. share of all deal counts and the U.S. share of all deal values for

⁴²Because bilateral FDI flows can be negative or zero, I use the inverse hyperbolic sine transformation.

⁴³Controls include free trade agreement (FTA) membership, destination GDP level and growth, $\text{asinh}(\text{bilateral trade})$ and diplomatic distance measured by UN-voting ideal-point distance.

each SWF country.⁴⁴ The estimates remain positive and significant. I use the event study mainly for timing and the decomposition table to separate entry, continuation, and intensive-margin responses behind the asinh outcome.

Table 6: FIRRMA Exposure and Third-Party FDI Inflows to SWF Home Countries

	High-Dilution Exposure				Count	Value
	(1)	(2)	(3)	Target Sources (4)	(5)	(6)
SWF US Exposure $\times$ Post	0.560*** (0.131)	0.643*** (0.137)	0.663*** (0.215)	1.574*** (0.581)		
SWF US Exposure (count) $\times$ Post					0.395** (0.163)	
SWF US Exposure (value) $\times$ Post						0.424*** (0.145)
US Trade Share $\times$ Post			-0.109 (1.278)	-0.991 (3.153)	1.241 (1.075)	1.259 (1.018)
Pre-FIRRMA FDI Level $\times$ Post			0.008 (0.075)	0.080 (0.152)	-0.053 (0.063)	-0.031 (0.065)
asinh(Bilateral Trade)		-0.707** (0.297)	-0.707** (0.302)	-1.209** (0.606)	-0.711** (0.302)	-0.694** (0.302)
Diplomatic Distance		-0.248 (0.210)	-0.249 (0.211)	-0.448 (0.505)	-0.243 (0.211)	-0.231 (0.210)
Other gravity controls	–	✓	✓	✓	✓	✓
Source $\times$ Year FE	✓	✓	✓	✓	✓	✓
Pair FE	✓	✓	✓	✓	✓	✓
SE clustered by pair	✓	✓	✓	✓	✓	✓
Wild-cluster p (destination)	0.007	0.024	0.015	0.056	0.168	0.092
Source countries	224	224	224	47	224	224
Destination countries	22	22	22	22	22	22
Observations	17,591	17,591	17,591	5,529	17,591	17,591
R^2	0.362	0.363	0.363	0.343	0.363	0.363

Note: Dependent variable: asinh(bilateral FDI inflow) from source j to destination i , World Bank Harmonized Bilateral FDI Database, 2014–2020; see Appendix B.3 for data construction. The estimation sample uses the full PitchBook SWF-home-country universe, excluding the United States as an FDI destination; the controlled regressions contain 22 effective destinations. Because bilateral flows include negatives and zeros, the inverse hyperbolic sine is used. Columns (1)–(3): high-dilution exposure (share of above-median-dilution pre-2018 SWF co-investment deals targeting the U.S.) with progressive controls; column (3) adds U.S. trade share $\times$ Post and standardized pre-FIRRMA FDI level $\times$ Post, where pre-FIRRMA FDI level is the destination’s 2014–2017 mean asinh bilateral inflow, standardized across destinations. Columns (4)–(6) include those balance controls: column (4) restricts source countries to the 47 nations where SWFs made high-dilution non-U.S. deals, and columns (5)–(6) use count-based and value-weighted exposure with full controls. All columns include source $\times$ year and pair FE. Gravity controls: FTA, log(GDP), GDP growth, asinh(bilateral trade), diplomatic distance. Standard errors are clustered by country pair. The wild-cluster row reports restricted wild-cluster bootstrap p -values clustered by SWF destination country. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

⁴⁴The exposure is counted at the SWF-home-country-by-deal level: if two SWFs from the same home country join the same deal, that deal contributes once to the country’s count-based and high-dilution exposure measures. The value-weighted exposure measure sums same-country SWF value within a deal before forming the country share.

Dynamic evidence. Figure 3 plots annual event-study coefficients from the balance-control specification in column (3) of Table 6. I use the figure as a timing check rather than as an estimate of the exact adjustment path, because annual FDI flows are volatile and the final post year, 2020, coincides with pandemic-era disruptions to global investment. The pre-FIRRMA leads are statistically indistinguishable from zero.⁴⁵ After FIRRMA, the coefficients are positive, with the largest estimate in 2020.




Figure 3: FIRRMA Event Study: FDI Inflows to SWF Home Countries

Notes: Annual event-study coefficients from the balance-control specification corresponding to column (3) of Table 6. The specification includes high-dilution exposure interacted with event time, U.S. trade share $\times$ Post, standardized pre-FIRRMA FDI level $\times$ Post, gravity controls, source $\times$ year FE, and pair FE. Sample: 22 effective non-U.S. SWF home destinations, 2014–2020. The omitted year is 2017. Error bars are 95% confidence intervals with standard errors clustered by country pair.

Margin decomposition. Because bilateral FDI flows can be negative, zero, or positive, the asinh specification mixes entry, continuation, and flow-size changes. Table 7 separates these margins following Haug, Nguyen and Owen (2023). The evidence points to stronger continuation among existing bilateral relationships rather than new relationships from previously inactive source countries. The entry margin does not increase for source–destination pairs without positive pre-FIRRMA inflow (column 1). In contrast, existing relationships strengthen on two margins. Conditional on positive current flows, the intensive-margin coefficient is positive, 0.480, but imprecisely estimated. Among pairs that were already active before FIRRMA, a one-unit increase in high-dilution exposure lowers the annual probability of falling to zero or negative flow by 15.8 percentage points. This pattern is consistent with the model’s friction-reduction channel: exposed SWF destinations

⁴⁵The joint pre-period test does not reject under pair-clustered standard errors ($p = 0.120$) or destination-clustered standard errors ($p = 0.288$).

do not mainly add new source-country relationships, but existing bilateral investment relationships become larger and less likely to disappear.

The decomposition also gives a cleaner economic scale. At observed exposure shares, the conditional log-flow coefficient in column (2) implies increases of about 5.0 percent for Singapore (exposure 0.105), 6.9 percent for China (exposure 0.143), and 8.0 percent for Australia (exposure 0.167), among source–destination pairs with positive current FDI flows. The exit-margin coefficient in column (3) implies that high-dilution exposure lowers the annual probability that an already-active pair falls to zero or negative inflows by 1.7 percentage points for Singapore, 2.3 percentage points for China, and 2.6 percentage points for Australia.

Table 7: FIRRMA Spillover: Entry / Intensive / Exit Decomposition

	(1) Entry	(2) Intensive	(3) Exit
SWF US Exposure (high-dilution) $\times$ Post	-0.028 (0.022)	0.480* (0.252)	-0.158*** (0.050)
SWF US Exposure (count) $\times$ Post	-0.043*** (0.014)	0.287* (0.166)	-0.092** (0.036)
SWF US Exposure (value) $\times$ Post	-0.041*** (0.013)	0.194 (0.132)	-0.103*** (0.028)
Outcome	$\mathbf{1}\{\text{flow} > 0\}$	$\log(\text{flow})$	$\mathbf{1}\{\text{flow} \leq 0\}$
Sample	Never-active pre	in_flow > 0	Ever-active pre
Gravity controls	✓	✓	✓
Balance controls	✓	✓	✓
Source $\times$ Year FE	✓	✓	✓
Pair FE	✓	✓	✓
SE clustered by pair	✓	✓	✓
Observations	6,683	7,056	10,908

Note: Each cell is a separate regression estimated on the SWF-home destination sample from Table 6. Pre-period is 2014–2017; the FIRRMA shock occurs in 2018. A pair is classified as *never-active pre* if its bilateral inflow was non-positive in every pre-period year, and *ever-active pre* otherwise. Column (1) estimates the *entry* margin: a linear probability model for $\mathbf{1}\{\text{in_flow} > 0\}$ on the never-active-pre cohort, asking whether HD-exposed pairs are more likely to begin receiving positive bilateral inflows post-FIRRMA. Column (2) estimates the *intensive* margin: an OLS regression of $\log(\text{flow})$ on all pair-year observations with positive current flow, regardless of pre-period cohort, asking whether bilateral flow levels rise conditional on remaining positive. Column (3) estimates the *exit* margin: a linear probability model for $\mathbf{1}\{\text{in_flow} \leq 0\}$ on the ever-active-pre cohort, asking whether already-active pairs are less likely to drop to zero or negative flow. Thus, the entry and exit cohorts are disjoint, but the intensive-margin sample overlaps with both because it conditions on current positive flow. All specifications include the full set of gravity controls (FTA membership, log destination GDP, GDP growth, asinh bilateral trade, diplomatic distance), the two balance controls from Table 6, source $\times$ year and pair fixed effects, and standard errors clustered by country pair. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Robustness. Appendix Table 26 reports the same specifications on a broader non-U.S. destination universe. It compares FDI inflows into the 9 positive-exposure SWF destinations with inflows into zero-exposure destinations in the full non-U.S. destination panel. The magnitudes of the coefficients fall by about half but remain significant. I present the SWF-home destination sample as my main

specification on conceptual grounds: the geofinance mechanism predicts crowd-in toward SWF home countries specifically, not toward non-SWF destinations.

Appendix Table 27 reports timing and oil-exporter checks using the balance-control specification. Placebo shocks in 2015 and 2016, estimated entirely before FIRRMA, yield small and statistically insignificant coefficients. Dropping oil-exporter SWF countries keeps the estimate positive and statistically significant. Leave-one-positive-exposure-SWF-country-out estimates remain positive throughout (Appendix Figure 14). Dropping China leaves the coefficient essentially unchanged, so the result does not capture the contemporaneous U.S.–China trade war.

FIRRMA identifies crowd-in in bilateral FDI, but exposure varies only across SWF destination countries. I next use pair-level variation to ask whether a similar pattern appears in cross-border banking claims.

5.2 Banking Claims after SWF Equity Injections

Similar crowd-in appears in the banking sector. Using BIS Locational Banking Statistics in a matched stacked-cohort design centered on SWF equity injections into banks, I find that lender-country claims on SWF home countries rise after treatment, build over the first two post-treatment years, and remain elevated through year 5.

Data and treatment. The BIS Locational Banking Statistics report bilateral cross-border claim stocks of banks resident in a reporting country on counterparties in each borrower country. These stocks measure outstanding banking claims rather than new loan originations, so I interpret the estimates as changes in bilateral banking exposure. The data cover claims beyond syndicated lending, including direct loans, debt securities, and deposits, and span 1995–2024 for 31 reporting countries and 218 borrower countries. Appendix B.4 describes the data construction.

I identify SWF bank equity injections from the Global SWF (GSWF) and Sovereign Wealth Fund Institute (SWFI) transaction databases, retaining direct bank equity injections and excluding asset-management mandates, insurance investments, and fund-of-fund commitments. The raw treatment universe contains 107 deals across 55 lender–SWF-country pairs, totaling \$127 billion from 13 SWF countries. After merging to the LBS panel and applying data-availability restrictions, 12 treated pairs remain in the stacked-DiD design.⁴⁶ These treated pairs are concentrated among developed-country lenders and SWF-country borrowers.

⁴⁶Of the 55 treatment pairs, 48 have disclosed deal sizes; of those 48, 12 have lender countries that report to the BIS LBS over the treatment window and enter the matched stacked-DiD design.

Specification and design. The treatment variable is *SWF investment intensity*: the total SWF equity injected, in USD billions, into banks in lender country l by SWFs from borrower country b . For example, disclosed Singaporean SWF injections into U.S. banks total \$19.8 billion across three events, generating an intensity of 19.8 for the USA→Singapore pair. This is a pair-level aggregate treatment: if the same lender–borrower pair receives multiple disclosed injections, I sum them into a single total intensity and assign that pair to a cohort at its first injection year.⁴⁷ The estimand is therefore the aggregate change in lender-country banking claims associated with the overall SWF–bank relationship in the pair, not the effect of only the first observed transaction. Within each cohort, I match up to 20 control pairs from the same lender country on pre-treatment log-loan levels and growth rates using Mahalanobis distance.⁴⁸ I then stack the cohorts and estimate

$$\log(Y_{lbt}) = \beta \cdot (\text{Intensity}_{lb} \times \text{Post}_t) + \alpha_{g(lb)t} + \alpha_{g(lb),lb} + \varepsilon_{lbt}, \quad (22)$$

where Y_{lbt} is a positive bilateral claim stock, Intensity_{lb} is the pair-specific treatment intensity, and Post_t equals one for periods at or after the cohort’s treatment date. Cohort×period fixed effects absorb time-varying shocks common within a cohort. Cohort×pair fixed effects absorb time-invariant bilateral heterogeneity. β estimates whether more intensely treated lender–borrower pairs experience larger post-treatment increases in bilateral claims.

Investment intensity is endogenous: SWFs may target lender countries with stronger pre-existing bilateral ties or better growth prospects. The key assumption is therefore parallel trends within each stacked cohort. Conditional on cohort×period and cohort×pair fixed effects, treated pairs would have followed the same path as their matched controls absent SWF bank injections. Two features make this assumption more plausible. First, Mahalanobis matching aligns controls with treated pairs on pre-treatment log-loan levels and growth rates. Second, treatment is staggered across twelve lender–borrower pairs over 2007–2018, which reduces the scope for any single common shock to explain the full pattern. Although the matched stacks include many control pairs, inference is still based on 12 treated cohorts, so I report wild-cluster-bootstrap p -values in the main table as a small-sample diagnostic.

Main results and magnitude. Table 8 reports the annual main specification using Q4 stocks. Cross-border banking claims from the lender country to the SWF home country rise after SWF

⁴⁷This biases the implied per-dollar effect downward in early years for pairs with later additional events.

⁴⁸Each treated cohort is matched to the $k = 20$ nearest controls within the same lender country by Mahalanobis distance on pre-treatment log-loan level and slope. Post-match absolute standardized differences range from 0.05 to 0.19, below the 0.25 benchmark commonly used to flag large covariate imbalance in matching diagnostics (Stuart, 2010); see Appendix Table 28.

bank equity injections. In column (1), total claims rise by 0.058 log points per \$1 billion of SWF equity. Column (2) shows a similar estimate at 0.057 for loans and placed deposits, and column (3) shows a similar estimate at 0.053 for non-bank loans. The similarity across columns suggests that the increase is concentrated in loan-and-deposit claims to SWF-country borrowers rather than in debt securities or interbank positions. Column (4) shows a positive and significant effect on winsorized standardized loan flows, although that margin is no longer precisely estimated under wild-cluster-bootstrap inference.

The economic magnitude is large. The loans estimate implies roughly 5.8% higher bilateral cross-border loans per additional \$1 billion of SWF equity. This corresponds to about \$345 million at the pre-treatment treated-pair median and about \$600 million at the treated-pair mean. This is a stock interpretation: it describes the average post-treatment outstanding claims, not new loan originations or cumulative annual flows.⁴⁹

Dynamic evidence. Figure 4 reports year-by-year event-study coefficients for total claims, loans and deposits, and non-bank loans. Across the three stock outcomes, the pre-treatment coefficients are centered near zero, which supports the parallel-trends assumption and argues against a simple selection story in which SWF-country claims were already rising before the equity injections. After treatment, the coefficients turn positive, build over the first two post-treatment years, and remain elevated through year 5. The similar timing across total claims and loan-based components suggests that the post-treatment increase reflects a broader rise in lender-country banking exposure to SWF home countries, not only a narrow claim category.

Cross-channel evidence. Table 9 asks whether the banking response is specific to SWF equity injections into banks. I re-estimate the stacked-cohort design around unlisted non-financial SWF investments, excluding pairs with SWF bank equity injections. The estimates are positive but much smaller than the direct bank-channel effects; only non-bank loans are statistically significant. This contrast is informative because a broad country-level information channel would predict similar banking responses after non-financial SWF investments. Instead, banking claims respond much more strongly when SWF capital enters banks directly, consistent with a relationship-friction channel within banking links.

⁴⁹The scale calculation uses, for each treated lender–borrower pair, the pair’s mean pre-treatment loans and placed deposits in the annual stacked sample. Across the 12 treated pairs, the median is \$5.9 billion and the mean is \$10.3 billion.

Table 8: Cross-Border Banking Claims Spillover

	(1)	(2)	(3)	(4)
	Total Claims	Loans & Placed Dep.	Non-Bank Loans	Wins. Std. Flow Loans
SWF Invest (\$bn) $\times$ Post	0.0581** (0.0195)	0.0566** (0.0193)	0.0530*** (0.0151)	0.0252** (0.0113)
Outcome transform	log	log	log	wins. flow/pre-stock
Observations	3,259	3,258	3,250	3,259
Cohort $\times$ Year FE	✓	✓	✓	✓
Cohort $\times$ Pair FE	✓	✓	✓	✓
Cluster	Cohort	Cohort	Cohort	Cohort

Note: Matched stacked-cohort DID using BIS Locational Banking Statistics (residence-based). Treatment intensity: total SWF equity injection in \$bn into banks domiciled in the reporting country. When a lender-country $\times$ SWF-country pair receives multiple disclosed injections, those amounts are summed into a single pair-level intensity and the cohort is anchored at the pair's first injection date. Columns (1)–(3): log annual cross-border claim stocks (including loans, deposits and debt securities), using positive Q4 end-of-year stocks; column (2) restricts to loans & deposits; column (3) further restricts to non-bank counterparties. Column (4): reported annual LBS loan flow, summed from quarterly flows, divided by pre-treatment mean loan stock and winsorized at the 1st and 99th percentiles. Controls are Mahalanobis-matched on pre-treatment log-loan level and slope. Annual main specification uses a ± 6 -year window around treatment; quarterly dynamics are reported in the appendix. SE clustered by cohort, defined by the treated lender-country $\times$ SWF-country pair and its first injection year. Wild-cluster-bootstrap- t p -values clustered by cohort are 0.007, 0.009, 0.014, and 0.216 for columns (1)–(4), respectively. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.




Figure 4: Cross-Border Claims: Annual Intensity-Weighted Event Study

Notes: Year-by-year stacked-cohort event study for BIS LBS cross-border claims. Panels report log total claims, log loans and deposits, and log non-bank loans. Treatment intensity: SWF equity injection (\$bn). Reference year: $t - 1$. Controls are Mahalanobis-matched using pre-treatment log-loan level and slope. Main annual design uses Q4 end-of-year stocks in a ± 6 -year window around treatment with cohort $\times$ year and cohort $\times$ pair FEs; SE clustered by cohort. Quarterly event-study dynamics are reported in the appendix figure.

Table 9: Cross-Channel Banking Claims Response from SWF Firm Investments

	(1) Total Claims	(2) Loans & Deposits	(3) Non-Bank Loans
Treatment intensity $\times$ Post	0.0128 (0.0091)	0.0133 (0.0093)	0.0135** (0.0064)
Observations	23,138	22,983	21,526
Cohort $\times$ Year FE	✓	✓	✓
Cohort $\times$ Pair FE	✓	✓	✓
Cluster	Cohort	Cohort	Cohort

Note: Stacked-cohort DiD built around firm-channel treatment events, defined as the first SWF investment in an unlisted non-financial firm in each lender–borrower pair over 2002–2022 (GSWF/SWFI source; finance and listed/public-market deals excluded). The sample is restricted to pairs with no SWF bank equity injection in the same country pair, so the firm-channel effect is estimated independently of the bank equity injections mechanism. Treatment intensity is total SWF firm investment in the pair (USD bn); it enters interacted with Post. Outcomes are log annual BIS LBS claims measured using positive Q4 end-of-year stocks: total claims, loans and deposits, and loans to non-bank counterparties. Treated pairs are matched within lender country to up to 20 controls on pre-treatment log loans and log-loan growth rates. Cohort $\times$ year and cohort $\times$ pair fixed effects; SE clustered by cohort. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Robustness. These checks address temporal aggregation, placebo outcomes, borrower concentration, and residence-based measurement. Appendix Table 29 re-estimates the design at quarterly frequency with cohort $\times$ quarter fixed effects. The coefficient remains close to the annual baseline, so the result is not an artifact of using Q4 stocks. Appendix Table 30 column (5) reports a liability-side test based on BIS liability positions; the estimate is positive but statistically indistinguishable from zero. The asymmetry between claim-side and liability-side responses indicates that the bilateral expansion is concentrated in cross-border lending toward the SWF country rather than having more deposits from the SWF country.

Figure 15 drops each treated lender–borrower cohort in turn and re-estimates the intensity-weighted specification for loans. The coefficient remains positive throughout, so no single cohort drives the result.

Appendix Table 30, columns (6)–(7), compares the LBS estimates with BIS Consolidated Banking Statistics (CBS), which assign claims by parent-bank nationality rather than office residence. The CBS estimates on total claims (0.043) and on international claims (0.060) are close to the LBS total-claims baseline (0.058), which suggests that the result is not solely explained by the residence-based reporting convention.

Taken together, the regulatory-shock and bank-injection exercises provide complementary evidence on the model’s crowd-in implication. Both show increases of roughly 5–8 percent toward SWF home countries. The responses appear mainly on continuation or intensive margins, where lower bilateral frictions should matter most.

6 Conclusion

Sovereign capital does more than finance firms. By exploiting financing rents, sovereign investors can direct recipients' cross-border activity toward their home countries. The resulting activity builds bilateral links that can lower frictions for later investors. The paper connects this micro contracting margin to the geography of cross-border capital flows.

Three limitations define the scope of the evidence. First, the reduced-form estimates identify directed responses to observed SWF investments, but they do not quantify counterfactual changes in sovereign capital policy. Second, SWF transactions are large and infrequent, which limits treatment variation in each empirical setting. Third, the evidence focuses on SWFs' private investments, while their public-market portfolios remain outside the analysis. The external validity of the mechanism to passive public equity holdings is therefore an open question.

The framework developed here, geofinance, extends beyond sovereign wealth funds. It applies whenever a financier is willing to trade financial return for directed activity. Development finance institutions, state-owned enterprises, and strategic corporate investors are natural settings in which the same logic may apply.

APPENDIX OVERVIEW

Appendix A Model Derivations and Extensions

A.1 Benchmark Proofs A.2 KKT Conditions A.3 First-Best A.4 Dynamic IC

Appendix B Data Construction (panels and summary statistics)

B.1 PitchBook B.2 DealScan B.3 Bilateral FDI B.4 BIS LBS/CBS

Appendix C Additional PitchBook Results

C.1 SWF Representation C.2 Country-Sector Tilts C.3 Follow-on Financing
C.4 Sector Heterogeneity C.5 Round-Level Pricing Diagnostics
C.6 Leave-One-Out C.7 Permutation Inference C.8 Winsorization
C.9 Alternative Windows C.10 Same-Country C.11 Alternative Mechanisms
C.12 Destination-Fixed Matched Comparison

Appendix D Additional Bank Syndicated Lending Results

D.1 Pre-Crisis Balance: WriteDownRatio D.2 Leave-One-Out
D.3 Alternative SWF Treatment Definitions

Appendix E FIRRMA Spillover Robustness

E.1 Full-Destination Main Spec E.2 Pre-Treatment Balance
E.3 Non-U.S. Deal Volume E.4 Exposure by SWF Country
E.5 Placebos E.6 Leave-One-Out

Appendix F BIS Spillover Robustness

F.1 Matching Balance F.2 Leave-One-Out
F.3 Quarterly Spec F.4 Annual Spec & Nationality Robustness

A Model Derivations and Extensions

A.1 Benchmark regime and proofs of the main results

The propositions in Section 2 use an interior benchmark for contracted recipients: funding, directed deployment, outside-firm responses, and the contract rate are interior. The aggregate-capacity constraint may bind, with multiplier λ . The home-market value $A_i(r_i)$ is additively separable and independent of k_{is} , r_{is} , and Z , so it drops out of the derivative calculations below. This subsection derives the formulas; Appendix A.2 gives the KKT system and corner cases.

The Lagrangian is

$$\begin{aligned} \mathcal{L} = & \sum_{i \in C} \left[(r_{is} - \rho)x_i - \frac{c}{2}x_i^2 \right] + \theta U(Z) \\ & + \sum_{i \in C} y_i \left[\Pi_i(x_i, r_{is}, k_{is}, Z) - v_i(Z) \right] \\ & + \lambda \left[\bar{X} - \sum_{i \in C} x_i \right] \\ & + \Psi_Z \left[\sum_{i \in C} k_{is} + \sum_{j \notin C} k_{js}^*(Z) - Z \right]. \end{aligned} \quad (23)$$

The inequality multipliers y_i and λ are nonnegative. The equality multiplier Ψ_Z is unrestricted in sign.

Contract rate r_{is} . From Equation (4), $\partial \Pi_i / \partial r_{is} = -x_i$. Stationarity gives

$$\frac{\partial \mathcal{L}}{\partial r_{is}} = x_i - y_i x_i = 0. \quad (24)$$

For any contracted recipient with $x_i > 0$ and an interior contract rate,

$$y_i = 1. \quad (25)$$

The participation constraint therefore binds, and the contract rate is pinned down by the reservation value.

Financing volume x_i . Using $\partial \Pi_i / \partial x_i = r_i - r_{is}$,

$$\frac{\partial \mathcal{L}}{\partial x_i} = (r_{is} - \rho) - cx_i + y_i(r_i - r_{is}) - \lambda = 0. \quad (26)$$

Substituting $y_i = 1$ yields

$$x_i^* = \frac{r_i - \rho - \lambda}{c}, \quad (27)$$

so the contract rate r_{is} drops out of the marginal funding condition.

External action k_{is} (Proposition 1). From the operating payoff in Equation (4),

$$\frac{\partial \Pi_i}{\partial k_{is}} = G'_i(k_{is}) - \mu(Z) - r_i. \quad (28)$$

The financier's first-order condition is

$$\frac{\partial \mathcal{L}}{\partial k_{is}} = y_i [G'_i(k_{is}) - \mu(Z) - r_i] + \Psi_Z = 0. \quad (29)$$

Using $y_i = 1$ and rearranging,

$$G'_i(k_{is}) = \mu(Z) + r_i - \Psi_Z. \quad (30)$$

Comparing this expression with Equation (8) gives

$$\tau_{is}^* = \Psi_Z. \quad (31)$$

Aggregate network stock Z (Proposition 2). Differentiating $\mathcal{L}$ with respect to Z gives

$$\frac{\partial \mathcal{L}}{\partial Z} = \theta U'(Z) + \sum_{i \in C} y_i \left(\frac{\partial \Pi_i}{\partial Z} - \frac{\partial v_i}{\partial Z} \right) + \Psi_Z \left[\sum_{j \notin C} \frac{\partial k_{js}^*(Z)}{\partial Z} - 1 \right] = 0. \quad (32)$$

Maintaining that Z enters firm value only through the cost term $-\mu(Z)k$, and applying the envelope theorem to the outside option,

$$\frac{\partial \Pi_i}{\partial Z} - \frac{\partial v_i}{\partial Z} = -\mu'(Z)(k_{is} - k_{is}^{\text{out}}). \quad (33)$$

Total differentiation of the outside-firm condition

$$G'_j(k_{js}^*) = \mu(Z) + r_j$$

yields

$$G''_j(k_{js}^*) \frac{\partial k_{js}^*}{\partial Z} = \mu'(Z) \quad \implies \quad \frac{\partial k_{js}^*}{\partial Z} = \frac{\mu'(Z)}{G''_j(k_{js}^*)}. \quad (34)$$

Substituting these expressions into the Z first-order condition and using $y_i = 1$ gives

$$\Psi_Z = \left[1 - \sum_{j \notin C} \frac{\mu'(Z)}{G_j''(k_{js}^*)} \right]^{-1} \left[\theta U'(Z) - \mu'(Z) \sum_{i \in C} (k_{is} - k_{is}^{\text{out}}) \right]. \quad (35)$$

This is the expression used in Proposition 2. $\square$

Benchmark validity conditions. The closed-form expressions above characterize the benchmark candidate when the following conditions hold: (i) contracted recipients have $x_i^* > 0$, equivalently $r_i > \rho + \lambda$; (ii) directed deployment is interior, with $G_i''(k_{is}) < 0$, so $\chi_i = 0$ and the solution to the k_{is} first-order condition satisfies $k_{is} > 0$; (iii) outside firms are on an interior differentiable response margin with $G_j''(k_{js}^*) < 0$; (iv) the aggregate feedback satisfies $0 \leq \sum_{j \notin C} \partial k_{js}^*(Z) / \partial Z < 1$; and (v) the contract rate is interior.

A.2 KKT conditions and corner regimes

To describe corner regimes, add the non-negativity constraints $x_i \geq 0$ and $k_{is} \geq 0$, with multipliers $v_i \geq 0$ and $\chi_i \geq 0$. Relative to Equation (23), define

$$\tilde{\mathcal{L}} = \mathcal{L} + \sum_{i \in C} v_i x_i + \sum_{i \in C} \chi_i k_{is}.$$

KKT system. Primal feasibility is given by the constraints in Equation (12), together with $x_i \geq 0$ and $k_{is} \geq 0$. The inequality multipliers satisfy

$$y_i \geq 0, \quad \lambda \geq 0, \quad v_i \geq 0, \quad \chi_i \geq 0, \quad \Psi_Z \in \mathbb{R}.$$

Complementary slackness is

$$\begin{aligned} y_i [\Pi_i(x_i, r_{is}, k_{is}, Z) - v_i(Z)] &= 0, \\ \lambda \left(\bar{X} - \sum_{i \in C} x_i \right) &= 0, \\ v_i x_i &= 0, \quad \chi_i k_{is} = 0. \end{aligned}$$

The stationarity conditions for the firm-specific controls are

$$\begin{aligned}\frac{\partial \tilde{\mathcal{L}}}{\partial r_{is}} &= x_i - y_i x_i = 0, \\ \frac{\partial \tilde{\mathcal{L}}}{\partial x_i} &= (r_{is} - \rho) - cx_i + y_i(r_i - r_{is}) - \lambda + v_i = 0, \\ \frac{\partial \tilde{\mathcal{L}}}{\partial k_{is}} &= y_i[G'_i(k_{is}) - \mu(Z) - r_i] + \Psi_Z + \chi_i = 0.\end{aligned}$$

The Z stationarity condition is unchanged except that y_i need not equal one:

$$\theta U'(Z) + \sum_{i \in C} y_i \left(\frac{\partial \Pi_i}{\partial Z} - \frac{\partial v_i}{\partial Z} \right) + \Psi_Z \left[\sum_{j \notin C} \frac{\partial k_{js}^*(Z)}{\partial Z} - 1 \right] = 0. \quad (36)$$

If outside firms are at an interior solution, their response is characterized by

$$\frac{\partial k_{js}^*}{\partial Z} = \frac{\mu'(Z)}{G_j''(k_{js}^*)}.$$

⁵⁰ Under the maintained concavity and regularity assumptions, this system provides the local characterization used in the benchmark regime and in the corner cases below.

Contracted-recipient benchmark regime. If $x_i > 0$ and $k_{is} > 0$, complementary slackness implies $v_i = \chi_i = 0$. Provided the contract rate is interior, the contract-rate condition then pins down $y_i = 1$, and the benchmark formulas above follow immediately. The aggregate-capacity constraint may still bind. This is the regime used for the propositions in the main text.

Firms outside the contracted set. If a candidate contract is not activated, the firm is excluded from C and is treated as an outside firm in the main text. No strategic capital is deployed, no directed action is contracted, and the contract rate is economically irrelevant. The propositions focus on contracted recipients because the empirical predictions concern firms and banks that receive strategic capital.

⁵⁰If an outside firm is at a boundary, the derivative is replaced by the relevant local one-sided derivative, or by zero when the boundary locally fixes k_{js}^* , as implied by that firm's own KKT conditions. The main benchmark assumes interior outside responses.

No induced action or corner deployment. Positive funding need not imply distorted real activity. If $\Psi_Z = 0$ and both the contract allocation and commercial benchmark are interior, the financier leaves the firm at $k_{is} = k_{is}^{\text{out}}$. If $\Psi_Z < 0$, the financier would prefer less external deployment than the commercial benchmark.

With no tied-financing constraint, $x_i > 0$ is compatible with the corner $k_{is} = 0$. Such a contract is funding without positive directed strategic-market activity, so it lies outside the directed-action benchmark.

Binding aggregate capacity. If $\lambda > 0$, aggregate funding capacity binds. A positive λ enters each contracted recipient's funding condition negatively, compressing funding all else equal:

$$x_i^* = \frac{r_i - \rho - \lambda}{c}.$$

The contract-rate cancellation survives because r_{is} still drops out of the funding first-order condition when $y_i = 1$.

A.3 First-best benchmark

The baseline model characterizes the strategic financier's contracting problem, not a welfare optimum. This subsection defines a first-best planner that directly chooses deployment and makes clear which externalities the strategic financier does and does not internalize.

Setup. A first-best planner chooses external deployment $k_{\ell s}$ for every firm $\ell \in \mathcal{I}$ and values the aggregate network stock at $W(Z)$, with $W' > 0$ and $W'' \leq 0$, separately from the friction-reduction channel $\mu(Z)$. As in the main model, ordinary home-market activity is pre-optimized into $A_\ell(r_\ell)$, which is constant in the external-deployment problem. The planner inherits the same financing-rent mechanism as the baseline model, so the comparison isolates the internalization of network externalities. Unlike the strategic financier, the planner problem does not separate contracted recipients from outside firms. The planner's external-deployment problem is

$$\max_{\{k_{\ell s}\}_{\ell \in \mathcal{I}}} \sum_{\ell \in \mathcal{I}} \left[A_\ell(r_\ell) + G_\ell(k_{\ell s}) - \mu(Z)k_{\ell s} - r_\ell k_{\ell s} \right] + W(Z), \quad (37)$$

where $Z = \sum_{\ell \in \mathcal{I}} k_{\ell s}$, subject to $k_{\ell s} \geq 0$ for all $\ell \in \mathcal{I}$.

First-best optimality. Let

$$Z^{FB} = \sum_{\ell \in I} k_{\ell s}^{FB}.$$

Consider an interior planner allocation satisfying the relevant second-order conditions, with $k_{\ell s}^{FB} > 0$ and $G_{\ell}''(k_{\ell s}^{FB}) < 0$ for each firm ℓ . Because the planner directly controls every firm's deployment, changing $k_{\ell s}$ changes Z one-for-one. The first-order condition internalizes how a higher Z lowers the bilateral-friction cost on aggregate strategic-market deployment:

$$G_{\ell}'(k_{\ell s}^{FB}) - \mu(Z^{FB}) - r_{\ell} - \mu'(Z^{FB})Z^{FB} + W'(Z^{FB}) = 0. \quad (38)$$

Using the policy-wedge notation

$$G_{\ell}'(k_{\ell s}) = \mu(Z) + r_{\ell} - \tau_{\ell s},$$

the first-best wedge is

$$\tau_{\ell s}^{FB} = -\mu'(Z^{FB})Z^{FB} + W'(Z^{FB}). \quad (39)$$

At a given aggregate stock, the social wedge is common across firms. The planner applies it to all firms, while the strategic financier contracts only with firms in C . Since $\mu'(Z) < 0$, the friction-reduction component is positive when $Z^{FB} > 0$. A positive wedge lowers the effective marginal cost of strategic-market deployment relative to the private benchmark.

Comparison with the strategic financier. The planner and the strategic financier differ in which payoffs enter the marginal choice. The strategic financier values the whole network stock through $\theta U(Z)$. It does not, however, maximize total firm surplus, so the friction-reduction term enters only through the participation constraints of contracted recipients. Ignoring outside-firm feedback for exposition, the strategic wedge is common across contracted recipients and, for any $i \in C$, is

$$\tau_{is}^* = \theta U'(Z) - \mu'(Z) \sum_{n \in C} (k_{ns} - k_{ns}^{\text{out}}), \quad \tau_{is}^{FB} = W'(Z) - \mu'(Z)Z.$$

The sum over $n \in C$ appears because a marginal increase in Z changes the participation value of all contracted recipients, not only recipient i . These objects need not be ordered. The planner internalizes the social value of lower bilateral frictions for all firms; the strategic financier internalizes its own network payoff and the compensation needed to contract with C . Which wedge is larger depends on $\theta U'(Z)$ versus $W'(Z)$, the contracted set, outside-firm feedback, and the allocation

being compared.⁵¹ The first-best benchmark is therefore a diagnostic of which externalities the financier does and does not internalize, not a monotone bound.

A.4 Dynamic incentive compatibility and relational contracts

The baseline model assumes that the strategic financier can enforce the required external action k_{is} through legally binding covenants. Let

$$\tau_{is}^{\text{static}} \equiv \tau_{is}^* = \Psi_Z$$

denote the wedge under perfect covenant enforcement. This subsection keeps the contracted-recipient benchmark regime, but replaces legal enforcement with relational enforcement through the threat of termination. In the binding dynamic incentive compatibility (DIC) regime, the financier may leave rents to sustain compliance, so the static rent-extraction result $y_i = 1$ need not apply.

Environment. The relationship repeats indefinitely with common discount factor $\beta \in (0, 1)$. I study a stationary extension of the static contract: the same $\{x_i, r_{is}, k_{is}\}$ is offered each period, and Z is not separately accumulated over time. A deviating firm takes the cheap financing but chooses k_{is}^{out} . The deviation is detected after one period, and future access to the contract ends.

The dynamic incentive compatibility constraint. Let

$$B_i(k, Z) = G_i(k) - \mu(Z)k - r_i k$$

denote the part of operating payoff that depends on strategic-market deployment after home-market activity has been pre-optimized. The one-period deviation gain from ignoring the directed-action requirement is

$$D_i(k_{is}, Z) \equiv B_i(k_{is}^{\text{out}}, Z) - B_i(k_{is}, Z).$$

Since k_{is}^{out} is the private optimum at fixed Z , $D_i(k_{is}, Z) \geq 0$, with equality only when the mandated action coincides with the private optimum. This is the operating payoff the firm recovers by ignoring the directed-action requirement. Financing terms cancel in D_i because the firm receives the same

⁵¹If a planner allocation is at a corner, the comparison is read through the corresponding KKT inequality rather than the interior first-order condition. A weak ranking can be obtained under restrictive conditions. If rate interiority holds, outside-firm feedback is shut down ($\sum_{j \notin C} \partial k_{js}^*(Z) / \partial Z = 0$), $k_{is}^{\text{out}} \geq 0$, and $\theta U'(Z) \leq W'(Z)$, then $\sum_{i \in C} (k_{is} - k_{is}^{\text{out}}) \leq \sum_{i \in C} k_{is} \leq Z$, so $\tau_{is}^{\text{FB}} \geq \tau_{is}^*$ at the common Z . Once these restrictions are relaxed, no global ranking follows.

financing rate whether it complies or deviates. Let

$$R_i = (r_i - r_{is})x_i$$

be the per-period financing rent. Stationary participation requires

$$R_i - D_i(k_{is}, Z) \geq 0, \quad [y_i] \quad (40)$$

with multiplier $y_i \geq 0$. Compliance gives present value $(R_i - D_i(k_{is}, Z))/(1 - \beta)$ relative to the outside option. A one-shot deviation gives R_i today and ends future surplus. The dynamic incentive compatibility constraint is therefore

$$\frac{R_i - D_i(k_{is}, Z)}{1 - \beta} \geq R_i,$$

or equivalently

$$\beta R_i = \beta(r_i - r_{is})x_i \geq D_i(k_{is}, Z). \quad [\xi_i] \quad (41)$$

Let $\xi_i \geq 0$ denote the DIC multiplier. When the DIC binds with $D_i(k_{is}, Z) > 0$, the financier leaves rents to preserve compliance: $R_i = D_i/\beta > D_i$. Participation is then slack, so complementary slackness gives $y_i = 0$. Static rent extraction would set $R_i = D_i$ and leave no continuation surplus, so costly directed action requires extra rent under relational enforcement. The relevant nontrivial relational regime is therefore the binding-DIC, slack-participation regime analyzed below.

If no admissible contract satisfies the DIC, induced action cannot be sustained and the financier must either relax the required action toward k_{is}^{out} or exit the relationship. The binding-DIC analysis assumes the required rent can be delivered through an interior rate choice; otherwise, the financier must relax the action or exit.

Interior wedge under binding DIC. Adding the stationary participation term $y_i[R_i - D_i(k_{is}, Z)]$ and the DIC term $\xi_i[\beta R_i - D_i(k_{is}, Z)]$ to the Lagrangian gives the rate condition

$$x_i - y_i x_i - \beta \xi_i x_i = 0, \quad (42)$$

so for a contracted recipient with $x_i > 0$,

$$y_i + \beta \xi_i = 1. \quad (43)$$

The x_i first-order condition is

$$(r_{is} - \rho) - cx_i + y_i(r_i - r_{is}) - \lambda + \beta\xi_i(r_i - r_{is}) = 0. \quad (44)$$

Using $y_i + \beta\xi_i = 1$, the contract rate again cancels and the funding formula is preserved:

$$x_i^* = \frac{r_i - \rho - \lambda}{c}. \quad (45)$$

Under the maintained interior-deployment assumption, so $\chi_i = 0$, the k_{is} condition is

$$y_i[G'_i(k_{is}) - \mu(Z) - r_i] + \Psi_Z - \xi_i \frac{\partial D_i(k_{is}, Z)}{\partial k_{is}} = 0. \quad (46)$$

Because

$$\frac{\partial D_i(k_{is}, Z)}{\partial k_{is}} = -[G'_i(k_{is}) - \mu(Z) - r_i] = \tau_{is},$$

this condition implies

$$\tau_{is}^{\text{dyn}} = \frac{\Psi_Z}{y_i + \xi_i}.$$

In the binding-DIC regime with positive deviation gain, participation is slack, so $y_i = 0$ and $\xi_i = 1/\beta$. Therefore the dynamic wedge is

$$\tau_{is}^{\text{dyn}} = \beta\Psi_Z = \beta\tau_{is}^{\text{static}}. \quad (47)$$

Thus relational enforcement scales the static wedge by β , holding the contracted set and static network multiplier Ψ_Z fixed. The extension does not re-solve the network fixed point under dynamic enforcement.

Proposition 3 (Local attenuation under relational contracts). *In this reduced-form relational-contract extension with common discount factor β , binding DIC, and slack participation, compliance sustained by the threat of relationship termination rather than by legal covenants attenuates the implementability wedge according to*

$$\tau_{is}^{\text{dyn}} = \beta\tau_{is}^{\text{static}} = \beta\Psi_Z.$$

The same funding formula and contract-rate cancellation are preserved.

Empirical support for the relational-contract assumption. The relational-contract interpretation requires persistent ties. I treat the available evidence as supportive but not identifying. Supplemental persistence tables show repeated SWF participation in PitchBook funding histories and add-on SWF injections in the bank sample. These patterns indicate durable investor–recipient relationships, while leaving the contractual mechanism itself unobserved.

Connecting the extensions. When the benchmark wedge is non-negative, the dynamic extension attenuates it:

$$\tau_{is}^{\text{dyn}} \leq \tau_{is}^{\text{static}}$$

because $\beta \in (0, 1)$. The first-best planner provides a contrasting benchmark because it directly controls all firms and internalizes the full network effect. Table 10 summarizes these regimes. The table is not an additional empirical test; it clarifies which assumptions strengthen or weaken the model’s directed-action prediction.

Table 10: Implementability Wedge Across Regimes

Regime	Wedge	Key determinant
First-best planner	$\tau_{ts}^{FB} = -\mu'(Z^{FB})Z^{FB} + W'(Z^{FB})$	Direct command over all firms
Static strategic financier	$\tau_{is}^{\text{static}} = \Psi_Z$	Covenant enforcement
Dynamic relational contract (binding DIC)	$\tau_{is}^{\text{dyn}} = \beta\Psi_Z$	Firm patience β
Purely commercial finance	$\tau_{ts} = 0$	No strategic objective or directed-action requirement

Note: The first-best wedge is evaluated at the first-best allocation. The first-best planner directly controls all firms; the strategic financier operates through contracts with a subset. The dynamic row reports the binding-DIC, slack-participation regime in an infinitely repeated relationship with discount factor β . The static and dynamic strategic-financier rows use the same static contracted set and static-equilibrium network multiplier Ψ_Z ; the dynamic row applies the reduced-form β attenuation rather than re-solving a separate dynamic network fixed point.

B Data Construction

B.1 PitchBook Co-Investment Panel

This subsection documents the PitchBook panel behind Section 3. The construction choices are part of the empirical design: the sample keeps deals with both SWF and non-SWF investors, drops same-country comparisons that would mechanically favor the SWF, and uses an exact normalized-name bridge to PitchBook investor IDs to avoid fuzzy-name false positives.

Source tables. I link five PitchBook tables: *Deal* (2.5M transactions with deal date, size, post-money valuation, revenue, type, and status), *Investor* (400K profiles with type, headquarters country, and AUM), *DealInvestorRelation* (3.0M investor–deal participation records), *CompanyInvestorRelation* (2.4M company–investor relationship records), and *Company* (1.0M firms with headquarters country, sector, founding year, and financials). All monetary values are in USD millions.

SWF classification. I start from PitchBook’s “Sovereign Wealth Fund” investor type, which includes 98 investors across 51 countries. I then exclude 28 entities that are not SWFs, such as state pension funds, regional development agencies, and charitable foundations, and add 7 government entities that function as de facto SWFs, including the Future Fund of Australia, the Hong Kong Monetary Authority, and the Monetary Authority of Singapore. The resulting classification is then filtered by the co-investment and valuation-discount regression sample; 20 SWF home countries contribute SWF investor–deal observations to the main regression sample. The replication files report the full investor list.

Co-investment identification. A deal enters the co-investment sample if it includes at least one SWF and at least one deep-pocket non-SWF investor, defined as PE/buyout, VC, asset manager, hedge fund, investment bank, university, family office, growth/expansion, government investor or mutual fund. Of 2.5M completed deals with valid dates, 6,789 satisfy this filter.

Panel filters.

1. *Investor join:* DealInvestorRelation and CompanyInvestorRelation are linked to Investor by the PitchBook investor identifier. Rows whose investor IDs do not appear in Investor are dropped rather than imputed, which is conservative but avoids assigning SWF status or investor home country using uncertain identifiers.
2. *Target-country assignment:* I use the company-level headquarters country when available. If it is missing, I fall back only to validated site-location tokens: exact country names are accepted, and exact US state names or postal abbreviations are mapped to the United States. Unvalidated strings remain missing.
3. *Same-country drop:* Remove investor–deal pairs where the target headquarters country equals the investor’s home country, avoiding the mechanical confound that firms do more business at home.

4. *Same-as-SWF drop*: Remove non-SWF investors from the same country as the deal's SWF, since these cannot contribute clean cross-country within-deal variation.
5. *Both-arm requirement*: After each filter, verify that every deal still contains at least one SWF and one non-SWF row; drop deals that lose an arm.
6. *Contamination filter*: For multi-round companies, drop later-round rows where the *same investor* (by ID) appeared in an earlier round, to avoid double-counting repeat relationships. Later-round rows with new investors are retained.
7. *Deduplication*: Remove exact duplicate investor–deal rows inherited from duplicates in the raw DealInvestorRelation table.

The final panel contains 9,646 investor–deal pairs across 2,362 deals and 2,308 companies, with 2,677 SWF rows and 6,969 non-SWF rows. Rows with unresolved target country are dropped before the cross-border filters because they cannot establish whether the investor and target are foreign to one another. Most of the attrition from the 6,789 candidate co-investment deals to the final 2,362 deals comes from enforcing clean cross-country comparisons, requiring both investor arms after those exclusions, and dropping contaminated later rounds with repeat investors. The regression samples are smaller because the financing proxies impose additional non-missingness restrictions: the main valuation-discount sample contains 3,456 investor–deal pairs across 833 deals, the common valuation-discount-plus-dilution sample contains 3,403 pairs across 821 deals, and the revenue subsample used in the P/Revenue horse race contains 1,614 pairs across 355 deals.

Outcome construction. I identify expansion events through two relationship routes after linking focal companies to PitchBook investor IDs by exact normalized name: lower-cased, trimmed company name must equal lower-cased, trimmed investor name. In route 1, the panel company appears as a new investor, acquirer, follow-on investor, or add-on sponsor in PitchBook's DealInvestorRelation table for expansion-type deals, including M&A, buyout/LBO, corporate asset purchase, joint venture, PE growth/expansion, corporate, seed, early-stage VC, later-stage VC, and PIPE. In route 2, the panel company appears in CompanyInvestorRelation with an active, former, acquirer, or add-on sponsor relationship. I combine the two routes and deduplicate them on company, target country, and date, which yields 12,221 unique expansion events. Focal IPO and PIPE rounds are excluded from the co-investment panel; PIPE can enter only as a later expansion event. I do not use fuzzy matching, alias expansion, or manual predecessor-name stitching, so any resulting measurement error should attenuate the estimates by missing true expansions rather than inflate them through false matches.

For each investor–deal pair, I construct two outcomes:

- *New entry* (Y_{ext}): equals 1 if the target expanded into the investor’s home country within 5 years after the deal, or through the end of the observed expansion window for later deals, and had no prior expansion activity in that country before the deal. Pairs with pre-deal activity are coded as zero (continuation, not new entry).
- *Count* ($\text{asinh}(\text{Count})$): inverse hyperbolic sine of the number of expansion events in the investor’s home country over the same observed post-deal window.

Valuation discount. For each deal d in industry sector s at year t :

$$\text{ValDiscount}_d = -\left(\log(\text{PostVal}_d) - \overline{\log(\text{PostVal})}_{s(d),t(d)}\right),$$

where the benchmark $\overline{\log(\text{PostVal})}_{s,t}$ is computed from the full PitchBook deal universe, not just the co-investment panel. The sign flip ensures that higher values correspond to lower valuations relative to peers, which I interpret as evidence of weaker outside financing options, that is, higher r_i . I winsorize the variable at the 1st and 99th percentiles and standardize it to unit SD. I construct dilution (= dealsize/postvaluation) analogously, winsorize it at the 99th percentile, and standardize it.

Pricing premium. The round-level pricing diagnostic in Table 18 uses the log post-money-valuation-to-revenue ratio as an alternative pricing measure. Revenue in PitchBook is deal-specific rather than taken from the static Company.dta snapshot. Some pre-revenue firms, mainly in biotech, mining exploration, and early-stage technology deals, report zero or negative revenue. I exclude those deals from the log price-to-revenue calculation and propagate the resulting missing values through that regression sample. This does not affect the main valuation-discount specification.

Pre-deal geographic orientation. For each investor–deal pair, I compute the target company’s pre-deal expansion share toward the investor’s home country: the number of expansion events in that country during the three years before the deal divided by the total number of expansion events across all countries in the same window. This variable controls for pre-existing geographic orientation and enters progressively interacted with the valuation discount.

Table 11 summarizes the final investor–deal panel used in the main PitchBook regressions. It connects the construction steps above to the paper’s within-deal design by showing the outcome

rates, SWF and non-SWF row counts, and the valuation-discount and dilution measures used in the mechanism checks.

Table 11: PitchBook Co-Investment Panel: Summary Statistics

	N	Mean	SD	p25	p50	p75
Panel A: Regression Variables (investor–deal level)						
New entry (5yr post-deal)	3,456	0.06	0.23	0.00	0.00	0.00
Expansion count (5yr)	3,456	0.48	2.94	0.00	0.00	0.00
asinh(expansion count)	3,456	0.16	0.57	0.00	0.00	0.00
SWF investor	3,456	0.28	0.45	0.00	0.00	1.00
Valuation discount (std.)	3,456	0.00	1.00	-0.72	-0.07	0.66
Dilution (std.)	3,403	-0.02	0.99	-0.89	-0.47	1.39
Deal size (\$M)	3,403	1025.11	2784.62	70.00	200.00	621.01
Pre-acq. geographic share	3,456	0.02	0.10	0.00	0.00	0.00
Baseline (country deal share)	3,456	0.15	0.19	0.01	0.03	0.44
Panel B: Panel Dimensions						
Investor–deal pairs (regression obs.)	3,456					
Deals	833					
Target companies	802					
SWF investor–deal pairs	954 (27.6%)					
Unique SWF investors	42					
SWF home countries	20					
Non-SWF investor–deal pairs	2,502					
<i>SWF Investment Breadth (per SWF investor):</i>						
Deals per SWF	23 (median 5, range 1–250)					
Target countries per SWF	5.9 (median 3)					
Sectors per SWF	3.4 (median 3)					
Panel C: Broader Clean Co-Investment Panel						
Investor–deal pairs	9,646					
Deals	2,362					
Target companies	2,308					
SWF investor–deal pairs	2,677					
Non-SWF investor–deal pairs	6,969					
SWF investors	48					
SWF home countries with any row	24					

Notes: Panel A reports investor–deal level variables used in the main regressions: new entry and expansion-count outcomes within 5 years of the deal, or through the end of the observed PitchBook expansion window for later deals, the SWF-investor indicator, the standardized valuation discount and dilution measures, deal size in USD millions, and the two pre-deal geographic-orientation controls. Panel B reports dimensions for the regression sample with non-missing valuation discount. Panel C reports the broader cleaned co-investment panel before imposing the valuation-discount restriction, with 9,646 investor–deal pairs, 2,362 deals, and 2,308 target companies. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

B.2 DealScan Syndicated Lending Panel

This subsection documents the construction of the bank–country–quarter syndicated lending panel from Refinitiv DealScan.

Raw data. The source file is the DealScan general-syndicate extract, which contains lender–tranche-level records for syndicated loan facilities worldwide. Each row represents one lender’s participation in one tranche. I retain the following fields: lender parent name and operating country, borrower country, deal and tranche identifiers, deal active date, tranche type, tranche amount (USD converted), lender share (USD), number of lenders, primary role, loan purpose fields, tranche maturity dates, tenor, average life, all-in spread drawn (bps), margin (bps), and fee variables (upfront, commitment, and annual fees in bps).

Lender parent assignment. DealScan assigns each lender to a parent entity. I apply one override: any row whose lender name contains “Merrill Lynch” (case-insensitive) is reassigned to the parent “Merrill Lynch,” regardless of DealScan’s current parent mapping. This preserves Merrill Lynch as a distinct treated unit in the pre-2008 period rather than merging Merrill-branded subsidiaries into post-acquisition parents such as BofA Securities.

Bank–country lending shares. The main outcomes do not use allocated loan volume. DealScan records lender participation directly, so new entry and deal count can be constructed from observed bank–tranche links. I use DealScan commitment amounts only to construct the pre-crisis bank–country lending-share control. For that control, I use reported lender shares when available and otherwise follow the standard lead-arranger allocation convention in DealScan ([Sufi, 2007](#); [Ivashina and Scharfstein, 2010](#)). Because this allocation rule is needed only for the pre-share control, not for the outcomes, the main bank results are not driven by allocated loan-volume measurement.

Aggregation to bank–country–quarter. Deal counts are summed within each (lender parent, borrower country, year-quarter) cell to produce the *book panel*. A *balanced panel* is then constructed by cross-joining all observed (bank, country) pairs with all observed quarters, filling missing cells with zero deal count.

Summary statistics. Table [12](#) documents the bank–country–quarter panel used in the crisis design. It reports the outcome variables, the ex ante SWF-status indicator, the bank-level write-down ratio, and the pair-level controls used in the fully controlled specification. The table connects the appendix construction to Table [5](#): the balanced panel contains many zero cells, so the paper reports new entry separately and estimates log count on the positive-lending subsample.

Table 12: DealScan Bank–Country–Quarter Panel: Summary Statistics

	N	Mean	SD	p25	p50	p75
Panel A: Regression Variables (bank–country–quarter level)						
Deal count	37,296	1.90	8.44	0.00	0.00	1.00
log(deal count), positive cells	10,761	1.148	1.058	0.000	1.099	1.792
New entry (LPM)	37,296	0.005	0.073	0.000	0.000	0.000
Lending share (within-bank quarter)	37,296	0.011	0.061	0.000	0.000	0.001
Write-down ratio (bank)	37,296	0.096	0.133	0.022	0.065	0.092
SWF status (country)	37,296	0.130	0.336	0.000	0.000	0.000
Post ($\mathbf{1}[t \geq 2007Q3]$)	37,296	0.429	0.495	0.000	0.000	1.000
WD $\times$ SWFStatus	37,296	0.013	0.062	0.000	0.000	0.000
WD $\times$ Pre-share	37,296	0.0012	0.0102	0.0000	0.0000	0.0001
WD $\times$ FTA	37,254	0.022	0.068	0.000	0.000	0.000
WD $\times$ Common language	36,960	0.018	0.073	0.000	0.000	0.000
WD $\times$ Political distance	34,804	0.164	0.324	0.012	0.060	0.165
Panel B: Panel Dimensions						
Bank–country–quarter observations	37,296					
Banks	29 (5 treated, 24 control)					
Borrower countries	163 total (14 SWF-status)					
Quarters in window	14 ($[-8, +5]$ around 2007Q3)					

Notes: Panel A reports bank–country–quarter level variables used in the main continuous DiD (Table 5). Outcomes are deal count, log deal count on positive cells, and the new-entry LPM. The bank-level write-down ratio and country-level SWF-status indicator define the treatment interaction WD $\times$ SWFStatus, and the additional pair-level controls (Pre-share, FTA, common language, political distance) enter the fully saturated specification interacted with Post. Panel B reports panel dimensions: 29 realized G-SIB bank parents (5 treated, 24 control), 163 borrower countries in total (14 SWF-status), and the $[-8, +5]$ quarter window around 2007Q3, that is, eight pre-treatment quarters and six quarters from 2007Q3 through 2008Q4. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

B.3 Bilateral FDI and Gravity Controls

The bilateral FDI data used in Sections 5.1 and related spillover tests are drawn from the World Bank Harmonized Bilateral FDI Database (Steenbergen et al., 2022). This subsection documents the cleaning steps applied to the raw database before estimation.

Source. The WBG-HBFDI database harmonizes four underlying sources: OECD Bilateral FDI Statistics (BMD4 and historical BMD3 series, 2005–2019), the IMF Coordinated Direct Investment Survey (CDIS, 2009–2019), UNCTAD Bilateral FDI Statistics (2001–2012), and China’s official FDI bulletins (2003–2019). The database covers 247 reporting economies and 247 partner economies over 2001–2021, with four variables per pair-year: inward/outward flows and inward/outward stocks, all in USD millions. I use inward flows as the outcome variable: for a reporting economy i (destination) and partner economy j (source), the inward-flow series $FDI_{ji,t}^{in}$ represents FDI flowing from j into i in year t .

Source deduplication. The harmonized database is constructed through a sequential augmentation process, where each source fills gaps left by higher-priority sources (Table 6 of Steenbergen et al. 2022). The priority hierarchy for inward flows is: (1) OECD-BMD4 from the host (reporting) economy; (2) OECD-BMD4 mirror from the partner economy; (3) China source; (4) UNCTAD; (5) IMF-CDIS. Despite this sequential design, the 2022 vintage of the database contains approximately 39,500 duplicate pair-year observations, concentrated entirely in 2018–2019 where OECD and CDIS coverage overlaps. Of these, roughly 7,500 pairs have conflicting values from different sources (e.g., OECD-BMD4 reports +532 USD million while CDIS reports –108,582 USD million for the same pair-year). I deduplicate by retaining the highest-priority source for each pair-year, breaking ties by preferring the reporter economy (“r”) over the partner economy (“p”) as the data origin. This reduces the raw database from 588,058 to 548,487 observations.

Estimation window and gravity controls. The cleaned WBG-HBFDI database covers 2001–2021, but the FIRRMA regressions use the 2014–2020 panel after imposing the event window around the 2018 shock and merging the required controls. Bilateral gravity variables are drawn from CEPII Gravity V2022 (Head and Mayer, 2014): FTA membership, destination GDP (used in log current-USD form), bilateral trade flows (Comtrade), and GDP growth. Diplomatic distance is measured as ideal-point distance from the UN General Assembly voting database (Bailey, Strezhnev and Voeten, 2017). CEPII trade coverage is the binding control and is available through 2020, so the main panel ends in 2020. The 2014 GDP growth observation uses the 2013–2014 lag. Pairs

missing gravity controls are dropped from the estimation sample.

Winsorization. Bilateral FDI flows contain extreme values driven by financial-center conduit transactions, such as Luxembourg→Ireland flows above \$100 billion in a single year. To limit the influence of those outliers without dropping observations, I winsorize inward flows at the 1st and 99th percentiles before applying the asinh transformation. I do this within each estimation window, including the main 2014–2020 panel and each placebo window.

Negative and zero flows. Bilateral FDI flows can be negative when disinvestment, profit repatriation, or intercompany loan repayment exceeds new investment in a given year. In the control-complete FIRRMA estimation sample (2014–2020, excluding US pairs), approximately 17% of observations have negative inward flows and 58% are exact zeros. The inverse hyperbolic sine transformation accommodates both: $\text{asinh}(x) = \ln(x + \sqrt{x^2 + 1})$ is defined for all real x , approximately linear near zero, and approximately $\ln(2|x|)$ for large $|x|$.

Summary statistics. Table 13 documents the source×destination×year panel used in the FIRRMA design. The table links the raw FDI data to Table 6 by reporting the transformed FDI outcome, the exposure measures, the gravity controls, and the two balance controls after the 2014–2020 sample restriction. It also makes the estimation support explicit: the crowd-in estimates are not estimated on the full WBG-HBFDI universe, but on the subset with the controls needed for the source×year and pair-fixed-effect specification.

B.4 BIS Locational and Consolidated Banking Statistics

The main spillover test in Table 8 uses BIS Locational Banking Statistics (LBS), which report claims by office location or residence. The robustness check in Table 30 uses BIS Consolidated Banking Statistics (CBS), which report claims by the nationality of the parent bank.

SWF investment treatment universe. SWF bank equity injections are identified from the Global SWF (GSWF) and Sovereign Wealth Fund Institute (SWFI) transaction databases, supplemented and cross-checked against the IFSWF member-fund deal disclosures, S&P Capital IQ, Mergermarket, and PrivCo. Across all sectors, the cleaned and deduplicated unlisted private-market SWF deal universe covers about \$1.28 trillion of cross-border investment over 2002–2022. Within that broader universe, I retain direct bank equity injections, dropping asset-management mandates,

Table 13: FIRRMA Bilateral FDI Panel: Summary Statistics

	N	Mean	SD	p25	p50	p75
Panel A: Regression Variables (source $\times$ destination $\times$ year, 2014–2020)						
Bilateral FDI inflow (\$M, winsorized)	17,591	104.3	649.9	0.0	0.0	2.2
asinh(bilateral FDI inflow)	17,591	0.639	3.248	0.000	0.000	1.520
Exposure: high-dilution (destination)	17,591	0.237	0.350	0.000	0.091	0.400
Exposure: count-based (destination)	17,591	0.332	0.375	0.000	0.244	0.667
Exposure: value-weighted (destination)	17,591	0.358	0.429	0.000	0.080	0.927
Post ($\mathbf{1}[\text{year} \geq 2018]$)	17,591	0.433	0.495	0.000	0.000	1.000
FTA membership	17,591	0.204	0.403	0.000	0.000	0.000
log(destination GDP)	17,591	20.279	1.487	19.509	19.849	21.228
Destination GDP growth	17,591	0.007	0.104	-0.046	0.015	0.067
asinh(bilateral trade)	17,591	0.510	0.967	0.001	0.034	0.501
Diplomatic distance (UN voting)	17,591	0.854	0.754	0.135	0.667	1.512
US trade share (destination)	17,591	0.090	0.065	0.042	0.079	0.118
Pre-FIRRMA FDI level (std., destination)	17,591	-0.082	0.891	-0.654	-0.204	0.724
Panel B: Panel Dimensions						
Source–destination–year observations	17,591					
Source–destination pairs	3,093					
Source countries	224					
Destination countries	22					
Destinations with positive HD exposure	9					
Years (FIRRMA shock 2018)	2014–2020 (4 pre + 3 post)					

Notes: Panel A reports source $\times$ destination $\times$ year variables used in the FIRRMA crowd-in regression (Table 6). The outcome is asinh of the winsorized bilateral FDI inflow. Three exposure measures are constructed at the destination level from pre-2018 SWF deal characteristics: high-dilution (above-median dilution), count-based, and value-weighted. Gravity controls follow CEPII Gravity V2022 plus UN-voting diplomatic distance. The balance controls are U.S. trade share and standardized pre-FIRRMA FDI level; the latter is the destination's 2014–2017 mean asinh bilateral inflow, standardized across destinations. Panel B reports panel dimensions: source–destination–year observations, source/destination country counts, and the count of destinations with positive high-dilution exposure. The 2014–2020 window provides four pre-FIRRMA years and three post years. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

insurance investments, and fund-of-fund commitments. The resulting bank-treatment master contains 107 bank-injection deals totaling \$127 billion across 55 lender–SWF-country pairs and 13 SWF source countries. The BIS spillover design treats these as pair-level aggregate relationships: if a lender-country $\times$ SWF-country pair receives multiple disclosed injections, I sum them into a single total intensity and anchor the cohort at the pair’s first injection date. After BIS reporting, disclosure, timing, and matching restrictions, 12 treated lender–borrower pairs enter the stacked DiD used in Table 8. The firm-channel treatment in Table 9 further restricts this universe to unlisted non-financial deals over 2002–2022, totaling \$1.13 trillion in Appendix Table 14.

Data cleaning. The BIS introduced the instrument breakdown (loans vs. debt securities) in 1995, so I restrict the sample to 1995–2024. Negative stock positions (fewer than 0.1% of observations), which arise from exchange-rate valuation adjustments and short positions in debt securities ([Bank for International Settlements, 2019](#)), are floored at zero because they are reporting and valuation artifacts rather than economically meaningful negative lending stocks. The main stock specifications use log claims on strictly positive Q4 stock positions; in the matched annual regression sample this restriction drops one loans-and-deposits observation and nine non-bank-loan observations. Annual aggregation uses the Q4 end-of-year stock and drops pair-years without a Q4 observation, following [Avdjiev et al. \(2020\)](#). For the flow column, I use reported LBS loan flows, sum quarterly observations to the pair-year level, and scale the annual flow by the pair’s pre-treatment mean loan stock; I do not construct flows from changes in stocks. Matching uses the pre-treatment mean and linear trend of $\log(\text{loans})$, standardized within the treated-control pool.

Summary statistics. Table 14 documents the matched stacked-cohort panel used in the BIS spillover regression. It connects the BIS data construction to Table 8 by reporting both the LBS outcomes and the SWF equity-injection treatment intensity after coverage, timing, and matching restrictions. The table also shows the cost of the design: the matched panel is cleaner than the raw treatment universe, but it covers only the country pairs for which BIS reporting and pre-treatment matching variables are available.

C Additional PitchBook Results

This appendix collects robustness checks and supporting evidence for the PitchBook co-investment results in Section 3. The checks address the main threats to the within-deal design: SWF country-sector composition, follow-on financing, sector heterogeneity, round-level pricing diagnostics,

Table 14: BIS LBS Stacked-Cohort Panel: Summary Statistics

	N	Mean	SD	p25	p50	p75
Panel A: Regression Variables (cohort × pair × year, ±6-year window)						
Total cross-border claims (\$M)	3,259	22976.2	48149.6	973.0	5896.0	20846.0
Loans & placed deposits (\$M)	3,258	18392.3	37497.6	684.0	5094.4	17571.5
Non-bank loans (\$M)	3,250	5442.6	11114.3	285.5	1922.5	5764.8
log(total claims)	3,259	8.325	2.275	6.880	8.682	9.945
log(loans & deposits)	3,258	8.093	2.332	6.528	8.536	9.774
log(non-bank loans)	3,250	7.141	2.072	5.654	7.561	8.660
SWF intensity (\$bn)	3,259	0.267	1.730	0.000	0.000	0.000
SWF intensity × Post (\$bn)	3,259	0.144	1.276	0.000	0.000	0.000
Post (1[year ≥ cohort treat])	3,259	0.541	0.498	0.000	1.000	1.000
Treated pair indicator	3,259	0.048	0.213	0.000	0.000	0.000
Panel B: Panel Dimensions						
Cohort–pair–year observations	3,259					
Cohorts (lender–borrower × first-treatment year)	12					
Reporting countries (full annual LBS panel)	31					
Borrower countries (full annual LBS panel)	218					
Treated lender–borrower pairs	12					
Control pairs (Mahalanobis-matched)	172					
Event-time window	[−6, 6] around cohort treatment					
Panel C: SWF Investment Treatment (treated pairs only)						
Bank-channel deals (raw)	107					
Bank-channel lender–SWF-country pairs (raw)	55					
SWF home countries with bank injections	13					
Total bank-channel SWF equity (\$bn)	127					
SWF equity per pair, post-disclosure filter (\$bn)	48	2.65	4.76	0.11	0.43	2.73
Pairs in stacked DiD (LBS-reporting + disclosed)	12					
SWF equity per stacked-DiD pair (\$bn)	12	5.57	5.99	0.73	4.15	8.32
Non-financial unlisted SWF deals (2002–2022, cross-channel)	2,636					
Non-financial unlisted firm-channel pairs	378					
Total non-financial unlisted SWF investment (\$bn)	1134					
Firm-investment intensity per pair (\$bn)	378	3.00	9.02	0.10	0.46	1.87

Notes: Panel A reports cohort $\times$ pair $\times$ year variables used in the BIS spillover regression (Table 8). Outcomes are total cross-border claims, loans and placed deposits, and non-bank loans, both in raw USD millions and as log-transformed positive stocks. The treatment intensity is total SWF equity injected (USD billions) into banks domiciled in the lender country, by SWFs from the borrower country, interacted with the post-treatment indicator. When a pair receives multiple disclosed injections, those amounts are summed into a single pair-level total and switched on from the pair's first injection date. Panel A pools treated and matched-control rows; conditional-on-treatment summaries are reported in Panel C. Panel B reports panel dimensions. Panel C summarizes the SWF investment treatment universe: bank-channel deals and lender–SWF-country pairs (raw, before LBS coverage and timing restrictions) with the per-pair SWF equity-investment distribution, and the analogous non-financial firm-investment universe used as treatment in the cross-channel test (Table 9). Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

country concentration, random SWF-label assignment, tail treatment, outcome-window choice, mechanically overlapping same-country co-investors, alternative scale and control-share mechanisms, and destination-fixed matched comparisons.

C.1 SWF Representation

Table 15 reports a composition check within the main co-investment sample. The table asks whether the relative-valuation proxy is associated with greater SWF presence among retained institutional investor slots in rounds that already contain both SWF and non-SWF investors. Higher valuation discount is associated with greater SWF representation: a one-standard-deviation increase raises the SWF share of investor slots by 2.4 percentage points and the probability that a given investor row is an SWF by 3.5 percentage points. Dilution, measured as deal size over post-money valuation, gives similar estimates. I interpret these estimates as evidence on relative syndicate composition, not as a direct dollar-volume test, because PitchBook reports investor-level ticket amounts too sparsely. The main text therefore uses the table only as supporting evidence; the central test asks whether the SWF–non-SWF directed-expansion differential loads on valuation discount.

Table 15: Valuation Discount, Dilution, and SWF Representation

	SWF Share (1)	is_SWF (2)	SWF Share (3)	is_SWF (4)
Val. Discount	0.0236*** (0.0058)	0.0359*** (0.0081)		
Dilution			0.0318*** (0.0053)	0.0417*** (0.0076)
Year FE	Yes	Yes	Yes	Yes
SE	Company	Company	Company	Company
Observations	833	3,456	821	3,403

Note: Columns (1) and (3) are deal-level (one row per deal); the dependent variable is the SWF share of investor slots in the deal. Columns (2) and (4) are investor–deal level; the dependent variable is an indicator for whether the investor slot is held by an SWF. Columns (1)–(2) use valuation discount; columns (3)–(4) use dilution. All columns include year fixed effects. SE are clustered by company in all columns. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.2 SWF country-sector portfolio tilts

Figure 5 reports sector tilts for the 20 SWF home countries in the PitchBook valuation-discount sample, using a common 2005–2025 base period. For each home country and primary sector, I compute the sector’s share of cross-border SWF investor–deal observations for SWFs headquartered in that country and subtract the corresponding share among cross-border investor–deal observations

backed by non-government deep-pocket private investors. The measure is count-based, not value-weighted, and is not residualized by year.

The tilts vary sharply across SWF home countries, but they do not map mechanically to home-country revealed comparative advantage. Figure 6 focuses on country–sector cells with large positive tilts and benchmarks each cell against the SWF home country’s early revealed comparative advantage (RCA) in the mapped sector. For country c and PitchBook sector s , I compute

$$RCA_{c,s} = \frac{X_{c,s}/X_c}{X_{W,s}/X_W},$$

where $X_{c,s}$ is the country’s mapped goods-plus-services exports in sector s , X_c is total goods-plus-services exports, and W denotes the corresponding world totals. RCA above one means that the sector accounts for a larger share of the country’s exports than of world exports. The benchmark uses the geometric mean of annual RCA over 2005–2007, the earliest service-export window with broad country coverage.⁵²

The sector mapping is deliberately non-overlapping. For mixed sectors such as IT, mapped goods and services components are summed before forming $X_{c,s}$ and $X_{W,s}$. Goods exports use four-digit SIC codes: Energy uses coal, oil and gas, and petroleum products; Materials uses natural resources, basic materials, and non-pharmaceutical chemicals; B2C uses consumer goods, cosmetics, household appliances, and audio/video goods; B2B uses industrial products, machinery, transport equipment, and non-medical instruments; IT uses computer hardware, communications equipment, and electronic components; and Healthcare uses pharmaceuticals and medical or ophthalmic instruments. Services exports enter through balance-of-payments categories: Financial Services uses financial services only; IT uses telecommunications, computer, and information services; B2B uses other business services; Healthcare uses health services plus health-related travel; and B2C uses travel and personal-cultural services net of health components. This construction is a revealed-comparative-advantage proxy, not a complete production taxonomy.

The benchmark separates high-tilt cells into aligned and non-aligned cases. Most high-tilt cells with available export data have RCA below one. A smaller set lines up with revealed comparative advantage: China overweights Materials, while Oman, Qatar, and the United Arab Emirates overweight Energy. The more common low-RCA cells include Kuwait’s finance overtilt, Saudi Arabia’s B2C overtilt, and China’s energy and B2B overtilts. The narrow reading is that

⁵²World totals are constructed from the same source files used for country totals. Sanity checks confirm that no mapped sector total exceeds the corresponding goods, services, or combined export total. The seven mapped PitchBook sectors cover about 88% of world goods-plus-services exports in 2005–2009, so the RCA denominator remains the full export basket rather than only mapped sectors.

sector tilts proxy SWF-specific sector weights, not a simple rule of investing only where the home country already has an export advantage. In strong sectors, directed business links can complement existing domestic clusters. In weak sectors, they can be read as consistent with diversification or capability-building motives, subject to the coarse PitchBook-sector aggregation.





Figure 5: SWF Country-Sector Portfolio Tilts

Notes: Rows are the 20 SWF home countries that contribute SWF investor–deal observations to the PitchBook valuation-discount sample. Colored cells report each country’s pooled 2005–2025 cross-border sector share minus the corresponding 2005–2025 cross-border benchmark share for non-government deep-pocket private investors, in percentage points. SWF and benchmark portfolios are counted at the investor $\times$ deal level; when multiple SWFs from the same home country join the same deal, each contributes one portfolio row before pooling to home-country shares. Deal-investor rows with Add-on Sponsor status are excluded from this descriptive portfolio-count exercise. The final column reports the number of cross-border SWF investor–deal observations used to compute each SWF home country’s sector shares. The measure is count-based.



Cells: 2005-2025 overtilt pp / 2005-2007 export RCA. Clean non-overlapping goods/service map. Dagger: fewer than 25 SWF deals.

Figure 6: High-tilt sectors and early revealed comparative advantage, all SWF home countries
Notes: Rows are the 20 SWF home countries that contribute SWF investor-deal observations to the PitchBook valuation-discount sample. Colored cells report 2005–2025 country–sector overtilts of at least 5 percentage points and indicate whether the SWF home country’s 2005–2007 export RCA in the mapped PitchBook sector is above or below one. France has no highlighted cell because all of its exact sector tilts are below 5 percentage points in absolute value. Grey cells indicate that no early export-RCA benchmark is available for the mapped sector. Cell labels report the portfolio overtilt in percentage points and the export-RCA value; a dagger (†) after the overtilt number marks cells from home countries with fewer than 25 total cross-border SWF investor-deal observations. Deal-investor rows with Add-on Sponsor status are excluded from the descriptive portfolio-count measure. Goods sectors use a clean non-overlapping four-digit SIC mapping; services use balance-of-payments services categories. Constructing the same export-RCA diagnostic with cumulative SWF portfolios through 2015 or 2018 gives qualitatively similar patterns, although very early sector-tilt support is limited. In the regression design, sector tilts are constructed from prior cumulative individual-SWF portfolios before the focal deal year, and the valuation-discount regressions retain the original sample for comparability.

C.3 Follow-on Financing

Table 16 reports three pieces of evidence on follow-on financing and investor continuation. Panel A asks whether an SWF investor is more likely to return to the same company. Let c index target companies, j investors, and d financing rounds. Columns (1)–(2) estimate

$$\text{Return}_{cjd}^{5y} = \alpha_d + \beta \text{SWF}_{jd} + X'_{cjd}\gamma + \varepsilon_{cjd},$$

where Return_{cjd}^{5y} equals one if the exact investor appears in any later completed PitchBook deal record for the same company within five years. This broad follow-on measure includes later private financings, IPOs, PIPEs, M&A, joint ventures, asset purchases, and corporate transactions when PitchBook records the same investor-company relationship again. It is therefore a continuation-relationship measure, not a fresh-capital-only reinvestment measure. The original-round fixed effect compares SWF and non-SWF investors in the same original financing round. Column (1) includes no controls. Column (2) adds investor j 's own prior follow-on rate, computed from earlier portfolio-company relationships whose five-year return window had closed before round d . This is an investor-history control, not a measure of the target firm's prior fundraising. Column (2) also controls for same-home-country co-investors in the original round. The SWF coefficient is 3.46 percentage points in column (1) and 2.51 percentage points in column (2), so SWF investors are more likely than same-deal non-SWF co-investors to return.

Panel B then classifies the same valuation-discount focal-deal universe by realized later deal activity. Among these 833 focal deals, 410 have no later completed private PitchBook deal with usable investor links, 256 have later private deals but no later SWF participation, and 167 have at least one later private deal with SWF participation. The re-up firms are selected: their median post-money valuation is \$1.0 billion, their mean valuation discount is lower, and they raise more later rounds. Their cross-border SWF-row expansion rate is 4.3 percent, compared with 5.4 percent among firms with later private deals but no SWF re-up and 1.4 percent among firms with no later private deal.

Panel C interacts the main within-deal specification with the full A1/A2/B split, using A1 as the omitted group. For outcome m , I estimate

$$\begin{aligned} Y_{idt}^m = & \beta_1 \text{SWF}_{idt} + \beta_2 (\text{SWF}_{idt} \times \text{ValDiscount}_{dt}) \\ & + \beta_3 (\text{SWF}_{idt} \times A2_d) + \beta_4 (\text{SWF}_{idt} \times B_d) \\ & + \beta_5 (\text{SWF}_{idt} \times \text{ValDiscount}_{dt} \times A2_d) + \beta_6 (\text{SWF}_{idt} \times \text{ValDiscount}_{dt} \times B_d) \\ & + \alpha_{dt} + \alpha_{it} + \alpha_{i \times h(d)} + \varepsilon_{idt}, \end{aligned}$$

where i indexes the investor's home country, d indexes the focal financing round, t is the deal year, and $h(d)$ is the target firm's headquarters country. The deal-by-year fixed effects α_{dt} are implemented as deal fixed effects because each deal occurs in one year, α_{it} are investor-home-country $\times$ deal-year fixed effects, and $\alpha_{i \times h(d)}$ are investor-home-country $\times$ target-HQ-country fixed effects. Because this follow-on margin concerns repeated activity around the SWF relationship, the first outcome is any post-deal expansion rather than first entry. The $SWF \times$ valuation-discount slope is 0.0368 in A1, 0.0969 in A2, and 0.1826 in B for any expansion; the largest difference is therefore in B, whose slope is 0.1458 above A1. The same ordering holds for aggregate count. I interpret these estimates as selected continuation heterogeneity rather than a causal re-up effect.

C.4 Sector Heterogeneity

Table 17 asks whether the PitchBook result is a sector artifact. It drops each primary industry sector in turn and re-estimates the main specification. The new-entry estimate remains positive under every omission, but it is most sensitive to Consumer Products and Services (B2C): omitting that sector lowers the coefficient from 0.048 to 0.023 and removes conventional statistical significance. The count-based appendix outcome remains positive and statistically significant under every omission. This pattern suggests that consumer-facing targets contribute meaningfully to the extensive-margin result, while aggregate post-deal activity is less sector-specific.

C.5 Round-Level Pricing Diagnostics

PitchBook is a round-level pricing setting: same-round investors generally share a headline valuation, and the data do not consistently report investor-level security terms or ticket amounts. This subsection therefore uses round-level pricing measures as diagnostics. It asks whether the directed-expansion result is explained by observable valuation and financing-package measures, or instead loads on the sector-year valuation discount used as the relative-valuation proxy.

SWF-backed deals have higher valuation multiples. Figure 7 plots average $\log(P/\text{Revenue})$ for SWF and non-SWF deals within sector $\times$ year cells.⁵³ In roughly 66% of cells, SWF deals are priced higher.⁵⁴

⁵³ $P/\text{Revenue}$ denotes the deal's post-money valuation divided by the target firm's latest-twelve-month revenue.

⁵⁴ Using deal size relative to revenue as an alternative pricing proxy also shows that SWFs tend to participate in larger financing rounds.

Table 16: Follow-on Financing, Re-up Anatomy, and Directed Expansion

Panel A: Exact investor continuation

	Same-Investor Return	
	Base	+ Controls
SWF investor	0.0346*** (0.0111)	0.0251** (0.0109)
Investor baseline return controls	—	✓
Other same-country investor controls	—	✓
Observations	3,456	3,456
Deals / firms	833 / 802	833 / 802
Original-round FE	✓	✓
SE clustered by company	✓	✓

Panel B: Focal SWF-backed deal groups

Group	(1) A1	(2) A2	(3) B
Definition	No later financing	Later, no SWF	Later, SWF re-up
Focal SWF deals	410	256	167
Share of focal deals	49.2%	30.7%	20.0%
Median post-money (\$M)	551	788	1,000
Median employees	678	400	420
Mean valuation discount	0.29	-0.00	-0.23
Mean later rounds	0.00	1.81	2.95
Later-stage VC share	20.2%	61.7%	66.5%
SWF-row expansion rate	1.4%	5.4%	4.3%

Panel C: Ex-post group heterogeneity


Outcome	(1) Any Expansion	(2) asinh(Count)
SWF × Val. Discount (A1)	0.0368** (0.0157)	0.0504** (0.0225)
SWF × Val. Discount × A2	0.0601* (0.0321)	0.0925* (0.0531)
SWF × Val. Discount × B	0.1458*** (0.0491)	0.1807*** (0.0608)
Implied A2 slope	0.0969*** (0.0301)	0.1429*** (0.0534)
Implied B slope	0.1826*** (0.0472)	0.2311*** (0.0577)
Largest VD slope	B (0.1826)	B (0.2311)
SWF × group controls	✓	✓
Control mean	0.1291	0.2101
Observations	3,456	3,456
Deals / firms	833 / 802	833 / 802
Deal-by-year FE (deal FE)	✓	✓
Investor-home-country × deal-year FE	✓	✓
Investor-home-country × target-HQ-country FE	✓	✓
SE clustered by company	✓	✓

Note: All panels use the valuation-discount sample. Panel A reports investor–deal regressions where the outcome equals one if the exact investor appears in any later completed PitchBook deal record for the same company within five years. This broad continuation measure includes later private financings, IPOs, PIPEs, M&A, joint ventures, asset purchases, and corporate transactions when PitchBook records the same investor-company relationship again. Panel B collapses the sample to focal SWF-backed deals and classifies them using later completed private PitchBook deal records within five years, excluding IPO and PIPE later rounds and requiring usable investor links. A1 has no such later private deal; A2 has a later private deal but no SWF participation; B has at least one later private deal with SWF participation. SWF-row expansion rate is measured on cross-border SWF investor rows. Panel C returns to the investor–deal regression sample; Any Expansion equals any post-deal expansion into the investor’s home country within five years. A1 is the omitted group; A2 and B interactions are differences relative to A1, and the implied slopes add the A2 or B interaction to the A1 slope. A2 and B are selected after the focal deal, so Panel C is descriptive heterogeneity, not a causal design. SE clustered by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 17: Sector Heterogeneity: Leave-One-Sector-Out

	New Entry	asinh(Count)
Baseline (all sectors)	+0.048***	+0.111***
Business (B2B)	+0.055***	+0.115***
Consumer (B2C)	+0.023	+0.104***
Energy	+0.049***	+0.105***
Fin. Services	+0.051***	+0.112***
Healthcare	+0.051***	+0.119***
Info. Technology	+0.054**	+0.090***
Materials	+0.048***	+0.112***

Notes: Each row drops one PitchBook primary industry sector (Business Products and Services (B2B), Consumer Products and Services (B2C), Energy, Financial Services, Healthcare, Information Technology, or Materials and Resources) and re-estimates the baseline specification (SWF $\times$ Valuation Discount with deal-by-year FE, implemented as deal FE, plus investor-home-country $\times$ deal-year FE and investor-home-country $\times$ target-HQ-country FE; SE clustered by company). New Entry = $\mathbf{1}[\text{no pre-deal activity in investor country AND post-deal activity}]$. Asinh(Count) uses all post-deal activity (not conditioned on pre=0). Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Figure 7: SWF-backed Deals Command Higher Valuations Across Sector $\times$ Year Cells

Notes: Each dot represents one sector $\times$ year cell. The x -axis is the average $\log(\text{Post-money valuation/Revenue})$ of non-SWF deals in that cell; the y -axis is the average of SWF deals in the same cell. Cells require at least one SWF deal and five non-SWF deals. Dot size is proportional to the number of SWF deals in the cell. Dots above the 45-degree line indicate that SWF deals in that sector-year were priced higher than non-SWF deals.

Round terms do not explain directed expansion. Table 18 compares the sector–year valuation discount with $\log(P/\text{Revenue})$, $\log(\text{DealSize}/\text{Revenue})$, and an up-round indicator. The alternative round-level measures have little independent power once valuation discount is included, while the $\text{SWF} \times \text{valuation-discount}$ coefficient remains positive. The result therefore loads on the relative-valuation proxy, not on observable valuation multiples or financing-package size.

Table 18: Valuation Discount and Round-Level Terms

	Proxy Only		Horse Race with Val. Discount	
	New Entry (1)	asinh(Cnt) (2)	New Entry (3)	asinh(Cnt) (4)
<i>Panel A: $\log(P/\text{Rev})$ as round-level valuation multiple — revenue subsample</i>				
$\text{SWF} \times \log(P/\text{Rev})$	-0.0234 (0.0206)	-0.0104 (0.0296)	-0.0100 (0.0195)	0.0120 (0.0261)
$\text{SWF} \times \text{Val. Discount}$			0.1025*** (0.0354)	0.1660*** (0.0430)
Observations	1,618	1,618	1,614	1,614
<i>Panel B: $\log(\text{DealSize}/\text{Rev})$ as financing quantity proxy — revenue subsample</i>				
$\text{SWF} \times \log(\text{DealSize}/\text{Rev})$	0.0279 (0.0271)	0.0459 (0.0326)	0.0119 (0.0219)	0.0204 (0.0288)
$\text{SWF} \times \text{Val. Discount}$			0.1134*** (0.0350)	0.1804*** (0.0442)
Observations	1,575	1,575	1,575	1,575
<i>Panel C: Up Round as round-level pricing direction — full sample</i>				
$\text{SWF} \times \text{Up Round}$	-0.0278 (0.0329)	-0.0502 (0.0518)	-0.0089 (0.0324)	-0.0061 (0.0500)
$\text{SWF} \times \text{Val. Discount}$			0.0473*** (0.0164)	0.1105*** (0.0245)
Observations	3,456	3,456	3,456	3,456
Deal-by-year FE (deal FE)	✓	✓	✓	✓
Investor-home-country $\times$ deal-year FE	✓	✓	✓	✓
Investor-home-country $\times$ target-HQ-country FE	✓	✓	✓	✓

Note: Investor–deal level regressions. Each panel horse-races valuation discount against one observable round-level measure. Panel A uses $\log(\text{post-money valuation} / \text{revenue})$, a valuation multiple available in the revenue subsample. Panel B uses $\log(\text{deal size} / \text{revenue})$, a financing-package-size measure relative to firm scale. Panel C uses an up-round indicator, equal to one if the round price exceeded the previous round. These variables do not recover investor-specific contract rates. Columns (1)–(2) include only the round-level measure interacted with the SWF indicator; columns (3)–(4) add the $\text{SWF} \times \text{valuation-discount}$ interaction. All specifications include deal-by-year FE, implemented as deal FE, investor-home-country $\times$ deal-year FE, and investor-home-country $\times$ target-HQ-country FE. SE clustered by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.6 Leave-One-SWF-Country-Out

Figure 8 drops each SWF investor home country in turn. This check addresses the main external-validity concern in PitchBook: the identifying variation is concentrated in a small number of active SWF countries. The coefficient remains positive under every omission, preserving the sign of the result, but the sensitivity to Singapore and the UAE qualifies how broadly the PitchBook estimate should be generalized.




Figure 8: Leave-One-SWF-Country-Out: SWF × Val. Discount Coefficient

Notes: Each row drops one SWF investor home country and re-estimates the full-FE valuation-discount specification (deal-by-year FE, implemented as deal FE, plus investor-home-country × deal-year FE and investor-home-country × target-HQ-country FE). Orange dashed line: full-sample baseline. Error bars: 95% CIs; SE clustered by company. Singapore and the UAE together account for 73% of SWF investor–deal rows and 88% of SWF new-entry events in the valuation-discount sample. Dropping the UAE attenuates the estimate, while dropping Singapore amplifies it; both omissions leave the coefficient positive and statistically significant.

C.7 Permutation Inference

Table 19 asks whether the PitchBook pattern could be generated by randomly labeling same-round investors as SWFs. The exercise leaves the realized expansion outcomes unchanged. It also leaves each deal’s investor set, investor home countries, valuation discount, and number of SWF investors

fixed. Within each deal, I randomly reassign the SWF labels and recompute two statistics. Panel A uses a model-free statistic: the difference between SWF and non-SWF new-entry rates after scaling each group by its baseline country deal share. Panel B re-estimates the main regression and stores the placebo coefficient on SWF $\times$ valuation discount.

Table 19: Permutation Inference

	Observed	Null Mean [SD]	RI $p_{\text{two-sided}}$
<i>Panel A: Model-free ratio statistic</i>			
SWF vs. non-SWF ratio	0.877	-0.034 [0.018]	0.000
<i>Panel B: Regression coefficient</i>			
New Entry	0.0481***	0.0001 [0.0082]	0.000
asinh(Count)	0.1110***	-0.0001 [0.0147]	0.000
Permutations		1,000	

Note: Within-deal random reassignment of SWF investor labels (1,000 permutations). Panel A tests a model-free ratio statistic: $E[Y|SWF = 1]/E[\text{baseline}|SWF = 1] - E[Y|SWF = 0]/E[\text{baseline}|SWF = 0]$; RI p = share of permuted ratios with $|\text{ratio}^{\text{perm}}| \geq |\text{ratio}^{\text{obs}}|$. Panel B permutes the SWF indicator and re-estimates the main specification (deal-by-year FE, implemented as deal FE, plus investor-home-country $\times$ deal-year FE and investor-home-country $\times$ target-HQ-country FE; SE clustered by company); RI p = share of permuted SWF $\times$ Val. Discount coefficients with $|\hat{\beta}^{\text{perm}}| \geq |\hat{\beta}^{\text{obs}}|$. Stars on observed coefficients reflect asymptotic p -values from the main regression. Two-sided p -values reported. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.8 Winsorization Sensitivity

Table 20 tests whether the valuation-discount result is driven by extreme valuation residuals. It re-estimates the main specification with no winsorization and with a tighter two-sided winsorization cutoff. The coefficient is stable across the non-winsorized and winsorized specifications, so the valuation-discount result is not an artifact of a small number of extreme observations.

C.9 Alternative Outcome Windows

Table 21 reports two checks on the timing of the PitchBook outcome. Panel A asks whether the new-entry result depends on the post-deal window used to construct the dependent variable, while keeping the main regression sample fixed. The 3-year, 5-year, and 7-year estimates are all positive and significant at 1%, with coefficients rising from 0.0424 to 0.0481 to 0.0502.

Panel B addresses right-censoring by holding the baseline five-year new-entry outcome fixed and restricting the sample to deals with at least 3, 5, or 7 years of observable post-deal expansion history. The SWF $\times$ valuation-discount coefficient remains positive and similar in magnitude: 0.0565 in the at-least-3-year sample, 0.0545 in the at-least-5-year sample, and 0.0680 in the at-

Table 20: Winsorization Sensitivity

	Outcome: New Entry (5yr)		
	Main: winsor p1/p99 (1)	No winsorization (2)	Winsor p5/p95 (3)
SWF $\times$ Val. Discount	0.0481*** (0.0164)	0.0466*** (0.0162)	0.0462*** (0.0163)
Deal-by-year FE (deal FE)	Yes	Yes	Yes
Investor-home-country $\times$ deal-year FE	Yes	Yes	Yes
Investor-home-country $\times$ target-HQ-country FE	Yes	Yes	Yes
Observations	3,456	3,456	3,456

Note: Each column re-standardizes valuation discount under a different tail treatment and re-estimates the main specification (deal-by-year FE, implemented as deal FE, plus investor-home-country $\times$ deal-year FE and investor-home-country $\times$ target-HQ-country FE; SE clustered by company). Column (1) is the main-text specification, which winsorizes valuation discount at the 1st and 99th percentiles. Column (2) uses the raw valuation discount without winsorization. Column (3) winsorizes valuation discount at the 5th and 95th percentiles. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

least-7-year sample. Precision declines as the restrictions reduce the regression sample to 2,460, 1,270, and 770 investor–deal rows, respectively.

C.10 Same-Country Non-SWF Co-Investors

The baseline sample drops non-SWF investors headquartered in the same country as the deal’s SWF because their outcome, expansion into the investor’s home country, is mechanically identical to the SWF outcome. If an SWF-backed firm expands into the SWF home market, a same-country non-SWF co-investor inherits the same outcome without adding independent within-deal variation. Table 22 reports two diagnostics: an expanded sample that keeps those rows, and a stricter sample that drops the entire financing round whenever it contains a same-country non-SWF co-investor. The estimates remain positive across entry and expansion count, so the result is not driven by how these mechanically overlapping co-investors are handled.

C.11 Alternative Mechanisms: Scale and Control Share

Table 23 asks whether the valuation-discount interaction is simply capturing firm scale or SWF control rights in the financing round. Panel A first compares valuation discount with scale proxies. Smaller firms, measured by employee counts or deal-level revenue, show the expected standalone pattern: the SWF-directed entry gradient is larger for smaller targets. But these scale interactions lose explanatory power once valuation discount enters the same regression, while the valuation-discount interaction remains positive. The panel then residualizes valuation discount on log

Table 21: Alternative Outcome Windows and Follow-up Restrictions

Panel A: Alternative post-deal outcome windows

	New Entry		
	3-year (1)	5-year (2)	7-year (3)
SWF $\times$ Val. Discount	0.0424*** (0.0159)	0.0481*** (0.0164)	0.0502*** (0.0165)
Deal-by-year FE (deal FE)	Yes	Yes	Yes
Investor-home-country $\times$ deal-year FE	Yes	Yes	Yes
Investor-home-country $\times$ target-HQ-country FE	Yes	Yes	Yes
SE clustered by company	Yes	Yes	Yes
Observations	3,456	3,456	3,456

Panel B: Minimum observable post-deal follow-up

	Baseline 5-year New Entry		
	≥ 3 -year (1)	≥ 5 -year (2)	≥ 7 -year (3)
SWF $\times$ Val. Discount	0.0565** (0.0230)	0.0545 (0.0367)	0.0680 (0.0668)
Deal-by-year FE (deal FE)	Yes	Yes	Yes
Investor-home-country $\times$ deal-year FE	Yes	Yes	Yes
Investor-home-country $\times$ target-HQ-country FE	Yes	Yes	Yes
SE clustered by company	Yes	Yes	Yes
Non-SWF mean	0.0815	0.0725	0.0838
Observations	2,460	1,270	770

Note: Panel A constructs the new-entry outcome using different post-deal windows (3, 5, or 7 years) and re-estimates the main specification on the same regression sample as the main table. Panel B uses the baseline five-year new-entry outcome and restricts the sample by minimum observable follow-up, measured relative to the last observed PitchBook expansion date (February 9, 2025). New entry equals one if the target expands into the investor's home country within the relevant window and there was no prior activity in that country before the focal deal. All columns include deal-by-year FE (implemented as deal FE), investor-home-country $\times$ deal-year FE, and investor-home-country $\times$ target-HQ-country FE; SE are clustered by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 22: Same-Country Non-SWF Co-Investors: Stability Check

	Expanded Sample		Drop Same-Country Deals	
	New Entry (1)	asinh(Count) (2)	New Entry (3)	asinh(Count) (4)
SWF $\times$ Val. Discount	0.0408*** (0.0134)	0.0923*** (0.0205)	0.0491*** (0.0170)	0.1131*** (0.0280)
SWF	-0.0081 (0.0149)	-0.0198 (0.0205)	0.0002 (0.0330)	0.0341 (0.0463)
Deal-by-year FE (deal FE)	Yes	Yes	Yes	Yes
Investor-home-country $\times$ deal-year FE	Yes	Yes	Yes	Yes
Investor-home-country $\times$ target-HQ-country FE	Yes	Yes	Yes	Yes
SE clustered by company	Yes	Yes	Yes	Yes
Observations	3,737	3,737	2,920	2,920

Notes: Investor–deal level panel. Dependent variables are new market entry and asinh(expansion count), measured within 5 years post-deal or through the end of the observed PitchBook expansion window for later deals. Columns (1)–(2) use the expanded sample that keeps non-SWF investors whose home country matches the deal’s SWF home country. Columns (3)–(4) drop the entire financing round if it contains any such same-country non-SWF co-investor. Outcomes are constructed from the same PitchBook expansion-event matching procedure used in the main specification. All columns use the sector-year valuation-discount benchmark, include deal-by-year FE, implemented as deal FE, investor-home-country $\times$ deal-year FE, and investor-home-country $\times$ target-HQ-country FE, and cluster SE by company. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

employees, log revenue, log firm age, and dilution. The residual component still predicts directed entry, including when the component interactions are included directly. This residual result is useful because those controls explain much of the raw valuation-discount variation; the remaining component is therefore not just firm size, maturity, or dilution.

Panel B tests control-share interpretations. Total dilution, equal-share dilution (dilution divided by the number of retained investors in the round), and SWF-bloc dilution (dilution multiplied by the SWF investor share in the round) each point in the predicted direction when entered alone. The bloc measure is a round-level sovereign-presence proxy, not an individual investor stake. Once valuation discount is included, however, these ownership/control proxies become small and statistically insignificant, while valuation discount remains positive. I read the table as evidence that relative valuation contains information about outside financing options that is not exhausted by simple size or control-share proxies.

C.12 Destination-Fixed Matched Comparison

The within-deal specification in Section 3 compares SWF and non-SWF co-investors in the same financing round, which absorbs deal-level heterogeneity. It does not compare entry into the same SWF destination across SWF-backed and non-SWF-backed firms. Table 24 provides that destination-fixed matched comparison with exact-cell regressions on cross-border SWF-backed

Table 23: Alternative Mechanisms Behind the Valuation-Discount Gradient

	Employees		Revenue		Residual VD	
	Alone (1)	+ VD (2)	Alone (3)	+ VD (4)	Alone (5)	+ Comp. (6)
<i>Panel A: Scale and residual valuation-discount information</i>						
SWF × Val. Discount		0.0299* (0.0173)		0.1071*** (0.0337)		
SWF × Small employee base	0.0520*** (0.0194)	0.0284 (0.0205)				
SWF × Small revenue			0.0546* (0.0304)	-0.0171 (0.0294)		
SWF × Residual Val. Discount					0.0486** (0.0225)	0.0385* (0.0196)
Component interactions	–	–	–	–	–	✓
Observations	3,122	3,122	1,614	1,614	1,472	1,472
Deals	718	718	355	355	319	319
	Total Dilution		Equal-Share Dilution		SWF-Bloc Dilution	
	Alone (1)	+ VD (2)	Alone (3)	+ VD (4)	Alone (5)	+ VD (6)
<i>Panel B: Ownership and control-share proxies</i>						
SWF × Val. Discount		0.0417*** (0.0159)		0.0471*** (0.0164)		0.0479*** (0.0163)
SWF × Total dilution	0.0350** (0.0145)	0.0168 (0.0126)				
SWF × Dilution / investors			0.0211* (0.0118)	0.0062 (0.0103)		
SWF × SWF-bloc dilution					0.0200* (0.0118)	0.0034 (0.0101)
Observations	3,403	3,403	3,403	3,403	3,403	3,403
Deals	821	821	821	821	821	821
Deal-by-year FE (deal FE)	✓	✓	✓	✓	✓	✓
Investor-home-country × deal-year FE	✓	✓	✓	✓	✓	✓
Investor-home-country × target-HQ-country FE	✓	✓	✓	✓	✓	✓
SE clustered by company	✓	✓	✓	✓	✓	✓

Notes: Investor–deal level regressions. The dependent variable is new entry into the investor’s home country within five years, conditional on no prior activity there before the focal deal. All columns include deal-by-year FE, implemented as deal FE, investor-home-country × deal-year FE, and investor-home-country × target-HQ-country FE, with standard errors clustered by company. Continuous proxies, including valuation discount in the horse-race columns, are winsorized and standardized within the relevant regression sample, so valuation-discount coefficients in this table are within-column standardized effects rather than direct numerical replications of the main-table coefficient. Columns (1)–(2) require non-missing employees and valuation discount; columns (3)–(4) require non-missing deal-level revenue and valuation discount; columns (5)–(6) require non-missing employees, revenue, firm age, dilution, and valuation discount. Small employee base is negative log company employees; small revenue is negative log deal-level revenue. Residual valuation discount is the residual from regressing standardized valuation discount on log employees, log revenue, log firm age, and dilution; column (6) also includes the interactions of those components with the SWF indicator. Total dilution is deal size divided by post-money valuation. Equal-share dilution divides total dilution by the number of retained institutional investors in the round. SWF-bloc dilution multiplies total dilution by the SWF share of retained investors in the round; it is a round-level sovereign-bloc proxy, not an individual investor’s own implied stake. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

(deal, SWF country) pairs.

The treatment unit is a (deal, SWF country) pair. Each cross-border SWF-backed deal contributes one treatment observation per distinct SWF country among its investors, so that a deal with two SWF co-investors from different countries generates two treatment observations with the same deal covariates but different target countries. Same-country SWF pairs, such as Bpifrance backing a French firm, are dropped. Columns (1)–(2) use a baseline control pool of non-SWF-backed deals paired to treated SWF countries within exact cells defined by treated SWF country, three-year deal-year bin, PitchBook primary sector, and a seven-region headquarters grouping. Columns (3)–(4) tighten the comparison by requiring the control deal to have at least one non-SWF investor from the same country as the SWF investor in the treated pair. The outcome is an indicator for expansion into the treated SWF country within five years, constructed using the same expansion types and active-investor statuses as the main analysis (Appendix B.1).

In the baseline control pool, 8.80% of SWF-backed pairs expand into the SWF country within five years, compared with 0.89% of matched control pairs. The exact-cell estimate is 7.36 percentage points (SE 2.10), and 6.32 percentage points (SE 1.91) after adding firm and deal controls. This shows the observational entry gap relative to similar non-SWF-backed deals with the same potential SWF destination.

The country-anchor control pool raises the control entry rate to 7.65%, indicating that part of the baseline gap reflects a general investor-country anchor. The SWF coefficient nevertheless remains positive: 5.61 percentage points (SE 2.60) with exact-cell fixed effects only, and 3.07 percentage points (SE 1.31) after adding controls. The tighter specification is therefore best read as a decomposition: SWF-backed firms enter SWF home countries more often than comparable firms, but the SWF-specific level difference is smaller than the gross baseline gap.

D Additional Bank Results

This appendix reports robustness checks for the bank syndicated lending results in Section 4. The checks address the main design risks: whether write-down exposure proxies for pre-crisis bank characteristics, whether any single SWF destination drives the estimate, and whether alternative SWF-treatment definitions change the estimand.

D.1 Pre-Crisis Balance: Write-Down Ratio

Figure 9 asks whether the bank write-down ratio is simply proxying for pre-crisis bank characteristics. It regresses WriteDownRatio on five standardized covariates computed from the bank–

Table 24: Destination-Fixed Matched Comparison: SWF-Backed vs. Observationally Similar Non-SWF-Backed Firms

	Baseline Control Pool		Country-Anchor Control Pool	
	FE Only	FE + Controls	FE Only	FE + Controls
	(1)	(2)	(3)	(4)
SWF-backed deal-country pair	7.36*** (2.10)	6.32*** (1.91)	5.61** (2.60)	3.07** (1.31)
SWF-backed mean (%)	8.80	8.80	11.33	11.33
Matched control mean (%)	0.89	0.89	7.65	7.65
Exact-cell FE	Yes	Yes	Yes	Yes
Firm/deal controls	No	Yes	No	Yes
Control has non-SWF investor from same SWF country	No	No	Yes	Yes
Observations	117,060	117,060	1,525	1,525
SWF-backed pairs	432	432	309	309
Exact cells	238	238	137	137

Notes: Estimation is on the PitchBook deal universe, restricted to deals with a computable post-money valuation discount and to firms whose name matches an investor name in PitchBook (so expansion events can be detected on both sides). Treatment unit is a (deal, SWF country) pair: a deal with SWF co-investors from multiple distinct countries contributes one treatment observation per country. Same-country SWF pairs (e.g., Bpifrance backing a French firm) are dropped. Controls are non-SWF-backed deals paired to treated SWF countries within exact cells defined by treated SWF country, three-year deal-year bins, the seven PitchBook primary sectors, and a seven-region HQ grouping (North America, Western Europe, Asia, Latin America, MENA, Oceania, Other). Columns (3)–(4) additionally require the control deal to have at least one non-SWF investor from the same country as the SWF investor in the treated pair. Coefficients are percentage points from linear probability models. Firm/deal controls are standardized valuation discount, asinh deal size, asinh employee count, year founded, and the cumulative count of distinct prior-investor home countries observed before the deal year (computed at the company-year level, so within-year deals share a value). Outcome: indicator for expansion into the treated SWF country within five years, using the same expansion types and active-investor statuses as the main analysis (Appendix B.1). All specifications restrict to fully observed five-year outcome windows and report SE clustered by company and treated SWF country. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

country–quarter panel over 2004Q3–2007Q2. Four covariates are pre-window levels, and the fifth is a within-window growth measure that compares 2004Q3–2005Q4 with 2006Q1–2007Q2. The balance check is reassuring overall: all five coefficients are statistically insignificant, including pre-crisis SWF-country lending share. The main specification additionally absorbs bank–country level differences through bank×country fixed effects and the $WD \times Pre\text{-}share \times Post$ control. Taken together, the figure argues against the most direct continuation story that distressed banks were already systematically more tilted toward SWF-country lending before the crisis.




Figure 9: Bank Pre-Crisis Balance: WriteDownRatio on Standardized Covariates

Notes: Cross-sectional OLS of the bank-level WriteDownRatio on five standardized pre-crisis covariates (N=29 realized G-SIB bank parents). Pre-crisis covariates aggregate the bank–country–quarter panel over 2004Q3–2007Q2: total cross-border lending (USD), deal count, number of borrower countries, share of total lending to SWF home countries, and a pre-crisis lending growth measure defined as the log of total lending in 2006Q1–2007Q2 relative to 2004Q3–2005Q4. Each covariate is standardized to mean zero and unit standard deviation. Error bars: 95% confidence intervals from OLS standard errors. Coefficients near zero indicate balance; any residual level differences are absorbed by the Bank×Country and $WD \times Pre\text{-}share \times Post$ controls in the main specification.

D.2 Leave-One-SWF-Country-Out

Figure 10 drops each SWF destination country in turn from the full-controls log-count DID specification. This addresses the bank-sample concentration concern: with a finite set of SWF destinations, the estimated reallocation could in principle be carried by one destination market. The coefficient

remains positive in every omission and remains statistically significant at the 5% level except when Australia is omitted. The exercise therefore rules out a sign reversal or a result driven mechanically by one country.




Figure 10: Leave-One-SWF-Country-Out: Deal Count

Notes: Each point is the WD coefficient from the $\log(\text{deal count})$ specification, restricted to cells with positive deal count, omitting one SWF country from the destination set. Dashed line: full-sample baseline. Error bars: 95% CIs; SE clustered by bank. Full-controls specification with three-way FEs.

D.3 Alternative SWF Treatment Definitions

The baseline bank specification uses an ex ante destination-level treatment. SWFStatus_c equals one for countries that had active cross-border sovereign investors before the crisis: the 7 countries whose funds later injected equity into treated banks plus 8 additional pre-2007 SWF countries in the defined set.⁵⁵ This definition is chosen before observing the realized crisis-era bank–SWF matches. It asks whether banks with greater demand for sovereign capital redirected lending toward sovereign-capital destinations, not only toward the specific countries that ultimately provided their equity.

⁵⁵The defined set contains 15 countries in total. Libya has no lending observations in the realized panel, leaving 14 effective SWF-status destinations in the estimation sample.

Table 25 replaces the baseline SWF-status set with the 20 PitchBook SWF home countries that contribute to the valuation-discount regression sample. This check uses the same full-controls specification as the main bank table. The coefficients remain positive across both reported outcomes. New entry is positive but imprecise, 0.0092; the log-count coefficient is 0.7863 and significant at the 1% level. The PitchBook-based definition therefore preserves the sign of the bank result, although the strongest response remains on the deal-frequency margin. I treat this as a destination-set diagnostic rather than a preferred specification, because the baseline bank definition is chosen ex ante for the pre-crisis syndicated-lending setting.

Table 25: Alternative SWF Treatment Definition: PitchBook SWF Home Countries

	New Entry (1)	log(Count) (2)
<i>Panel A: PitchBook SWF home countries (20 countries)</i>		
WD $\times$ SWF (PitchBook) $\times$ Post	0.0092 (0.0072)	0.7863*** (0.2794)
Bank $\times$ Country FE	✓	✓
Bank $\times$ Quarter FE	✓	✓
Country $\times$ Quarter FE	✓	✓
SE clustered by bank	✓	✓
Observations	34,510	10,010

Note: Bank–country–quarter panel around 2007Q3, $[-8, +5]$ quarters. The table replaces the baseline ex ante SWF-status destination set with the 20 PitchBook SWF home countries that contribute to the valuation-discount regression sample. All columns use the full-controls specification with bank $\times$ country, bank $\times$ quarter, and country $\times$ quarter fixed effects; SE clustered by bank. Columns include WD $\times$ Post interactions with pre-lending share, FTA, common language, and UN-voting political distance. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

E FIRRMA Spillover Robustness

This appendix reports robustness checks for the FIRRMA crowd-in results in Section 5.1. The checks address the main limits of the design. Appendix E.1 expands the destination universe beyond the PitchBook SWF-home destinations used in the main specification. Appendix E.2 reports pre-treatment balance on destination characteristics, Appendix E.3 documents non-U.S. high-dilution SWF deal volume around FIRRMA, and Appendix E.4 displays the cross-country exposure distribution. Appendix E.5 tests placebo shock dates and drops oil-exporting SWF countries, and Appendix E.6 reports leave-one-exposed-SWF-country-out estimates.

E.1 Full-Destination Main Specification

Table 26 repeats the FIRRMA specification on the broader non-U.S. destination universe rather than restricting destinations to the SWF home countries used in the main text. This is a demanding comparison because the additional destinations have zero SWF exposure by construction and need not be on the same investment margin as SWF home countries. The estimates remain positive and statistically significant across the high-dilution, count-based, and value-weighted exposure measures. The full-control high-dilution coefficient is smaller than in the SWF-home destination design, 0.318 rather than 0.663, which is consistent with dilution from adding many zero-exposure destinations outside the main support. I therefore read this table as a robustness check on sign and statistical stability, not as the preferred magnitude.

Table 26: FIRRMA Exposure and Third-Party FDI Inflows: Full Destination Sample

	High-Dilution Exposure				Count	Value
	(1)	(2)	(3)	Target Sources (4)	(5)	(6)
SWF US Exposure $\times$ Post	0.310*** (0.098)	0.333*** (0.098)	0.318*** (0.098)	0.795*** (0.306)		
SWF US Exposure (count) $\times$ Post					0.224*** (0.086)	
SWF US Exposure (value) $\times$ Post						0.238*** (0.083)
US Trade Share $\times$ Post			0.178 (0.116)	0.102 (0.244)	0.202* (0.116)	0.202* (0.116)
Pre-FIRRMA FDI Level $\times$ Post			-0.088*** (0.019)	-0.154*** (0.039)	-0.092*** (0.019)	-0.091*** (0.019)
asinh(Bilateral Trade)		0.105 (0.155)	0.127 (0.155)	0.157 (0.256)	0.129 (0.155)	0.133 (0.155)
Diplomatic Distance		0.007 (0.059)	0.015 (0.059)	0.082 (0.117)	0.015 (0.059)	0.016 (0.059)
Other gravity controls	–	✓	✓	✓	✓	✓
Source $\times$ Year FE	✓	✓	✓	✓	✓	✓
Pair FE	✓	✓	✓	✓	✓	✓
SE clustered by pair	✓	✓	✓	✓	✓	✓
Source countries	224	224	224	47	224	224
Destination countries	172	172	172	172	172	172
Observations	112,541	112,541	112,541	38,495	112,541	112,541
R^2	0.296	0.296	0.296	0.291	0.296	0.296

Note: Same specifications as Table 6, estimated on the broader non-U.S. destination universe (172 destinations) as a robustness check. Dependent variable: asinh(bilateral FDI inflow), 2014–2020. Columns (3)–(6) include U.S. trade share $\times$ Post and standardized pre-FIRRMA FDI level $\times$ Post. All columns include source $\times$ year and pair FE. Standard errors are clustered by country pair. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

E.2 Pre-Treatment Balance: Exposure

Figure 11 asks whether the high-dilution exposure measure is proxying for destination-level fundamentals within the SWF-home destination sample. It regresses exposure on six pre-2018 destination covariates: GDP, GDP growth, U.S. trade share, the level and slope of pre-2018 FDI inflows, and the number of source partners. This balance check is necessarily at the destination level because the treatment varies at that level; bilateral covariates are absorbed by source $\times$ year and pair fixed effects in the main regression. The figure shows two imbalances. Exposed SWF destinations have greater pre-2018 U.S. trade exposure, which is expected because the treatment is built from U.S.-targeted SWF deals. They also have lower pre-2018 bilateral FDI inflow levels. The other covariates, including destination GDP, GDP growth, pre-2018 FDI growth, and the number of source partners, are not statistically distinguishable across exposure levels. I therefore read the balance test as a reason to include U.S.-trade-share and pre-FDI-level post interactions in the main table, and to read the estimates together with the dynamic and placebo checks.

E.3 Non-U.S. high-dilution deal volume

Figure 12 reports annual non-U.S. high-dilution SWF deal counts around FIRRMA. Panel A scales deal counts by the number of SWF home countries in each exposure group; Panel B reports raw annual counts. Non-U.S. high-dilution deal volume rises after 2018 in both groups, including among SWF countries with positive pre-FIRRMA U.S. high-dilution exposure. I read the figure as descriptive evidence that the relevant non-U.S. deal margin expanded after FIRRMA.

E.4 Exposure by SWF Country

Figure 13 plots the cross-country distribution of high-dilution FIRRMA exposure for the 23 non-U.S. PitchBook SWF home countries. Nine countries have positive exposure and fourteen have zero. Within the positive group, three countries reach 67–100% mechanically: Ireland and New Zealand sit at 100% on a single pre-2018 SWF deal each, and South Korea at 67% on four deals. The substantive variation comes from France (40% on 18 deals), Australia (17% on 9), China (14% on 11), Qatar (13% on 29), Malaysia (9% on 17), and Singapore (10.5% on 176). The zero-exposure comparison group includes UAE (n=24) and Kuwait (n=11) alongside smaller zero-exposure SWF home countries. This dispersion across positive-exposure countries motivates the leave-one-out exercise in Appendix Figure 14.




Figure 11: FIRRMA Exposure Balance: Pre-2018 Destination Covariates

Notes: Cross-sectional OLS of destination-level high-dilution exposure on six standardized pre-2018 destination covariates. Sample: 22 effective non-U.S. SWF home destinations in the main FIRRMA panel (Table 6); pre-2018 means computed over 2014–2017. Each covariate is standardized to mean zero and unit standard deviation. Error bars: 95% confidence intervals. The U.S. trade-share coefficient is positive, consistent with the fact that exposure is built from pre-2018 SWF deals targeting the United States. The negative coefficient on pre-2018 FDI inflow levels means that more-exposed destinations start from lower bilateral inflows on average; pair fixed effects absorb fixed bilateral levels in the main regression, while the event-study and placebo tests address differential timing.




Figure 12: Non-U.S. High-Dilution SWF Deal Volume around FIRRMA

Notes: Annual PitchBook counts of SWF-backed high-dilution deals targeting non-U.S. companies, 2014–2022. High-dilution deals are those with dilution above the median among SWF co-investment deals. Positive-exposure countries are SWF home countries with at least one U.S.-targeted high-dilution deal in 2014–2017; zero-exposure countries have none. Panel A divides annual counts by the number of SWF home countries in the exposure group. Panel B reports raw annual counts. The dashed vertical line marks 2018, the FIRRMA enactment year. The figure is descriptive and does not condition on controls or fixed effects.




Figure 13: FIRRMA High-Dilution Exposure by SWF Country

Notes: Bars report each SWF home country's pre-FIRRMA high-dilution exposure, defined as the share of 2014–2017 SWF co-investment deals with observed dilution that both targeted the United States and involved above-median dilution. Countries are sorted by exposure. Labels report the total number of pre-2018 SWF home-country deal observations. High-dilution exposure shares are computed over observations with non-missing dilution.

E.5 Placebo Shocks and Oil-Exporter Drop

Table 27 tests two confounds in the FIRRMA design. The placebo columns move the shock date earlier while keeping the panel entirely pre-FIRRMA and holding fixed the realized 2014–2017 exposure measure, asking whether that realized exposure predicts similar pre-2018 FDI shifts when the shock date is moved earlier. The placebo estimates are much smaller than the baseline and statistically insignificant: -0.008 for a 2015 placebo shock and -0.130 for a 2016 placebo shock, compared with the baseline estimate of 0.663. The oil-exporter exclusion asks whether the result is driven by concurrent oil-related shocks in Gulf and other commodity-linked SWF countries, including Norway. Dropping oil-exporter SWF countries leaves a positive and statistically significant coefficient of 0.561. The table therefore separates the FIRRMA interpretation from two obvious alternatives: earlier timing patterns in the realized exposure and oil-linked SWF shocks.

Table 27: FIRRMA Spillover: Robustness

	Baseline (1)	Placebo 2015 (2)	Placebo 2016 (3)	No Oil Exporters (4)
SWF US Exposure (high-dilution) $\times$ Post	0.663*** (0.215)	-0.008 (0.238)	-0.130 (0.262)	0.561** (0.235)
asinh(Bilateral Trade)	-0.707** (0.302)	-0.372 (0.255)	-0.637** (0.290)	-0.945*** (0.322)
Diplomatic Distance	-0.249 (0.211)	-0.134 (0.169)	-0.270 (0.196)	-0.324 (0.217)
Gravity controls (GDP, FTA)	✓	✓	✓	✓
Balance controls	✓	✓	✓	✓
Source $\times$ Year FE	✓	✓	✓	✓
Pair FE	✓	✓	✓	✓
SE clustered by pair	✓	✓	✓	✓
Observations	17,591	17,147	14,787	14,310
R^2	0.363	0.377	0.386	0.393

Note: Column (1): baseline with full gravity and balance controls (high-dilution exposure). Columns (2)–(3): placebo shocks at 2015 and 2016, respectively, using the same realized 2014–2017 exposure measure as the baseline but restricting the estimation window to pre-FIRRMA years ending in 2017. Column (4): oil-exporter SWF countries dropped (UAE, Saudi Arabia, Qatar, Kuwait, Norway, Oman, Bahrain, Brunei). Dependent variable: asinh(bilateral FDI inflow). Balance controls are U.S. trade share $\times$ Post and standardized pre-FIRRMA FDI level $\times$ Post. All columns include source $\times$ year and pair FEs. SE clustered by country pair, as in the main FIRRMA specification; because treatment varies at the destination-exposure level, inference should be read together with the placebo-timing and leave-one-country-out checks. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

E.6 Leave-One-SWF-Country-Out

Figure 14 drops each positively exposed high-dilution SWF destination country in turn. This check addresses the finite-country concern in the macro FDI design while focusing on the countries that

identify the high-dilution exposure coefficient. The coefficient remains positive across all omissions, so the result is not driven by a single exposed destination. The movement across omissions is still informative: countries whose omission shifts the estimate most are the destinations carrying more of the identifying variation, so the figure should be read as a sensitivity diagnostic rather than as a claim that every SWF country contributes equally.




Figure 14: Leave-One-SWF-Country-Out: FIRRMA Spillover

Notes: Each row drops one positively exposed high-dilution SWF destination country and re-estimates the balance-control high-dilution specification from Table 6, column (3). Orange dashed line: full-sample column (3) baseline. Error bars: 95% CIs; SE clustered by country pair.

F BIS Spillover Robustness

This appendix reports robustness checks for the BIS cross-border banking-claims spillover results in Section 5.2. Appendix F.1 reports matching balance, Appendix F.2 tests whether one treated lender–borrower cohort drives the result, and Appendix F.3 checks whether the result survives quarterly timing. Appendix F.4 reports a liability-side test and compares the LBS residence-based

estimates with CBS nationality-based claims.

F.1 Matching Balance

Table 28 reports the balance achieved by the cohort-by-cohort Mahalanobis matching used in the main BIS specification. The matching uses pre-treatment levels and slopes of log loans and deposits, the outcome closest to the main banking-claims channel. Standardized differences fall sharply after matching, including for auxiliary outcomes that were not used in the match. The stacked-cohort estimates therefore compare treated pairs with controls that had similar pre-treatment banking relationships, rather than with the full pool of unmatched country pairs.

Table 28: BIS Mahalanobis Matching Balance

	Treated	Pre-Match Pool		Matched Controls (k=20)	
	Mean	Mean	Std. Diff.	Mean	Std. Diff.
Pre-mean log(loans & deposits)*	8.228	5.013	1.136	7.975	0.114
Pre-slope log(loans & deposits)*	0.268	0.150	0.265	0.221	0.187
Pre-mean log(total claims)	8.433	5.228	1.111	8.218	0.095
Pre-mean log(non-bank loans)	6.881	4.407	1.008	6.964	0.045
Number of pairs (across cohorts)	12	765		172	

Notes: Cohort-aggregated balance for the matched stacked-cohort design used in Table 8. Pre-treatment statistics are computed within the three pre-treatment years before each cohort’s first-treatment date. “Treated” pools the 12 treated lender–borrower pairs; “Pre-Match Pool” pools the full set of eligible non-treated pairs in the same lender country (765 unique pairs across cohorts); “Matched Controls (k=20)” pools the cohort-specific Mahalanobis-matched controls (172 unique pairs across cohorts). Because controls can be reused across cohorts, the number of unique matched-control pairs is smaller than the mechanical upper bound of 240 cohort-control slots (12 cohorts times 20 controls), implying 68 control-pair reuses across cohorts. The matching routine uses up to 20 controls per cohort, and in the realized matched sample all 12 cohorts use 20 controls. Absolute standardized differences are computed as $|\bar{x}_{\text{treated}} - \bar{x}_{\text{control}}| / \sqrt{(s_{\text{treated}}^2 + s_{\text{control}}^2)/2}$. The two starred variables are used in the Mahalanobis distance; total claims and non-bank loans are auxiliary checks not used in the match. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

F.2 Leave-One-Cohort-Out

Figure 15 drops each treated lender–borrower cohort in turn, including that cohort’s matched controls. This is the direct concentration check for the BIS design: it asks whether the pooled coefficient is coming from a single treated pair. The loans coefficient remains positive in every omission. The estimates range from 0.0475 to 0.0768, compared with the full-sample estimate of 0.0566.




Figure 15: Leave-One-Cohort-Out: LBS Cross-Border Loans Spillover

Notes: Each row drops one treated lender-country $\times$ SWF-country cohort, including its matched controls, and re-estimates the intensity-weighted DiD from Table 8 (column 2: log loans & deposits). Orange dashed line: full-sample baseline ($\hat{\beta} = 0.056$). Error bars: 95% CIs; SE clustered by cohort. All estimates remain positive; the LOO range is [0.0475, 0.0768]. Eleven of twelve estimates are significant at the 5% level, and the remaining estimate is significant at the 10% level.

F.3 Quarterly Specification

Table 29 re-estimates the BIS design at quarterly frequency. The annual specification is the baseline because BIS stocks are cleaner at year end and match the main stacked-cohort window. The quarterly version is a timing check: similar estimates would mean the crowd-in pattern is not created by annual aggregation. Treatment timing uses the exact deal quarter when it is observed; in the raw post-filter treatment universe, 33 of 48 lender–borrower pairs have dated quarters and 15 are imputed to Q1 of the treatment year. The quarterly analog using all 12 matched cohorts gives a log-loans coefficient of 0.0515, close to the annual baseline of 0.0566. One matched cohort uses a Q1-imputed treatment quarter; dropping that cohort gives the same estimate at displayed precision, so treatment-quarter imputation is not driving the result.

Table 29: Cross-Border Banking Claims Spillover: Quarterly Specification

	(1) Total Claims	(2) Loans & Placed Dep.	(3) Non-Bank Loans
SWF Invest (\$bn) $\times$ Post	0.0531*** (0.0165)	0.0515** (0.0167)	0.0557*** (0.0127)
Outcome transform	log	log	log
Observations	12,276	12,270	12,217
Matched cohorts	12	12	12
Cohort $\times$ Quarter FE	✓	✓	✓
Cohort $\times$ Pair FE	✓	✓	✓
Cluster	Cohort	Cohort	Cohort

Note: Quarterly analog of the annual specification in Table 8. Same matched stacked-cohort design using BIS Locational Banking Statistics (residence-based). Treatment intensity: total SWF equity injection in \$bn. The treatment quarter equals the observed deal quarter when available and Q1 of the treatment year otherwise. Pre-matching window: 12 quarters. Control pairs are Mahalanobis-matched on pre-treatment log-loan levels and slopes. SE clustered by cohort. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

F.4 Liability Test and Nationality Robustness

Table 30 compares the main LBS residence-based estimates with CBS nationality-based claims. This check matters because LBS assigns claims to the location of the reporting banking office, not to the nationality of the banking group. For example, a loan booked by the London branch of a U.S. bank is a U.K. claim in LBS, not a U.S. claim. CBS instead follows the nationality of the parent banking group, so it provides a complementary check that the result is not created by the residence-based booking convention or by activity routed through international banking centers. Column (5) first reports a liability-side test based on BIS liabilities; the estimate is small and statistically insignificant, which is inconsistent with a symmetric expansion of both sides of the bilateral balance sheet. The CBS estimates are the same sign and of similar order: the total-claims

estimate is smaller than the LBS total-claims baseline (0.0426 versus 0.0581), while the CBS international-claims estimate is slightly larger (0.060). The comparison therefore suggests that the macro spillover result is not an artifact of the LBS reporting convention, while still leaving some scope for differences between residence and nationality measurement.

Table 30: Cross-Border Banking Claims Spillover: Liability Placebo and Nationality Robustness

	LBS (Residence)					CBS (Nationality)	
	(1) Total Claims	(2) Loans & Placed Dep.	(3) Non-Bank Loans	(4) Wins. Std. Flow Loans	(5) Liabilities	(6) Total Claims	(7) International Claims
SWF Invest (\$bn) $\times$ Post	0.0581** (0.0195)	0.0566** (0.0193)	0.0530*** (0.0151)	0.0252** (0.0113)	0.0225 (0.0197)	0.0426* (0.0210)	0.0599*** (0.0181)
Outcome transform	log	log	log	wins. flow/pre-stock	log	log	log
Observations	3,259	3,258	3,250	3,259	3,255	4,325	4,055
Matched cohorts	12	12	12	12	12	17	17
Cohort $\times$ Year FE	✓	✓	✓	✓	✓	✓	✓
Cohort $\times$ Pair FE	✓	✓	✓	✓	✓	✓	✓
Cluster	Cohort	Cohort	Cohort	Cohort	Cohort	Cohort	Cohort

Note: Same matched stacked-cohort design as Table 8 but at annual frequency (± 6 years), using Q4 end-of-year annual stocks. Control pairs are Mahalanobis-matched on pre-treatment log-loan levels and slopes. Columns (1)–(3) are BIS Locational Banking Statistics (residence-based) log stock outcomes: total claims, loans and placed deposits, and non-bank loans. Column (4) is the reported annual LBS loan flow, summed from quarterly flows, divided by pre-treatment mean loan stock and winsorized at the 1st and 99th percentiles. Column (5) is the log BIS liability-side position, used as a placebo outcome. Columns (6)–(7) are BIS Consolidated Banking Statistics (nationality-based, immediate counterparty) log stock outcomes: total claims and international claims. CBS total claims include local claims in local currency; CBS international claims exclude these, making column (7) the closest CBS analog to the LBS cross-border measure. The CBS estimates are the same sign and of similar order as the LBS baseline: CBS total claims are smaller, while CBS international claims are slightly larger. The CBS sample includes 17 cohorts (vs. 12 for LBS) because CBS covers additional reporting countries. SE clustered by cohort. Significance levels: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

References

- Anderson, James E., and Eric van Wincoop.** 2003. “Gravity with Gravitas: A Solution to the Border Puzzle.” *American Economic Review*, 93(1): 170–192.
- Avdjiev, Stefan, Leonardo Gambacorta, Linda S. Goldberg, and Stefano Schiaffi.** 2020. “The Shifting Drivers of Global Liquidity.” *Journal of International Economics*, 125: 103324.
- Bailey, Michael A., Anton Strezhnev, and Erik Voeten.** 2017. “Estimating Dynamic State Preferences from United Nations Voting Data.” *Journal of Conflict Resolution*, 61(2): 430–456.
- Bank for International Settlements.** 2019. “Reporting Guidelines for the BIS International Banking Statistics.” Bank for International Settlements Technical Report. Monetary and Economic Department, July 2019.
- Berger, Daniel, William Easterly, Nathan Nunn, and Shanker Satyanath.** 2013. “Commercial Imperialism? Political Influence and Trade During the Cold War.” *American Economic Review*, 103(2): 863–896.
- Bernstein, Shai, Josh Lerner, and Antoinette Schoar.** 2013. “The Investment Strategies of Sovereign Wealth Funds.” *Journal of Economic Perspectives*, 27(2): 219–238.
- Bortolotti, Bernardo, Veljko Fotak, and William L. Megginson.** 2015. “The Sovereign Wealth Fund Discount: Evidence from Public Equity Investments.” *The Review of Financial Studies*, 28(11): 2993–3035.
- Cetorelli, Nicola, and Linda S. Goldberg.** 2011. “Global Banks and International Shock Transmission: Evidence from the Crisis.” *IMF Economic Review*, 59(1): 41–76.
- Chaney, Thomas.** 2014. “The Network Structure of International Trade.” *American Economic Review*, 104(11): 3600–3634.
- Chhaochharia, Vidhi, and Luc Laeven.** 2008. “Sovereign Wealth Funds: Their Investment Strategies and Performance.” CEPR Discussion Paper 6959.
- Clayton, Christopher, Amanda Dos Santos, Matteo Maggiori, and Jesse Schreger.** 2025. “Internationalizing Like China.” *American Economic Review*, 115(3): 864–902.
- Clayton, Christopher, Matteo Maggiori, and Jesse Schreger.** 2025. “Putting Economics Back Into Geoeconomics.” NBER Working Paper 33681.

- Clayton, Christopher, Matteo Maggiori, and Jesse Schreger.** 2026. “A Framework for Geoeconomics.” *Econometrica*, 94(1): 105–136.
- Corsetti, Giancarlo, Amil Dasgupta, Stephen Morris, and Hyun Song Shin.** 2004. “Does One Soros Make a Difference? A Theory of Currency Crises with Large and Small Traders.” *The Review of Economic Studies*, 71(1): 87–113.
- Daude, Christian, and Marcel Fratzscher.** 2008. “The Pecking Order of Cross-border Investment.” *Journal of International Economics*, 74(1): 94–119.
- De Haas, Ralph, and Neeltje Van Horen.** 2013. “Running for the Exit? International Bank Lending During a Financial Crisis.” *The Review of Financial Studies*, 26(1): 244–285.
- Dewenter, Kathryn L., Xi Han, and Paul H. Malatesta.** 2010. “Firm Values and Sovereign Wealth Fund Investments.” *Journal of Financial Economics*, 98(2): 256–278.
- Egger, Peter, and Valeria Merlo.** 2007. “The Impact of Bilateral Investment Treaties on FDI Dynamics.” *The World Economy*, 30(10): 1536–1549.
- Farrell, Henry, and Abraham L. Newman.** 2019. “Weaponized Interdependence: How Global Economic Networks Shape State Coercion.” *International Security*, 44(1): 42–79.
- Hadlock, Charles J., and Joshua R. Pierce.** 2010. “New Evidence on Measuring Financial Constraints: Moving Beyond the KZ Index.” *The Review of Financial Studies*, 23(5): 1909–1940.
- Haug, Alfred A., Anh T. N. Nguyen, and P. Dorian Owen.** 2023. “Do the Determinants of Foreign Direct Investment Have a Reverse and Symmetric Impact on Foreign Direct Divestment?” *Empirical Economics*, 64(2): 659–680.
- Head, Keith, and John Ries.** 2008. “FDI as an Outcome of the Market for Corporate Control: Theory and Evidence.” *Journal of International Economics*, 74(1): 2–20.
- Head, Keith, and Thierry Mayer.** 2014. “Gravity Equations: Workhorse, Toolkit, and Cookbook.” In *Handbook of International Economics*. Vol. 4, , ed. Gita Gopinath, Elhanan Helpman and Kenneth Rogoff, 131–195. Elsevier.
- Holmström, Bengt, and Jean Tirole.** 1998. “Private and Public Supply of Liquidity.” *Journal of Political Economy*, 106(1): 1–40.

- International Monetary Fund.** 2008. “Global Financial Stability Report: Containing Systemic Risks and Restoring Financial Soundness.” International Monetary Fund, Washington, DC. Chapter 1, Box 1.2: SWFs provided \$41 billion of the \$105 billion in capital raised by major financial institutions between November 2007 and March 2008.
- Ivashina, Victoria, and David Scharfstein.** 2010. “Bank Lending During the Financial Crisis of 2008.” *Journal of Financial Economics*, 97(3): 319–338.
- Kabir, Poorya, Adrien Matray, Karsten Müller, and Chenzi Xu.** 2025. “Exim’s Exit: The Real Effects of Trade Financing by Export Credit Agencies.” NBER Working Paper 32019. Revised March 2025.
- Karolyi, G. Andrew, and Rose C. Liao.** 2017. “State Capitalism’s Global Reach: Evidence from Foreign Acquisitions by State-Owned Companies.” *Journal of Corporate Finance*, 42: 367–391.
- Khwaja, Asim Ijaz, and Atif Mian.** 2008. “Tracing the Impact of Bank Liquidity Shocks: Evidence from an Emerging Market.” *American Economic Review*, 98(4): 1413–1442.
- Knill, April M., Bong-Soo Lee, and Nathan Mauck.** 2012. “Bilateral Political Relations and Sovereign Wealth Fund Investment.” *Journal of Corporate Finance*, 18(1): 108–123.
- Kotter, Jason, and Ugur Lel.** 2011. “Friends or Foes? Target Selection Decisions of Sovereign Wealth Funds and Their Consequences.” *Journal of Financial Economics*, 101(2): 360–381.
- Meggison, William L., and Veljko Fotak.** 2015. “Rise of the Fiduciary State: A Survey of Sovereign Wealth Fund Research.” *Journal of Economic Surveys*, 29(4): 733–778.
- Pellegrino, Bruno, Enrico Spolaore, and Romain Wacziarg.** 2025. “Barriers to Global Capital Allocation.” *The Quarterly Journal of Economics*, 140(4): 3067–3131.
- Portes, Richard, and Hélène Rey.** 2005. “The Determinants of Cross-Border Equity Flows.” *Journal of International Economics*, 65(2): 269–296.
- Schmidheiny, Kurt, and Sebastian Siegloch.** 2023. “On Event Studies and Distributed-Lags in Two-Way Fixed Effects Models: Identification, Equivalence, and Generalization.” *Journal of Applied Econometrics*, 38(5): 695–713.
- Segal, Ilya.** 1999. “Contracting with Externalities.” *The Quarterly Journal of Economics*, 114(2): 337–388.

- Steenbergen, Victor, Yan Liu, Mauricio Pinzon Latorre, and Xiao’ou Zhu.** 2022. “The World Bank’s Harmonized Bilateral FDI Databases: Methodology, Trends, and Possible Use Cases.” World Bank. Background paper for the Global Investment Competitiveness Report 2021/2022.
- Stuart, Elizabeth A.** 2010. “Matching Methods for Causal Inference: A Review and a Look Forward.” *Statistical Science*, 25(1): 1–21.
- Sufi, Amir.** 2007. “Information Asymmetry and Financing Arrangements: Evidence from Syndicated Loans.” *The Journal of Finance*, 62(2): 629–668.
- Wei, Shang-Jin.** 2000. “How Taxing Is Corruption on International Investors?” *The Review of Economics and Statistics*, 82(1): 1–11.
- Zhang, Jinyuan, and Ran Zhou.** 2025. “The Geoeconomic Role of Sovereign Wealth Funds: Effects on Global Trade.” SSRN. Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5448434.